

Deep Learning II

Itthi Chatnuntawech

Nanoinformatics and Artificial Intelligence (NAI)

National Nanotechnology Center (NANOTEC)

National Science and Technology Development Agency (NSTDA)

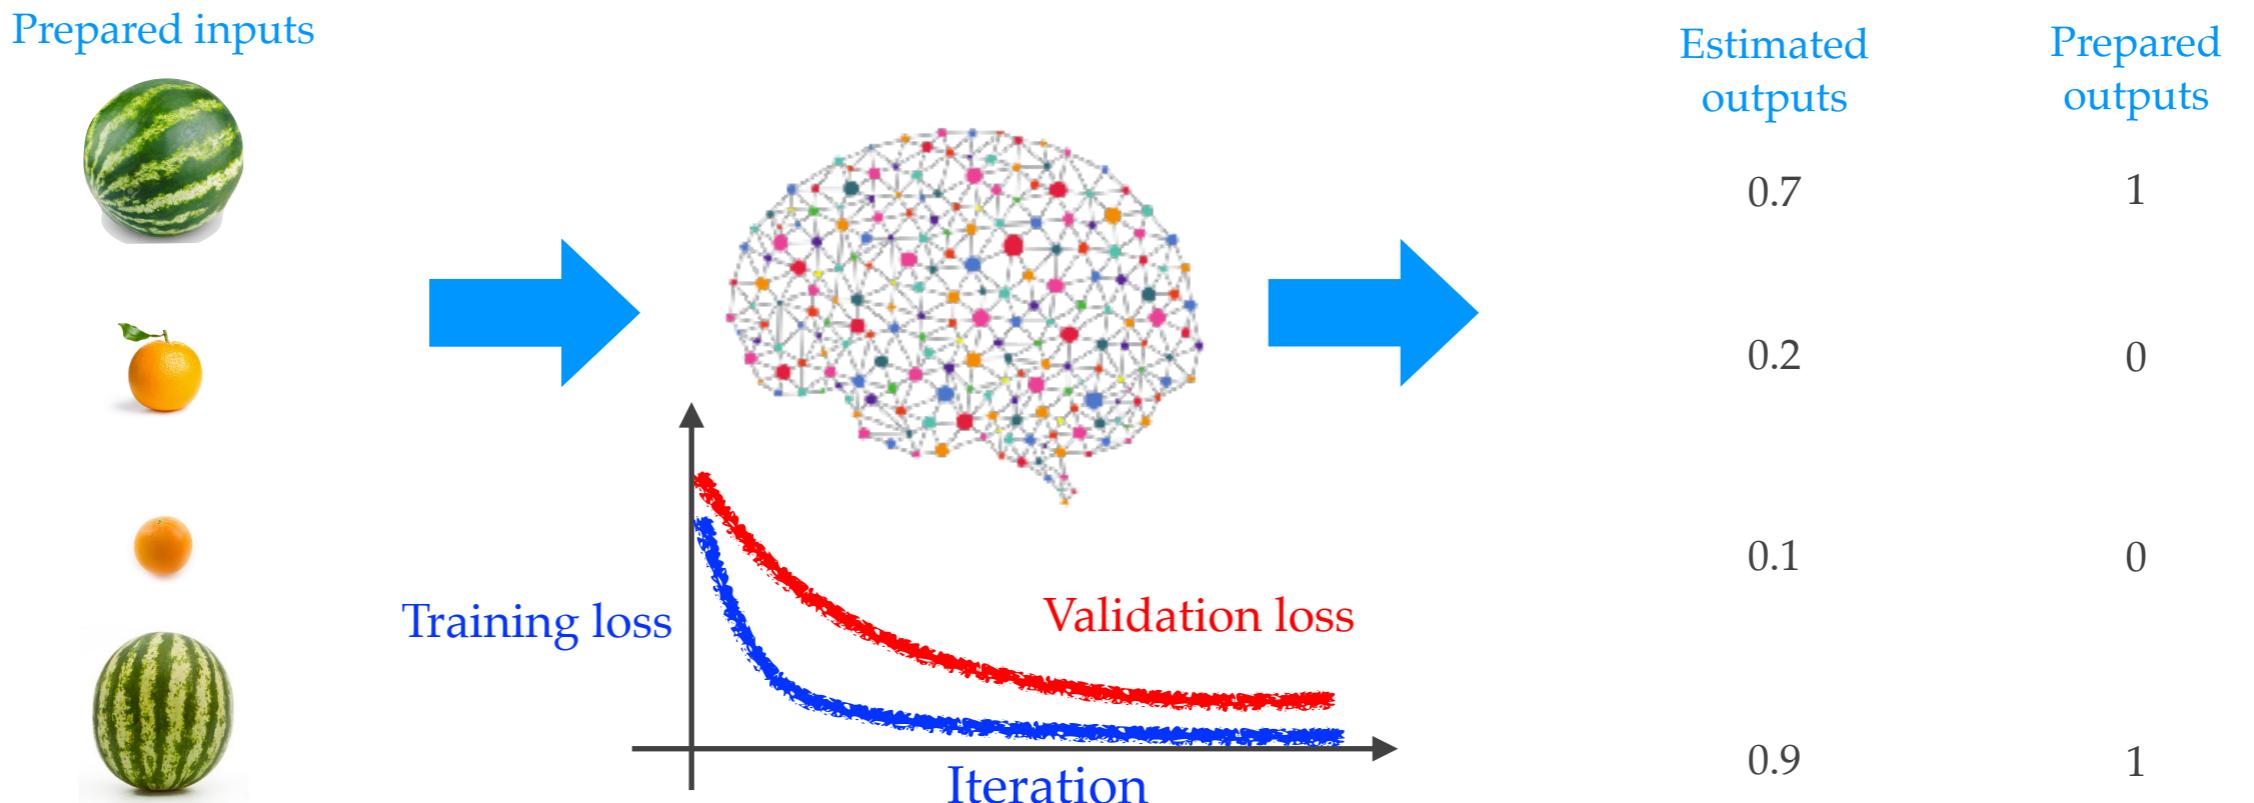
February 7, 2023

Outline

- ❖ Deep Learning Components
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)
- ❖ Hands-on: MRI Tumor Classification Using CNN
 - ❖ <https://github.com/ichatnun/CMU-medical-imaging-deep-learning>

From Last Time: Image Classification Using Deep Learning

- ❖ Training phase: Optimize the weights of a deep neural network



- ❖ Test phase: Perform prediction using the trained neural network





Fully connected layer (Dense)

Optimizer
SGD
Adam
RMSprop

Evaluation metric
accuracy
F1-score
AUC
confusion matrix



Convolutional layer
Conv1D, 2D, 3D, ...
separable Conv

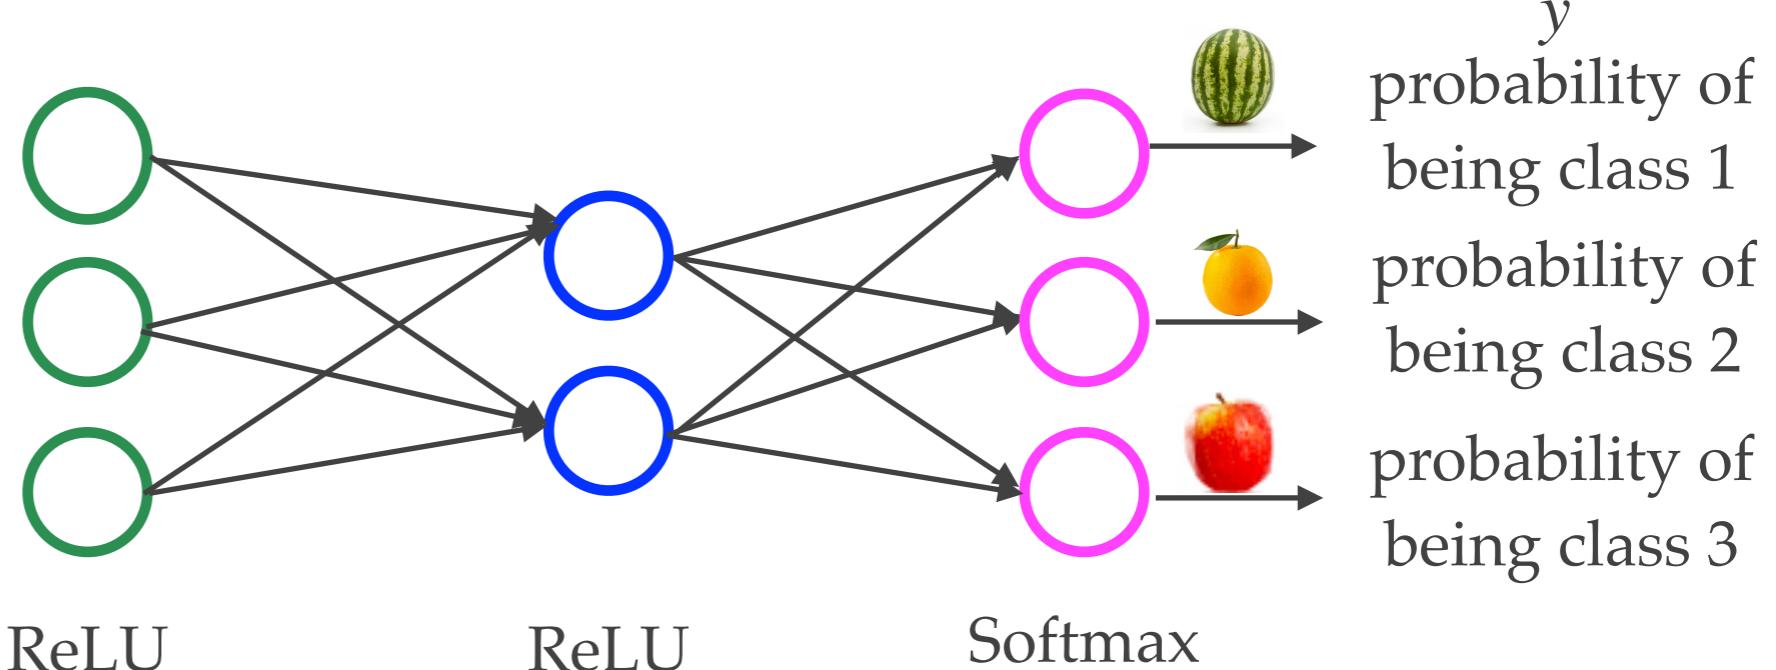
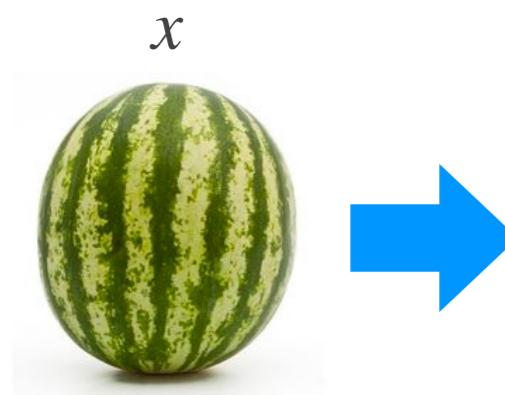
Pooling layer
max-pooling
average-pooling

Loss function
categorical crossentropy
binary crossentropy
mean squared error
mean absolute error

Regularization
Dropout
Data augmentation
 l_1, l_2 regularizations

Activation function
sigmoid
softmax
ESP (swish)
ReLU

- ❖ **Combine basic components to build a neural network**
 - More components → “More” representative power

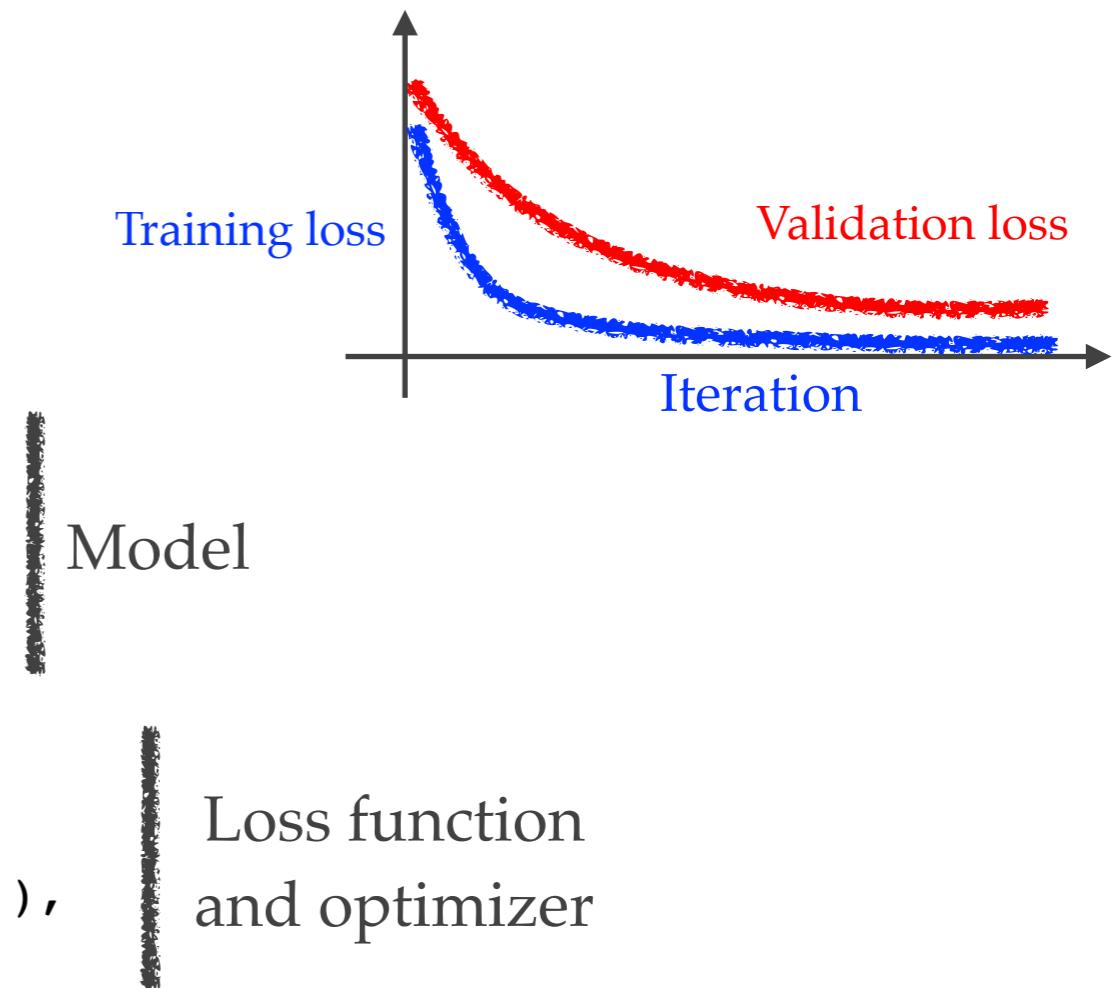


```
# Import necessary modules
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

# Create the model
model = keras.Sequential()
model.add(layers.Dense(3, activation="relu"))
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(3, activation="softmax"))

# Compile the model
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(),
)

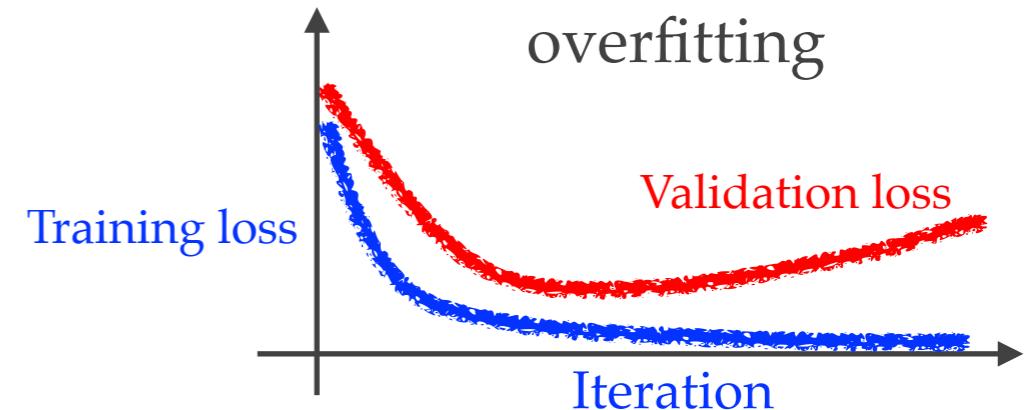
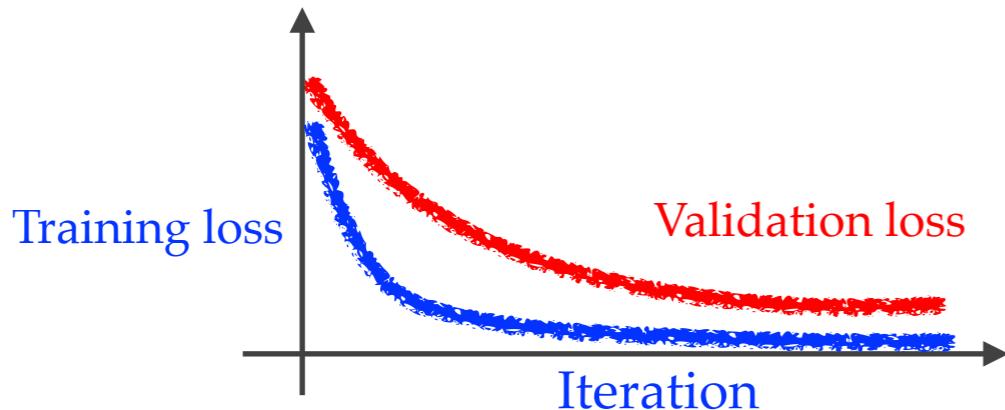
# Train the model for 100 epochs with a batch size of 32
model.fit(x_train, y_train, batch_size=32, epochs=100, validation_data=(x_val,y_val))
```



Outline

- ❖ Deep Learning Components
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)
- ❖ Hands-on: MRI Tumor Classification Using CNN
 - ❖ <https://github.com/ichatnun/CMU-medical-imaging-deep-learning>

Regularization



- ❖ Regularization is frequently used to mitigate overfitting
 - ❖ Add a regularization term to the loss function
 - ❖ ℓ_2 regularization

$$\min \sum_{i=1}^N L(y_i, f_W(x)) \rightarrow \min \sum_{i=1}^N L(y_i, f_W(x)) + \lambda \|w\|_2$$

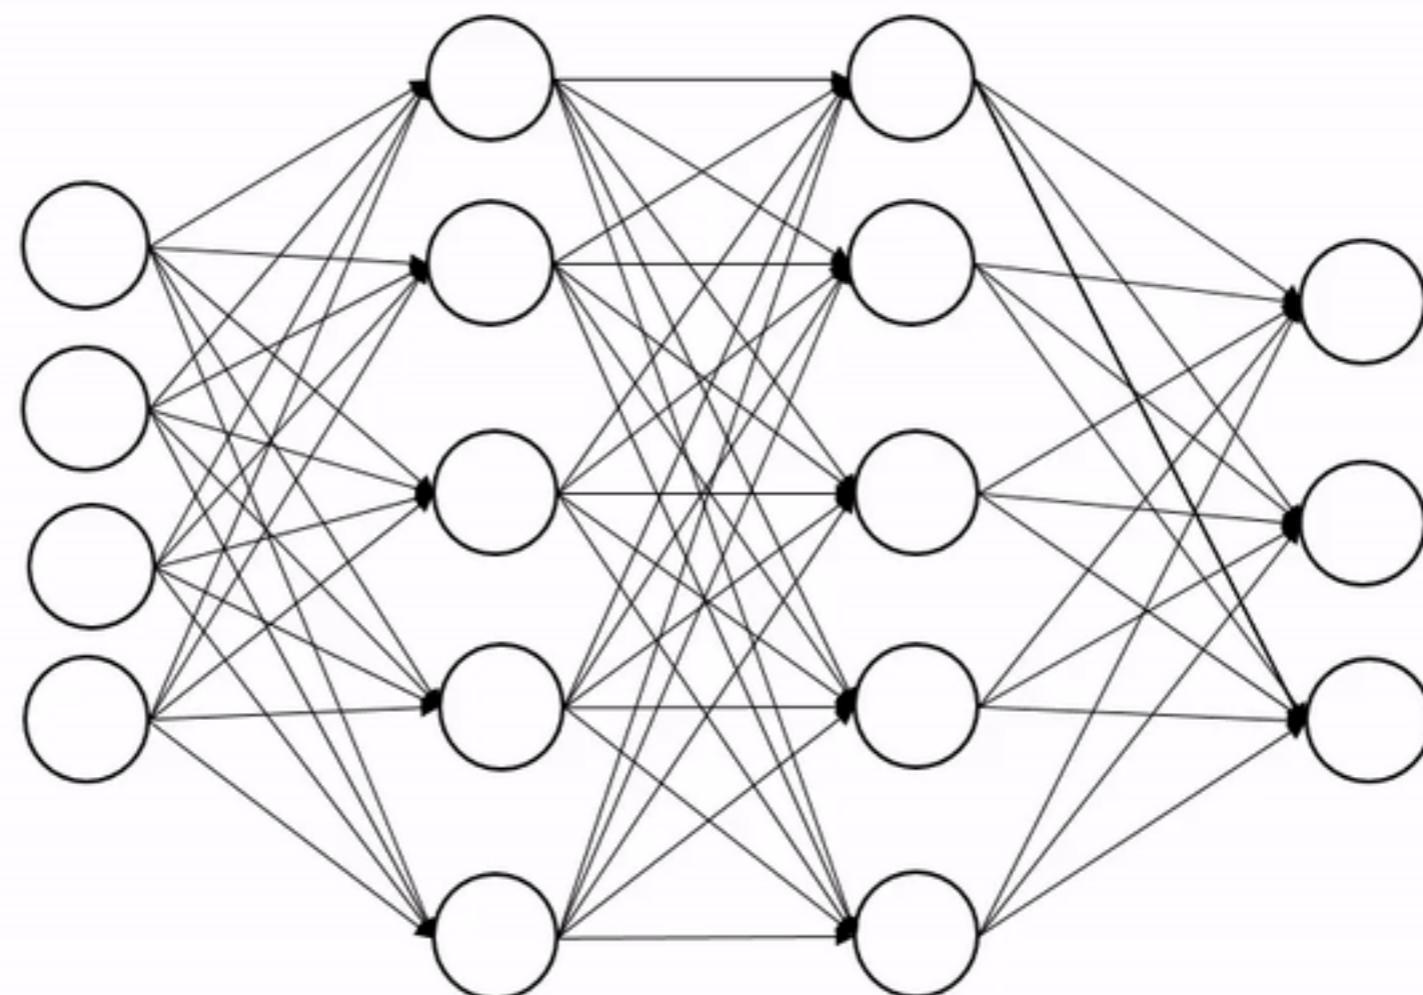
- ❖ ℓ_1 regularization

$$\min \sum_{i=1}^N L(y_i, f_W(x)) \rightarrow \min \sum_{i=1}^N L(y_i, f_W(x)) + \lambda \|w\|_1$$

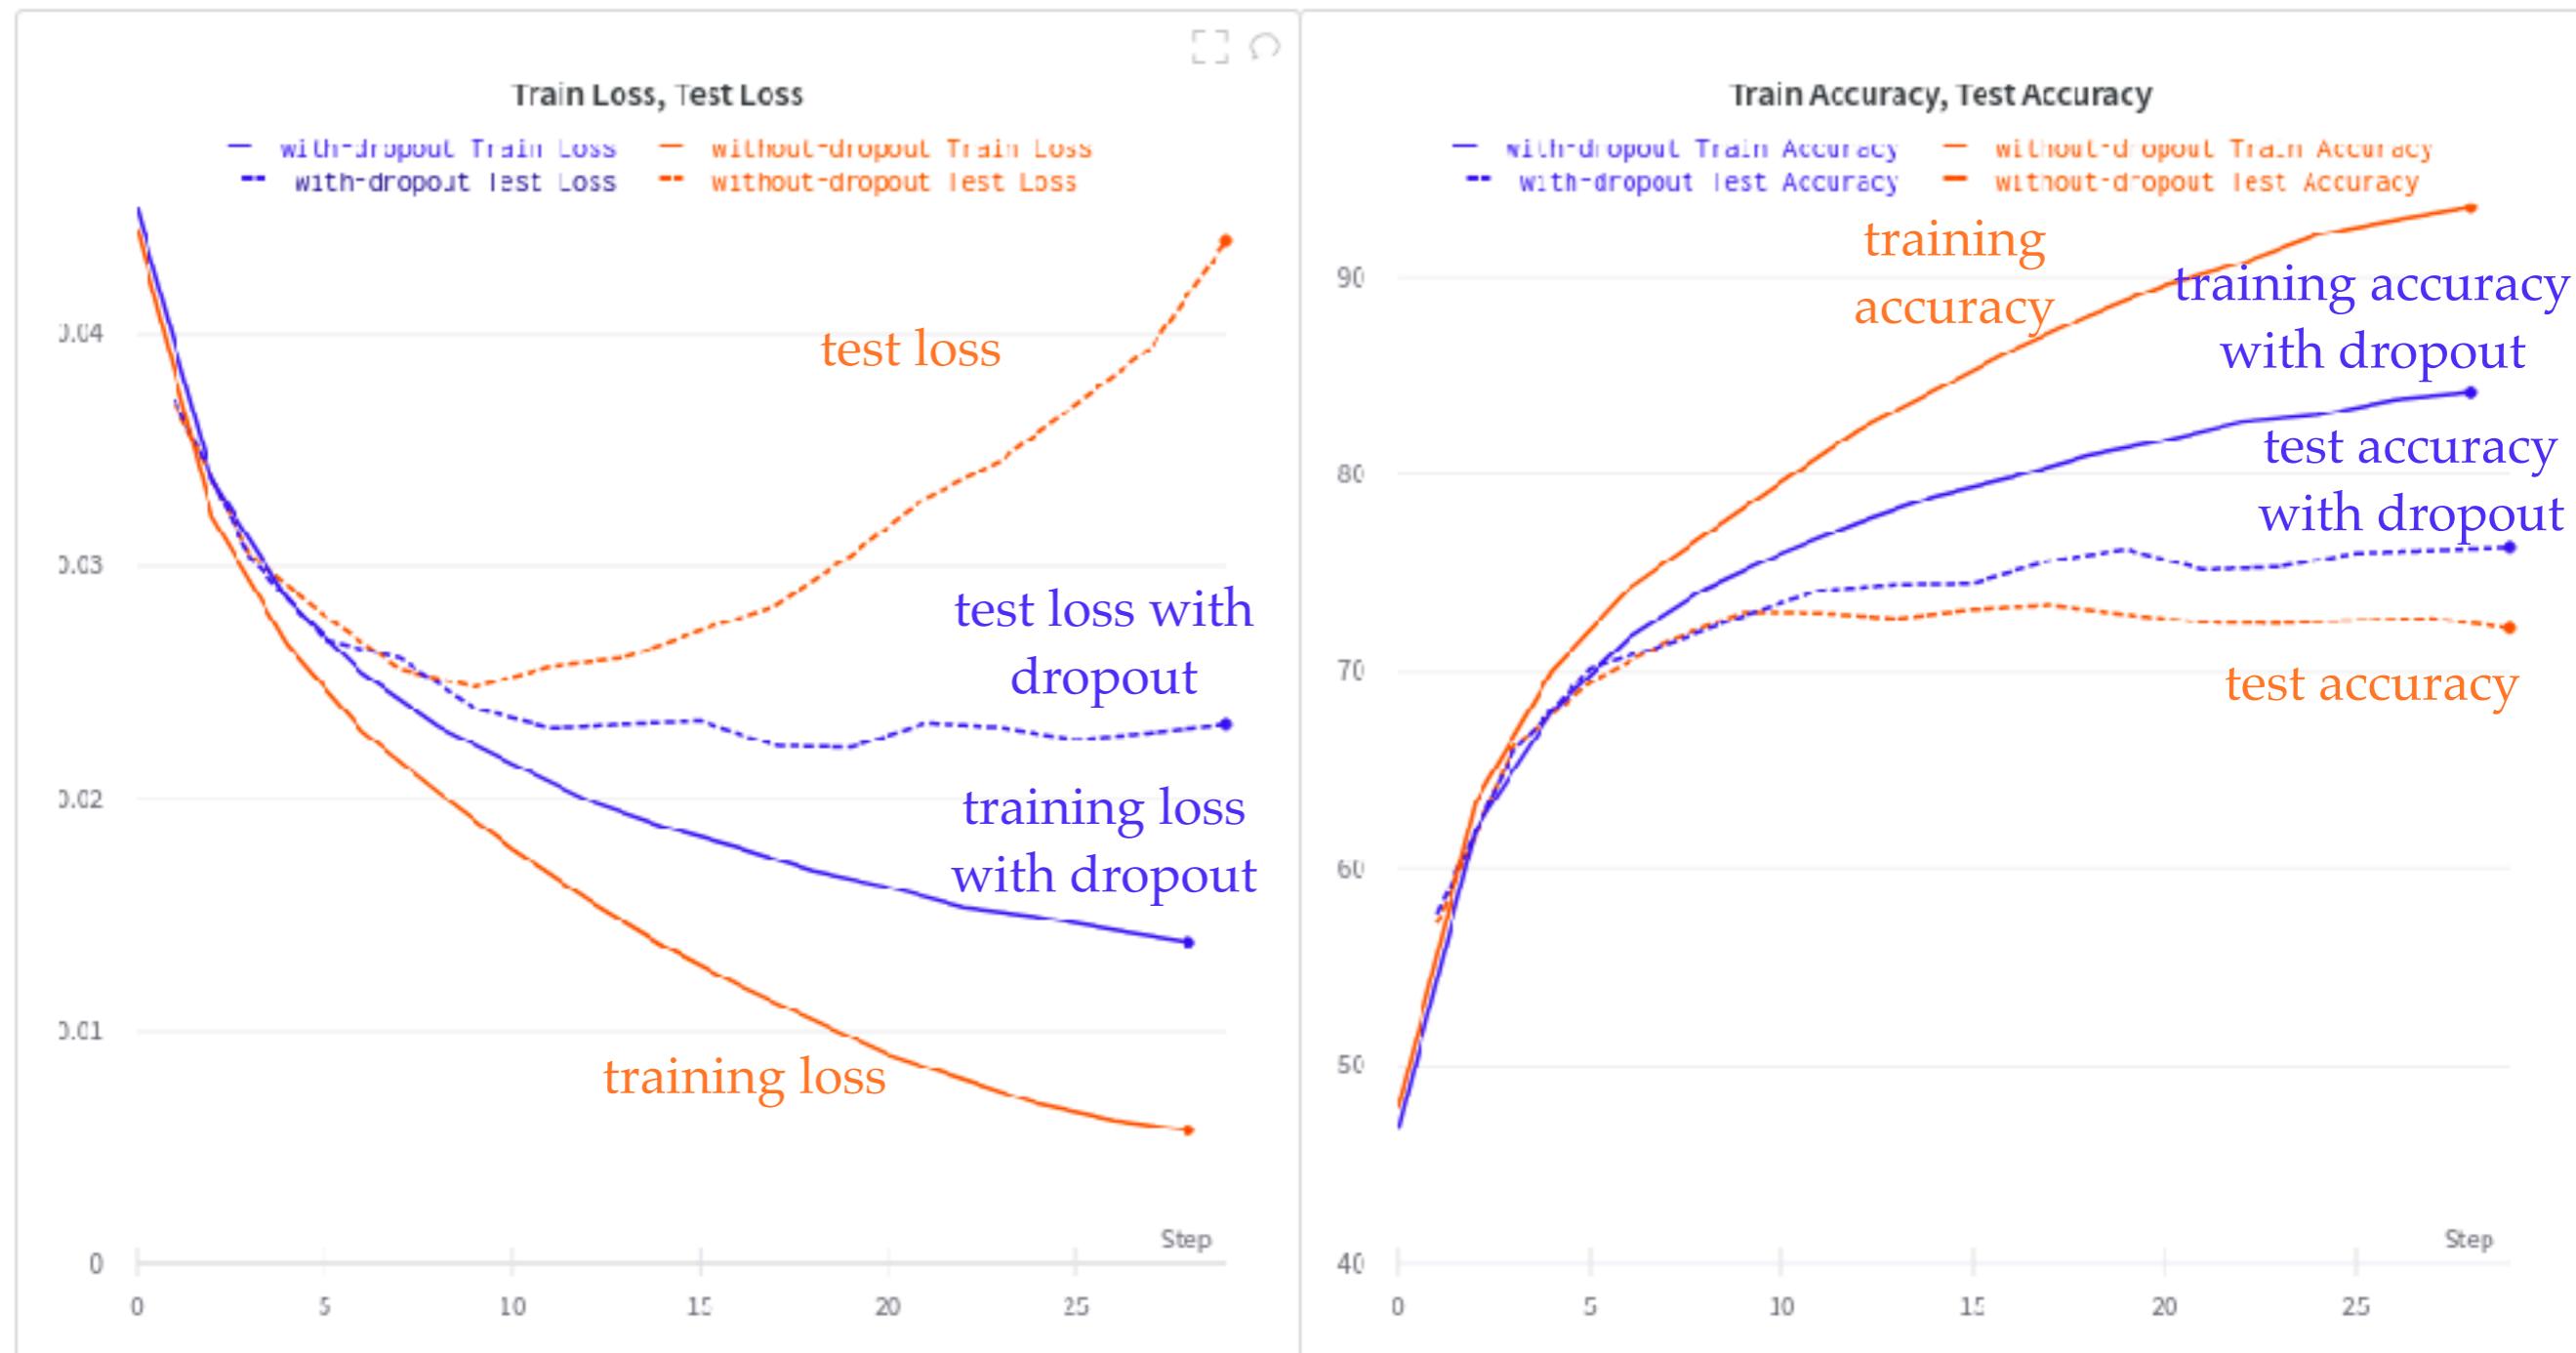
- ❖ Use dropout - much more popular

Dropout

- ❖ Dropout can be used to reduce overfitting
 - ❖ Randomly omit some neurons / units

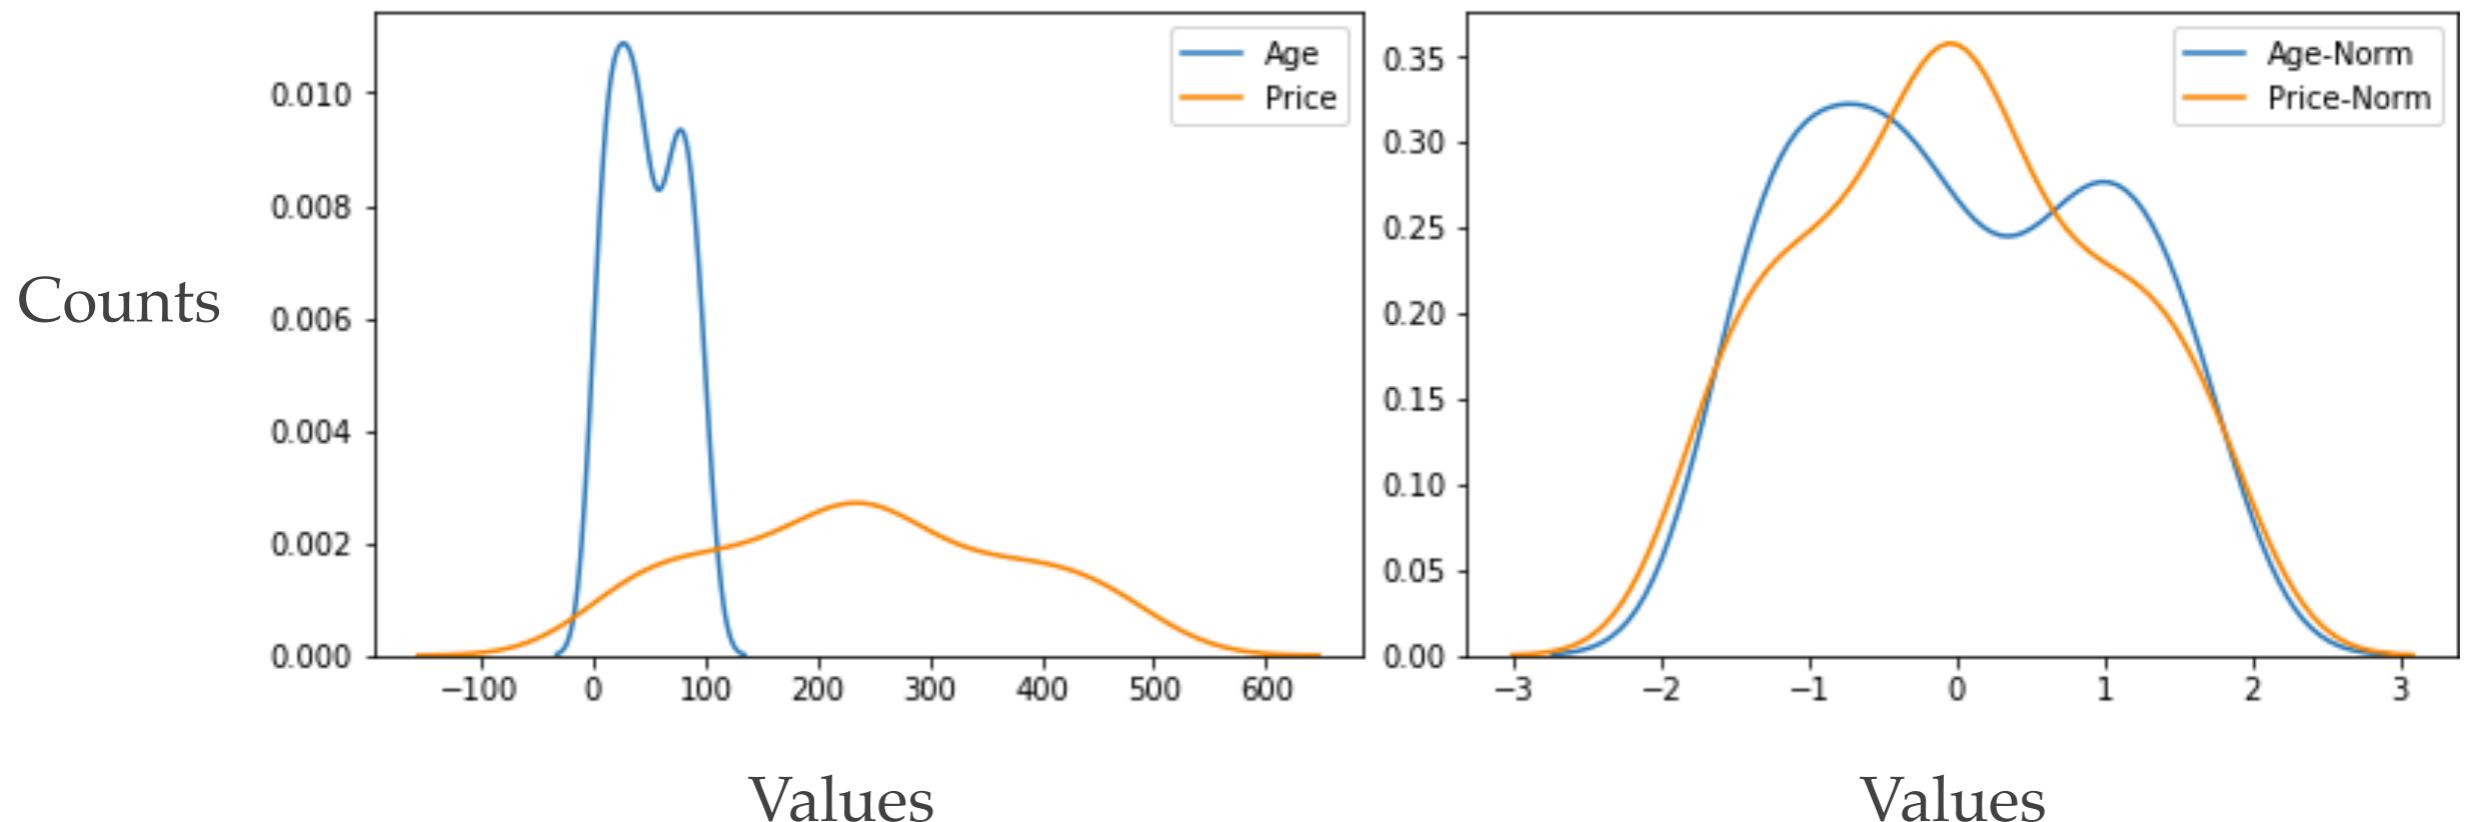


Dropout



Batch Normalization

- ❖ Normalize every batch
 - ❖ reduce internal covariate shift problems*
- ❖ Can be thought of as a form of implicit regularization
 - ❖ can help reduce overfitting

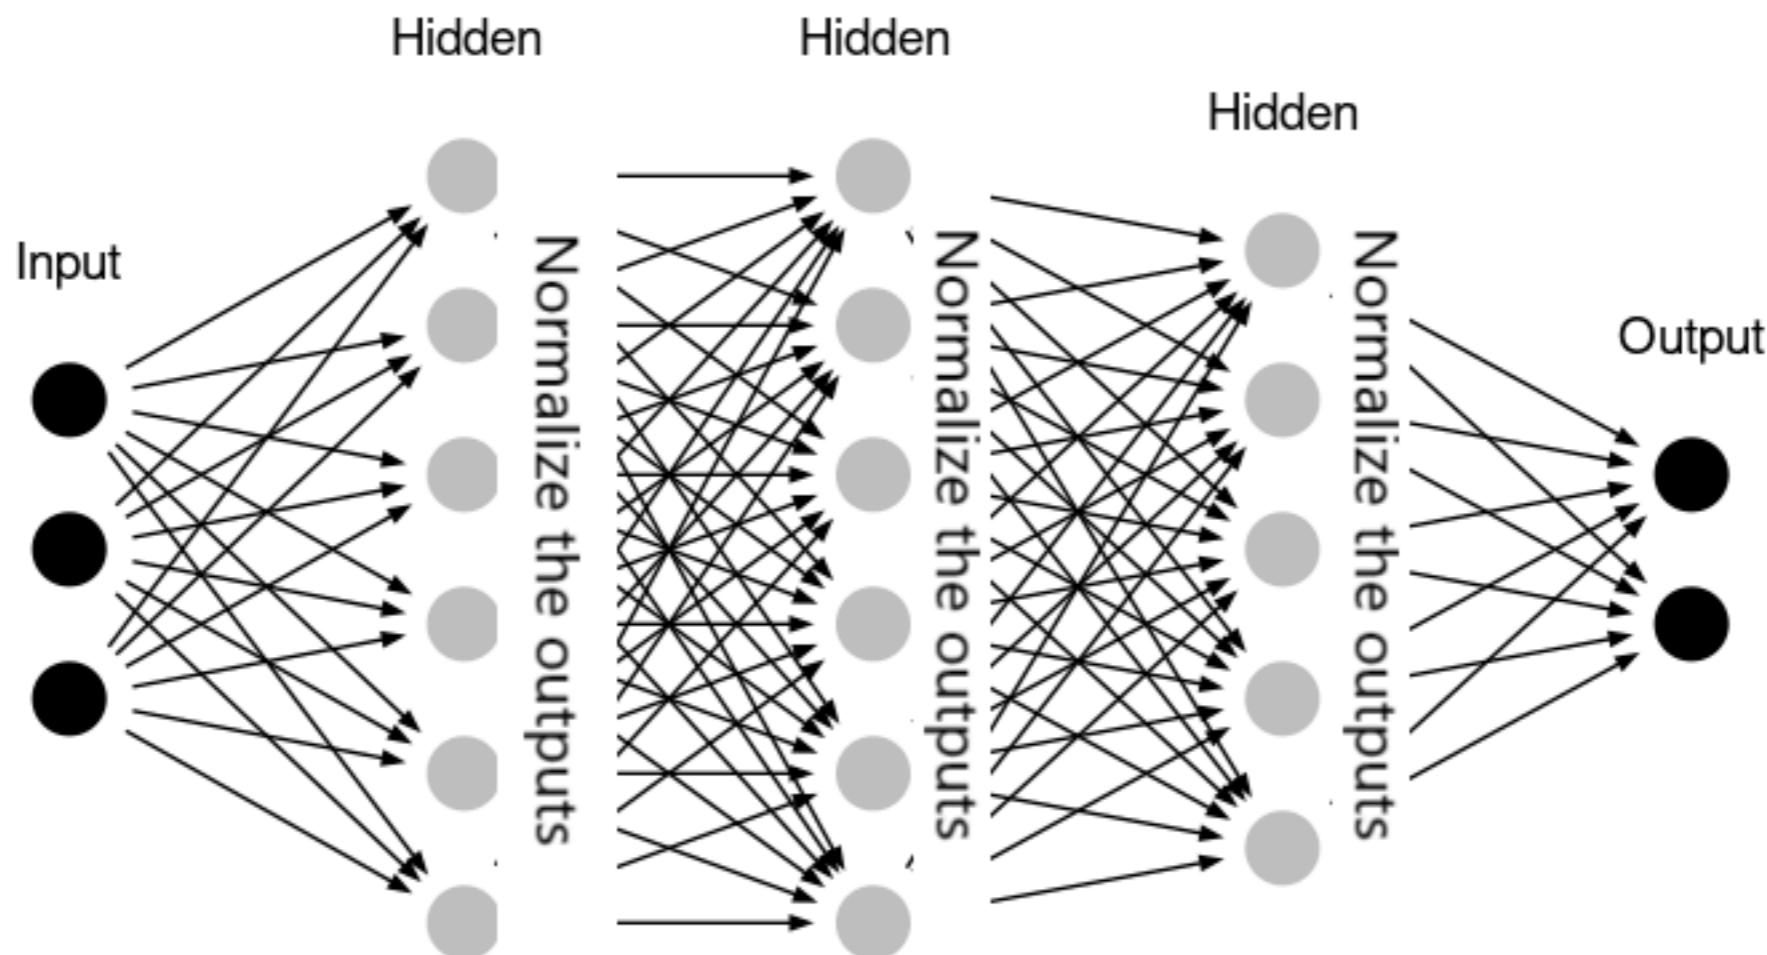


Batch Normalization — Speed up Neural Network Training

*This has been challenged by more recent work

Batch Normalization

- ❖ Normalize every batch
 - ❖ reduce internal covariate shift problems*
- ❖ Can be thought of as a form of implicit regularization
 - ❖ can help reduce overfitting



Batch Normalization — Speed up Neural Network Training

*This has been challenged by more recent work

Batch Normalization

- ❖ Reduce internal covariate shift problems*

batch size # of features

Input: $x : N \times D$

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Learnable params:

$\gamma, \beta : D$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Intermediates: $\mu, \sigma : D$
 $\hat{x} : N \times D$

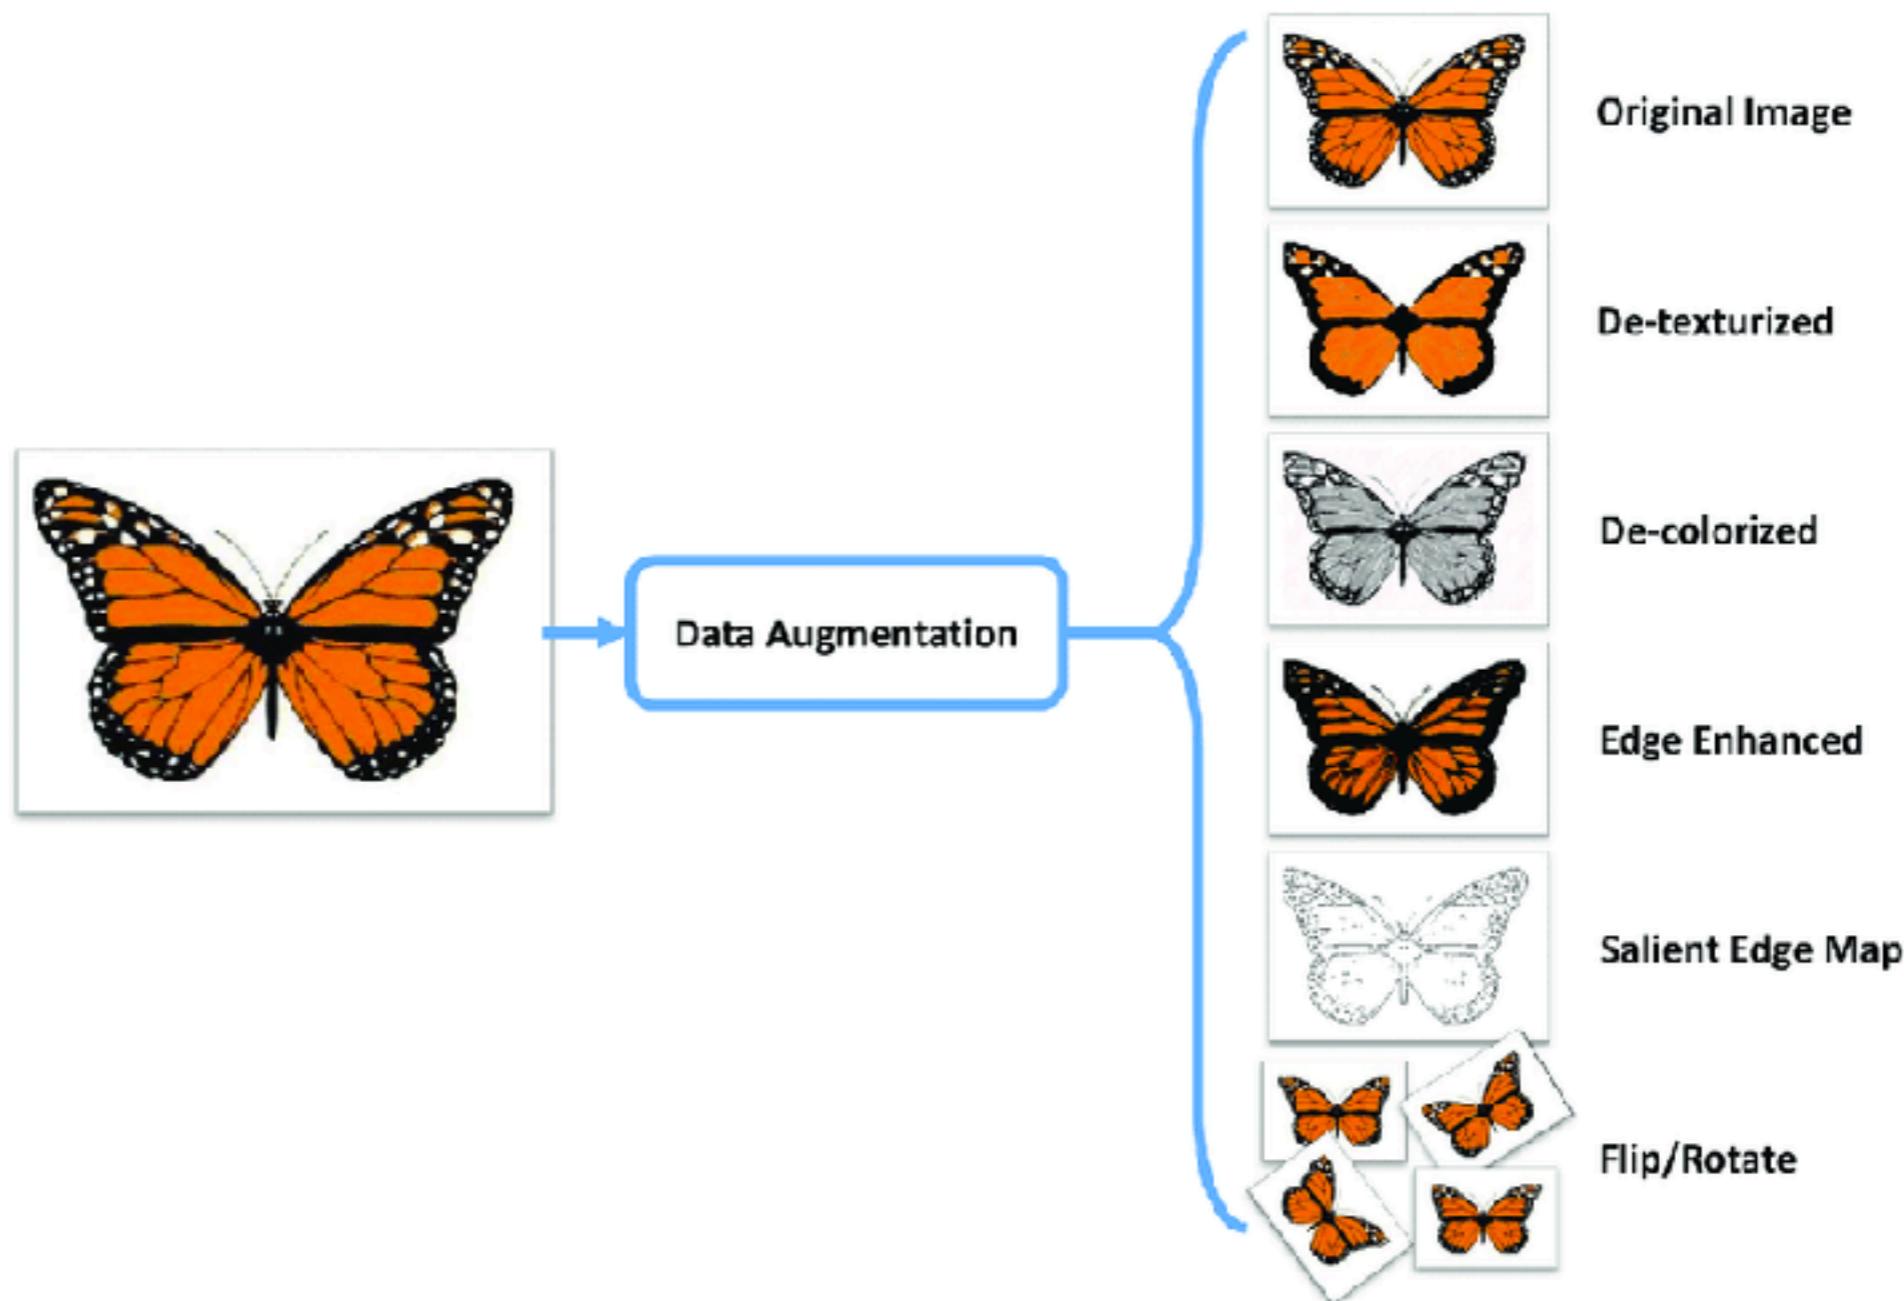
$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Output: $y : N \times D$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

- ❖ More robust to bad initialization

Data Augmentation



Outline

- ❖ Deep Learning Components
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)
- ❖ Hands-on: MRI Tumor Classification Using CNN
 - ❖ <https://github.com/ichatnun/CMU-medical-imaging-deep-learning>

Fully connected layer (Dense)

Convolutional layer
Conv1D, 2D, 3D, ...
separable Conv

Optimizer

SGD

Adam

RMSprop

Evaluation metric

accuracy

F1-score

AUC

confusion matrix

Loss function

categorical crossentropy

binary crossentropy

mean squared error

mean absolute error

Deep Learning

Pooling layer
max-pooling
average-pooling

Activation function

sigmoid

softmax

ESP (swish)

ReLU

Regularization

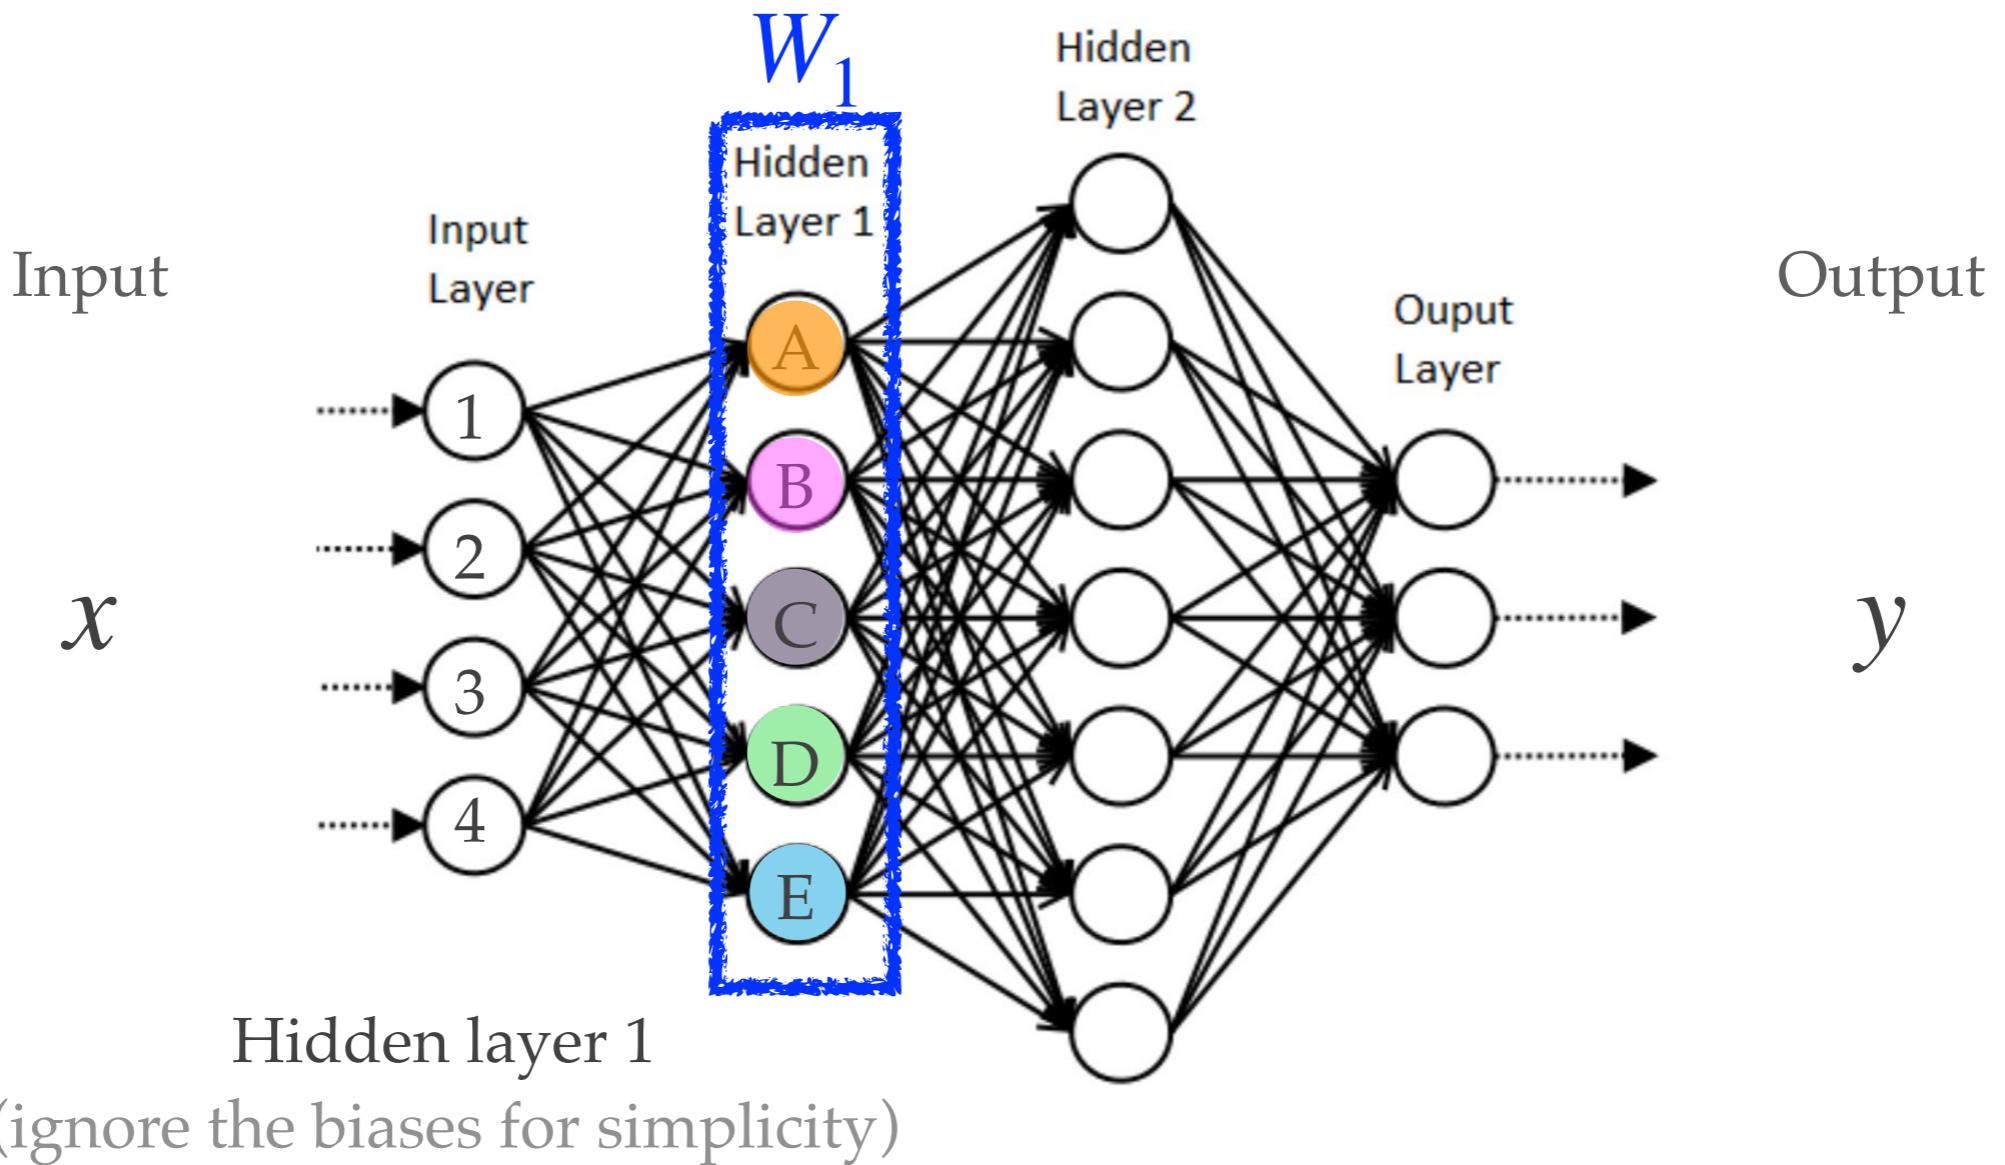
Dropout

Data augmentation

l_1, l_2 regularizations



- ❖ **Combine basic components to build a neural network**
 - More components → “More” representative power



$$\begin{bmatrix} out_A \\ out_B \\ out_C \\ out_D \\ out_E \end{bmatrix} = a_1(W_1x) = a_1 \left(\begin{bmatrix} w_{A,1} & w_{A,2} & w_{A,3} & w_{A,4} \\ w_{B,1} & w_{B,2} & w_{B,3} & w_{B,4} \\ w_{C,1} & w_{C,2} & w_{C,3} & w_{C,4} \\ w_{D,1} & w_{D,2} & w_{D,3} & w_{D,4} \\ w_{E,1} & w_{E,2} & w_{E,3} & w_{E,4} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} \right)$$

of weights for hidden layer 1
 $= 4 \times 5 = 20$

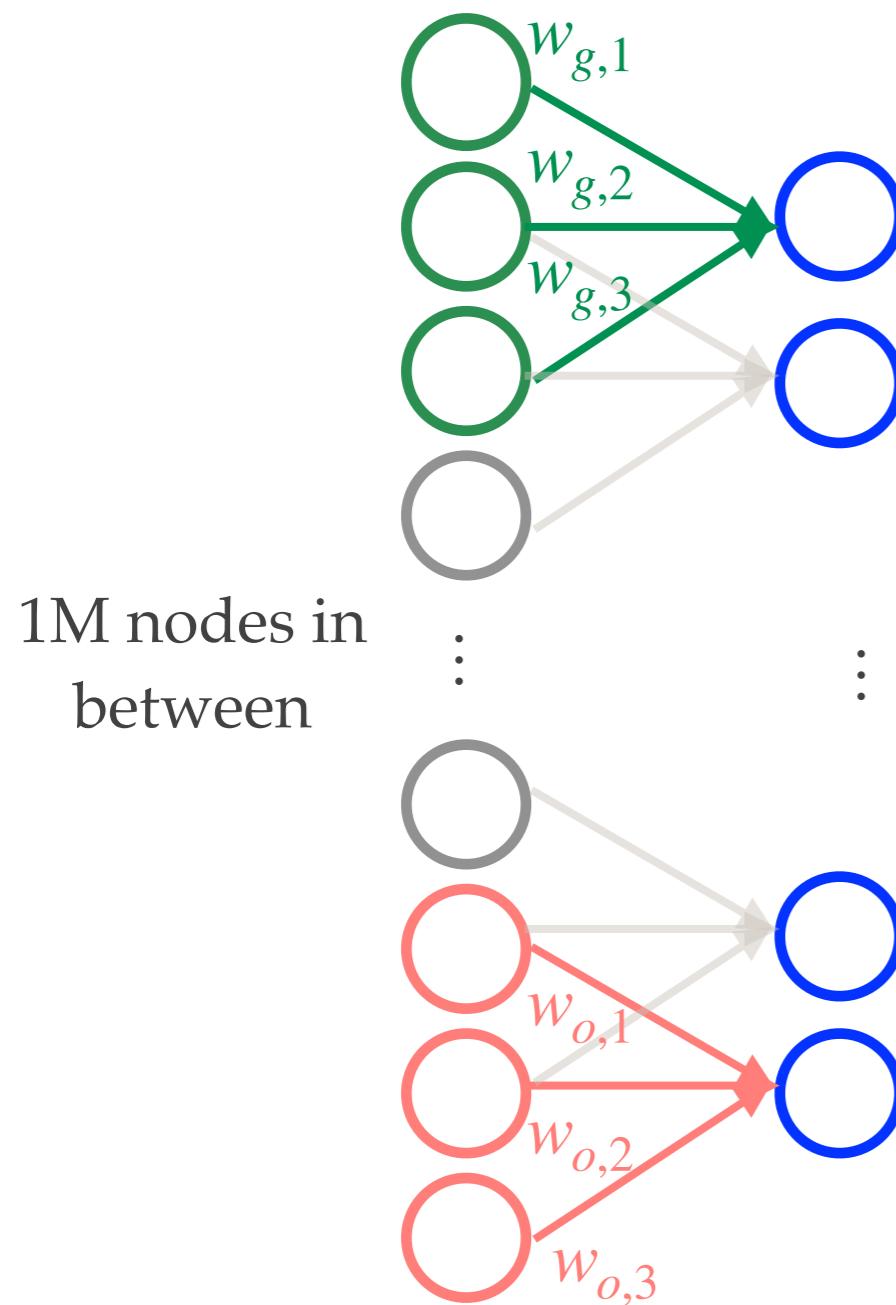
of weights per layer = # of incoming nodes x # of hidden nodes

What if our input dimension is $256 \times 256 = 65,536$ and we use 1024 hidden nodes?

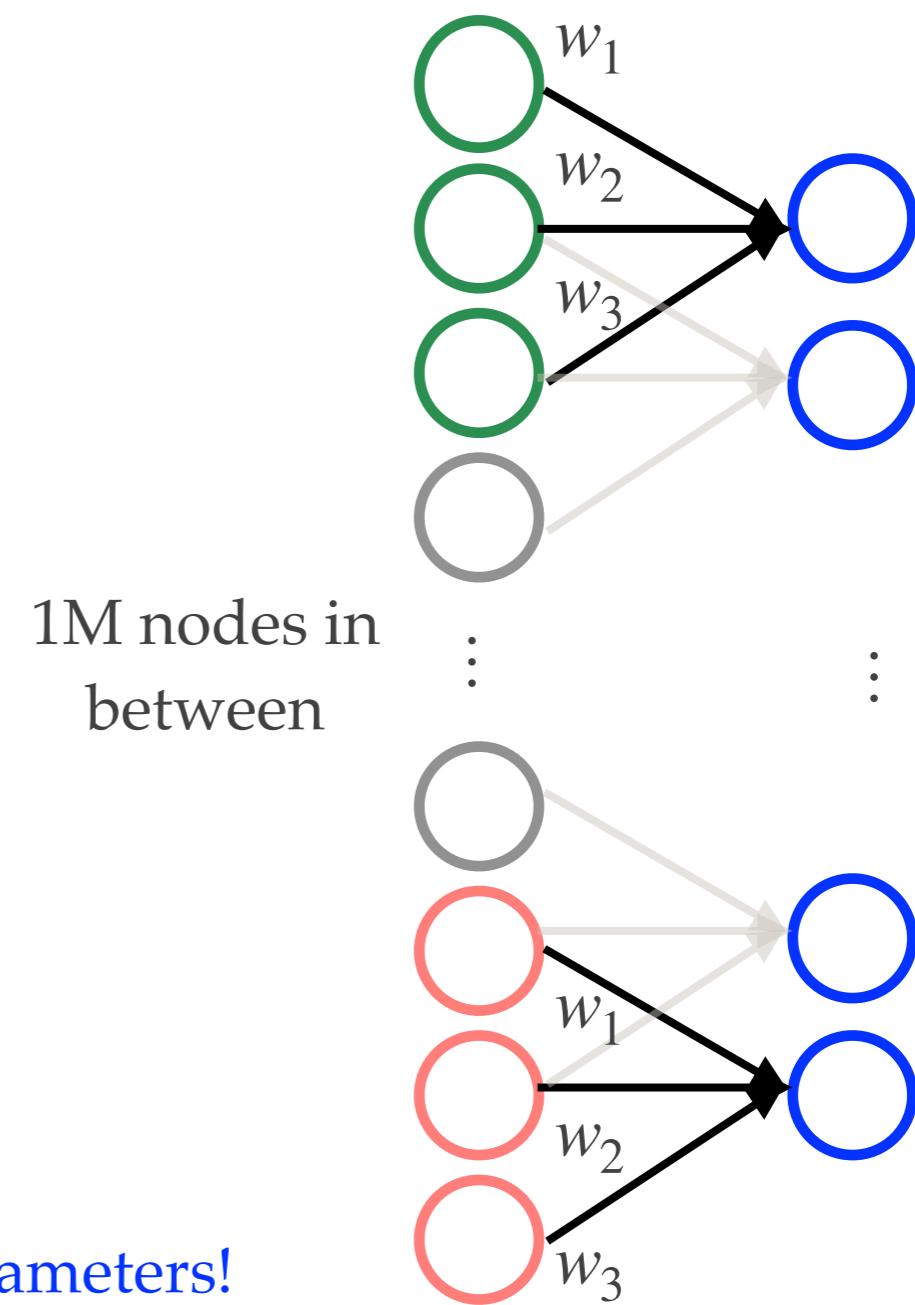
67 M trainable parameters for just one layer!

Do we need that many parameters?

- ❖ Would it be the case that
 - ❖ Only neighboring neurons talk to each other?
 - ❖ Furthermore, they talk to their neighbors in the same way?

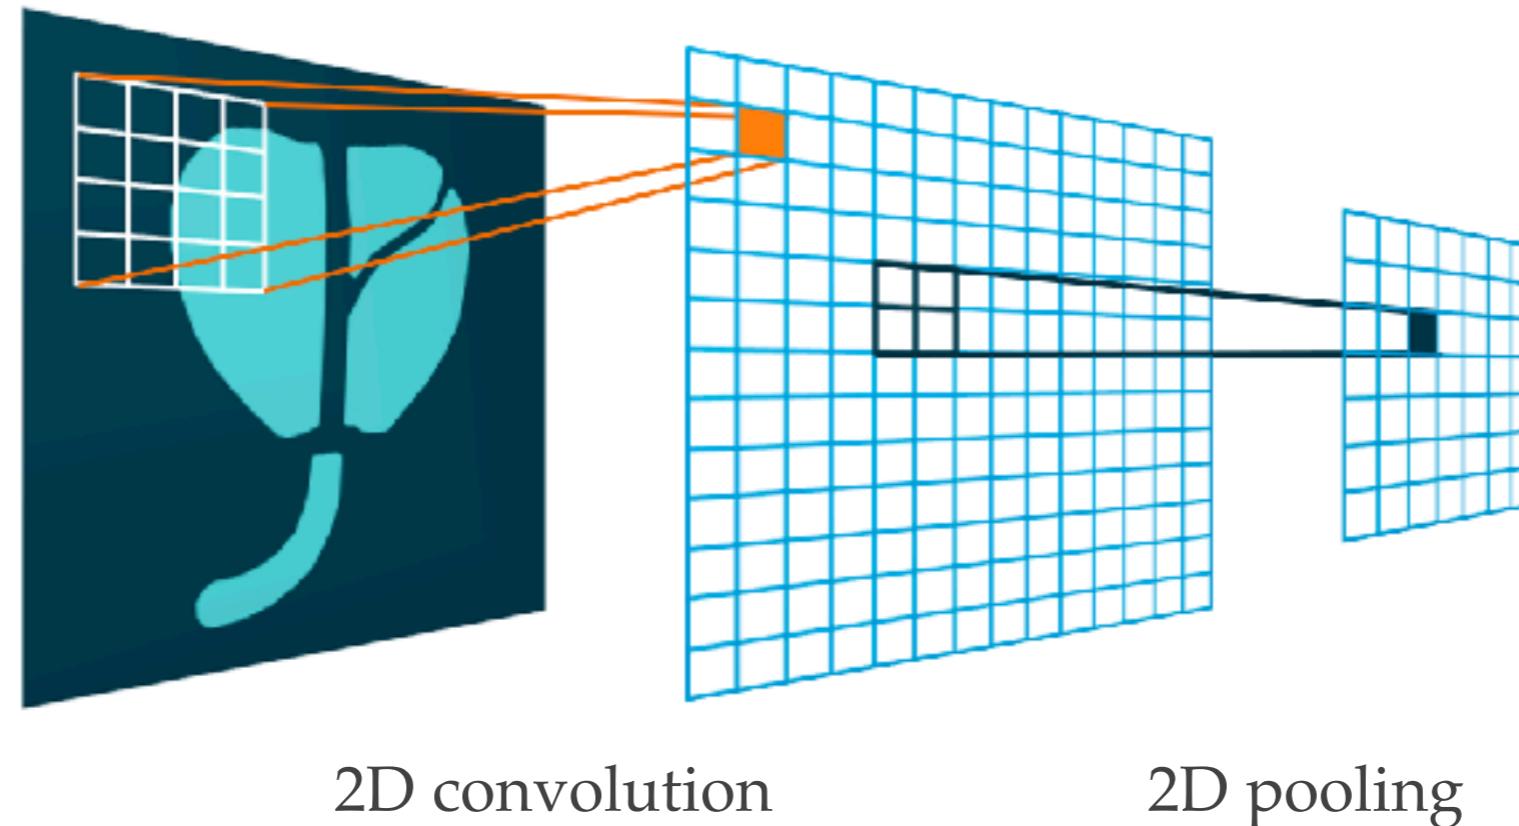


Much fewer parameters!



2D Convolution

- ❖ Similarly, for images, maybe only neighboring pixels talk to each other?



2D Convolution

$$y[n_1, n_2] = x[n_1, n_2] * h[n_1, n_2] = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h[n_1 - k_1, n_2 - k_2]$$

↑
2D Conv

$$0^*0 + 2^*1 + 1^*0 + 0^*1 + 1^*1 + 1^*1 + -3^*0 + -1^*1 + 1^*0$$

0	2	1	1	0
0	1	1	1	0
-3	-1	1	0	1
0	6	0	0	1
0	4	1	7	0

Input

x

0	1	0
1	1	1
0	1	0

h

Filter/kernel

defines the relationship
between neighboring pixels

		3		

y

2D Convolution

$$y[n_1, n_2] = x[n_1, n_2] * h[n_1, n_2] = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h[n_1 - k_1, n_2 - k_2]$$

↑
2D Conv

$$2*0 + 1*1 + 1*0 + 1*1 + 1*1 + 1*1 + -1*0 + 1*1 + 0*0$$

0	2	1	1	0
0	1	1	1	0
-3	-1	1	0	1
0	6	0	0	1
0	4	1	7	0

Input

x

0	1	0
1	1	1
0	1	0

h

Filter/kernel

defines the relationship
between neighboring pixels

3	5

y

2D Convolution

$$y[n_1, n_2] = x[n_1, n_2] * h[n_1, n_2] = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h[n_1 - k_1, n_2 - k_2]$$

↑
2D Conv

$$1^*0 + 1^*1 + 0^*0 + 1^*1 + 1^*1 + 0^*1 + 1^*0 + 0^*1 + 1^*0$$

0	2	1	1	0
0	1	1	1	0
-3	-1	1	0	1
0	6	0	0	1
0	4	1	7	0

Input

0	1	0
1	1	1
0	1	0

Filter / kernel

defines the relationship
between neighboring pixels

3	5	3

y

2D Convolution

$$y[n_1, n_2] = x[n_1, n_2] * h[n_1, n_2] = \sum_{k_1=-\infty}^{\infty} \sum_{k_2=-\infty}^{\infty} x[k_1, k_2] h[n_1 - k_1, n_2 - k_2]$$

↑
2D Conv

0	2	1	1	0
0	1	1	1	0
-3	-1	1	0	1
0	6	0	0	1
0	4	1	7	0

Input

x

0	1	0
1	1	1
0	1	0

h

Filter / kernel

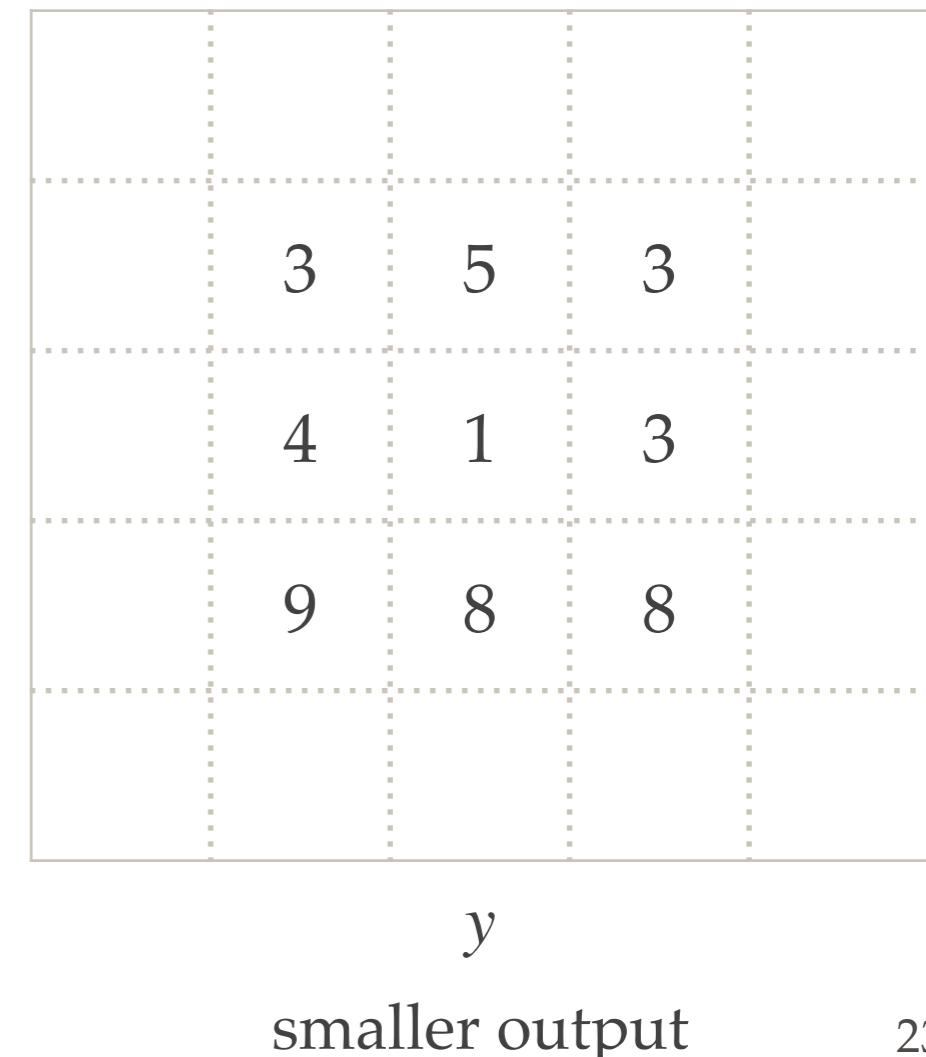
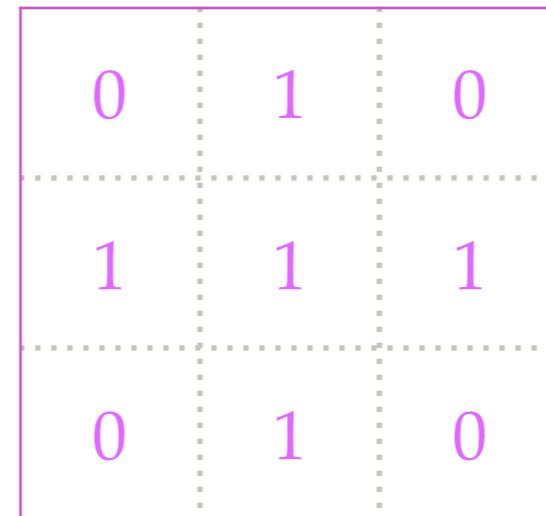
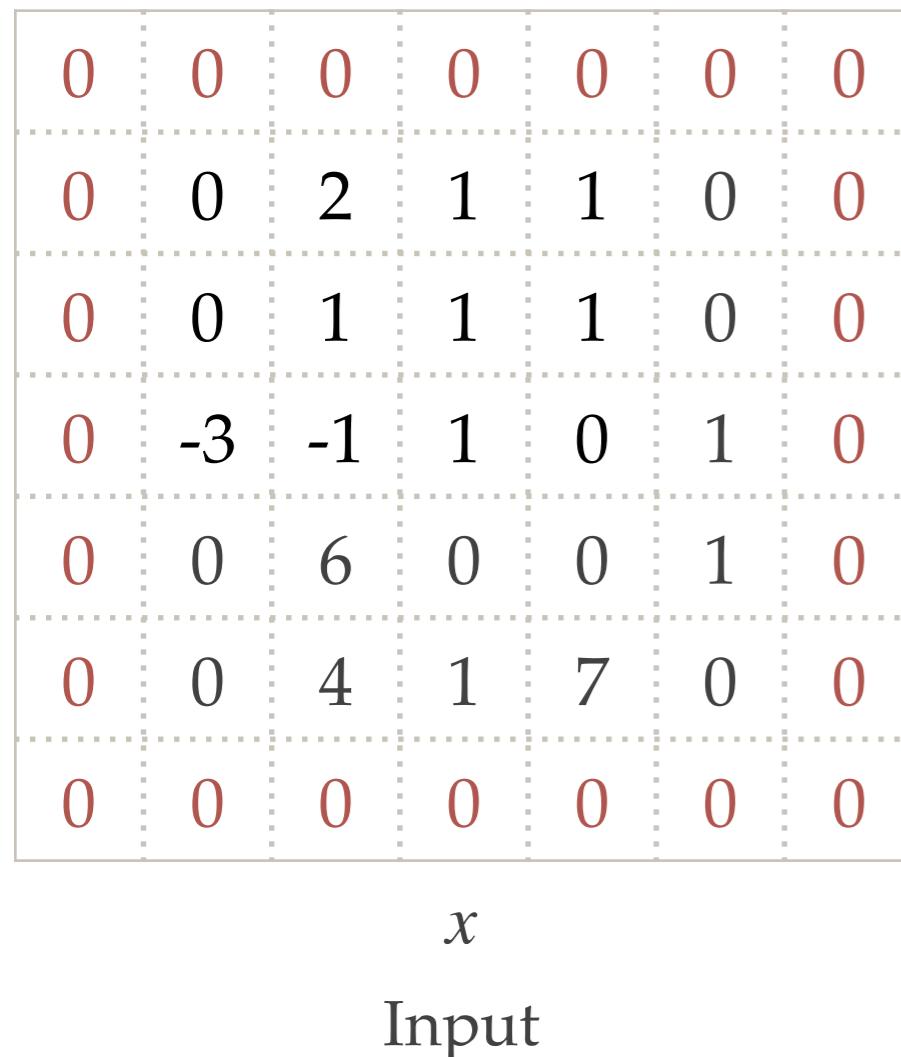
defines the relationship
between neighboring pixels

smaller output

y

2D Convolution

Zero-padding



defines the relationship
between neighboring pixels

2D Convolution

Zero-padding

x	0	0	0	0	0	0	0
0	0	0	2	1	1	0	0
0	0	1	1	1	0	0	0
0	-3	-1	1	0	1	0	0
0	0	6	0	0	1	0	0
0	0	4	1	7	0	0	0
0	0	0	0	0	0	0	0

0	1	0
1	1	1
0	1	0

h

Filter/kernel

defines the relationship
between neighboring pixels

2	4	5	3	1
-2	3	5	3	2
-4	4	1	3	2
3	9	8	8	2
4	11	12	8	8

y

output of the same size

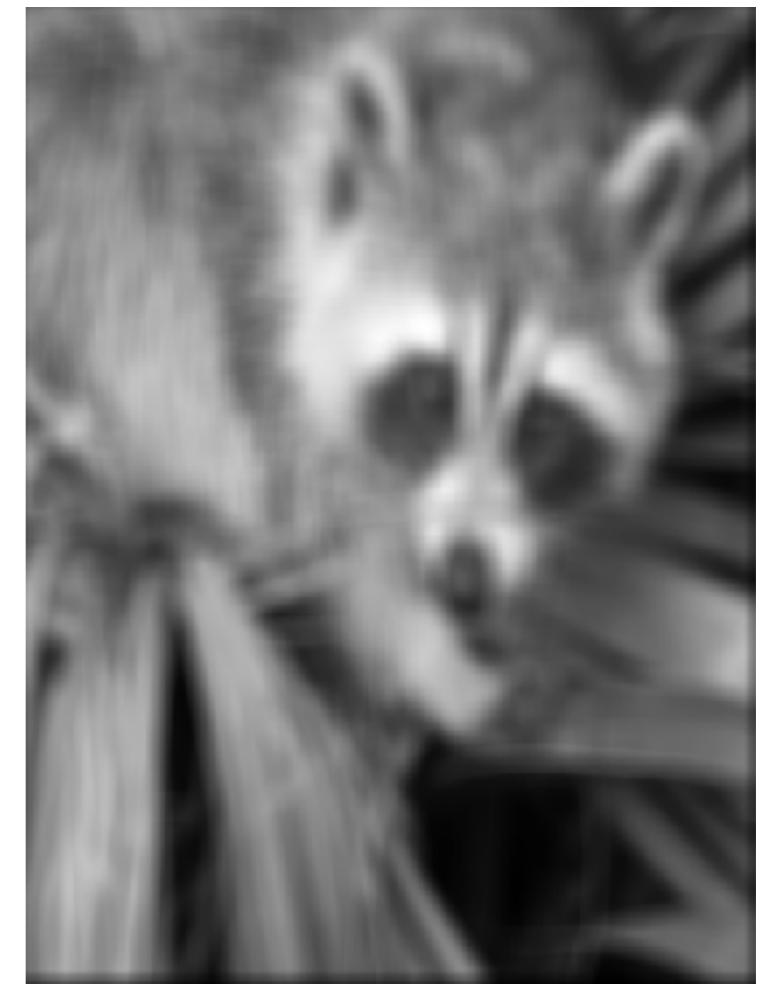
2D Convolution

- ❖ Low-pass Filter



$$\text{Input} * \begin{array}{|c|c|c|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ \hline \end{array} = \text{Output}$$

The diagram illustrates the convolution process. On the left is the input image of a raccoon. In the center is a 7x7 filter kernel filled with the value 1. An asterisk (*) indicates the convolution operation, followed by an equals sign (=). Below the filter is the label "Filter / kernel". To the right is the resulting output image, which is a blurred, low-pass filtered version of the original raccoon face.

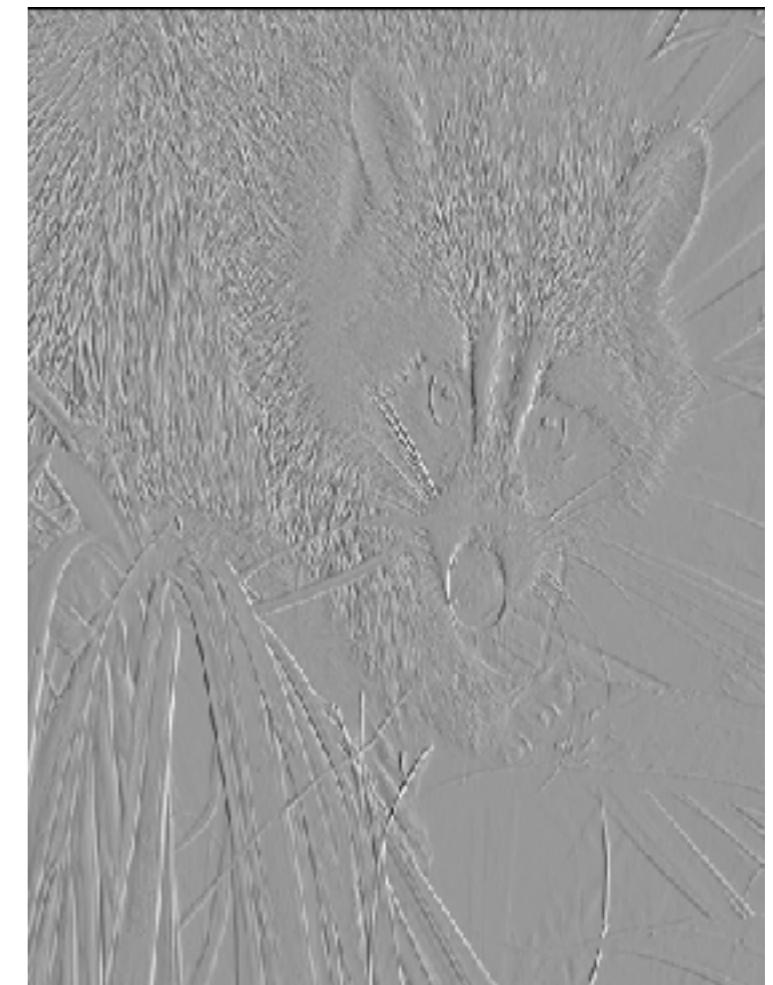


2D Convolution

- ❖ “Gradient image”



$$\ast \quad \begin{array}{|c|c|} \hline -1 & 1 \\ \hline -1 & 1 \\ \hline \end{array} = \quad \text{Filter/kernel}$$

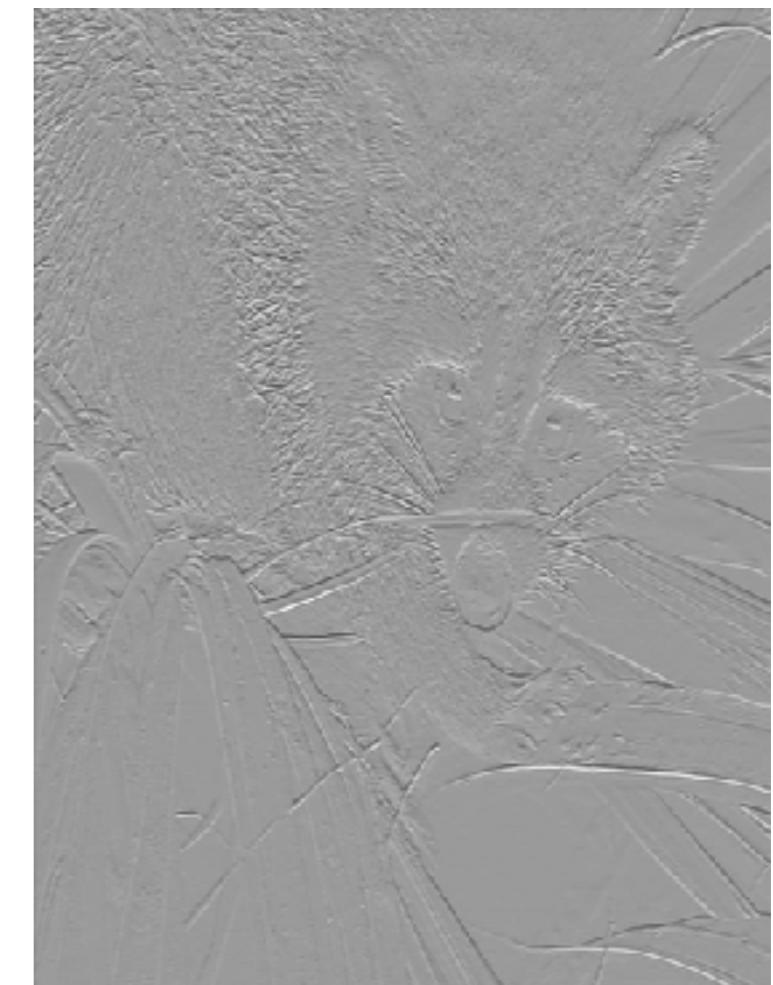


2D Convolution

- ❖ “Gradient image”

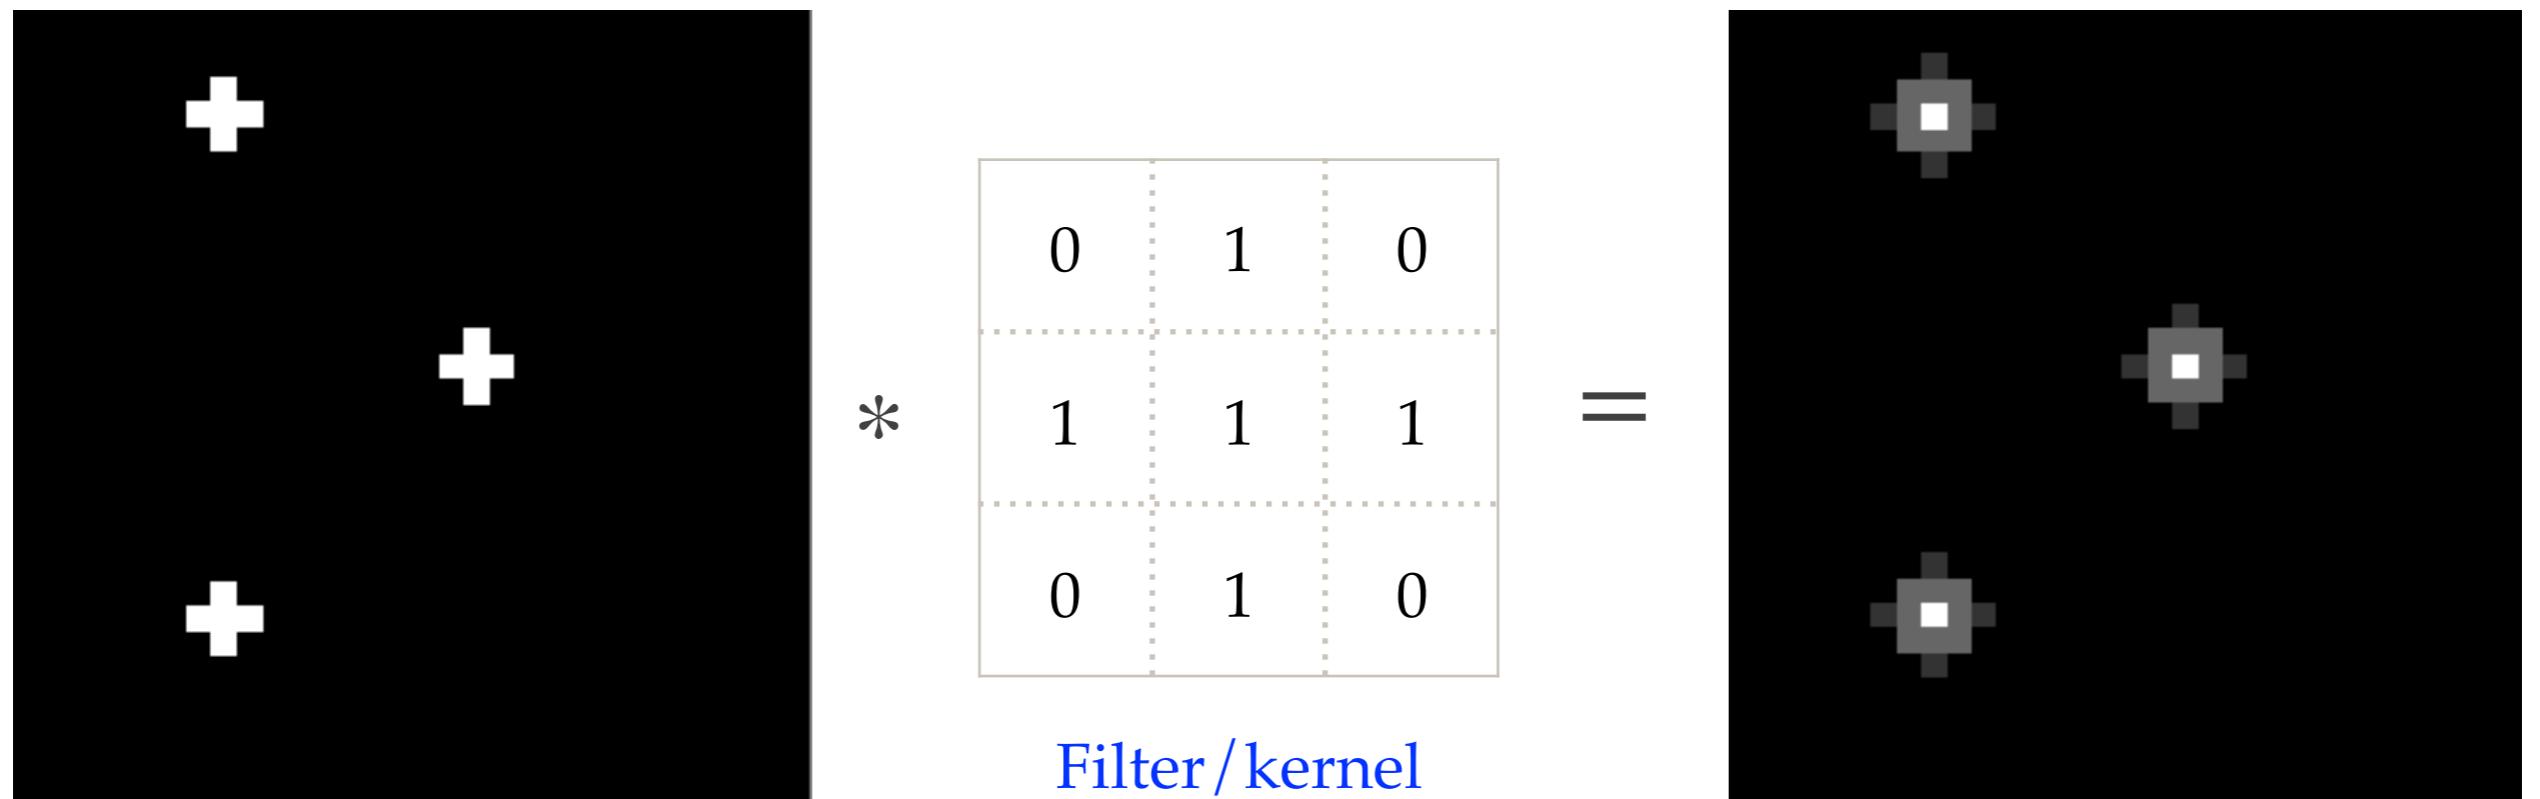


$$\begin{matrix} * & \begin{array}{|c|c|} \hline -1 & -1 \\ \hline 1 & 1 \\ \hline \end{array} & = \\ & \text{Filter / kernel} & \end{matrix}$$



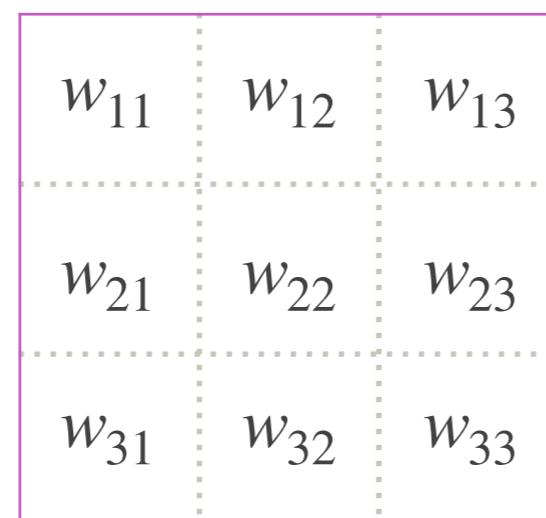
2D Convolution

- ❖ “Plus-sign Detector”



2D Convolution

- ❖ Different filters/kernels process images differently
 - ❖ Low-pass filter (blurring/denoising)
 - ❖ High-pass filter (sharpening)
 - ❖ Shape detector
 - ❖ Refer to [Image Kernels](#) for a good demo
- ❖ Instead of defining the filter ourselves, we let neural network come up with good ones for us

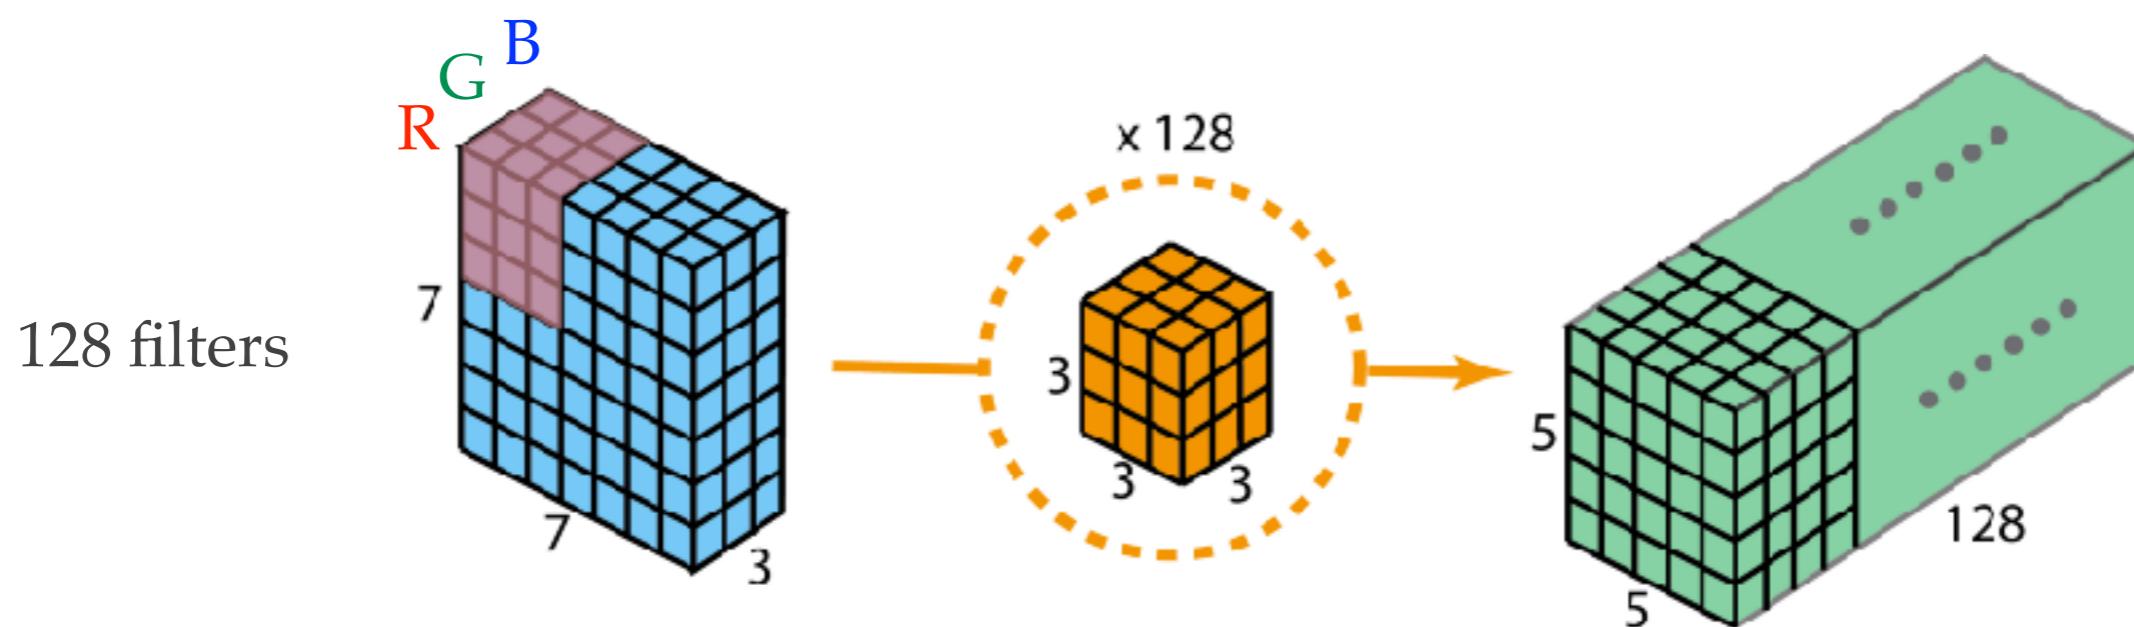
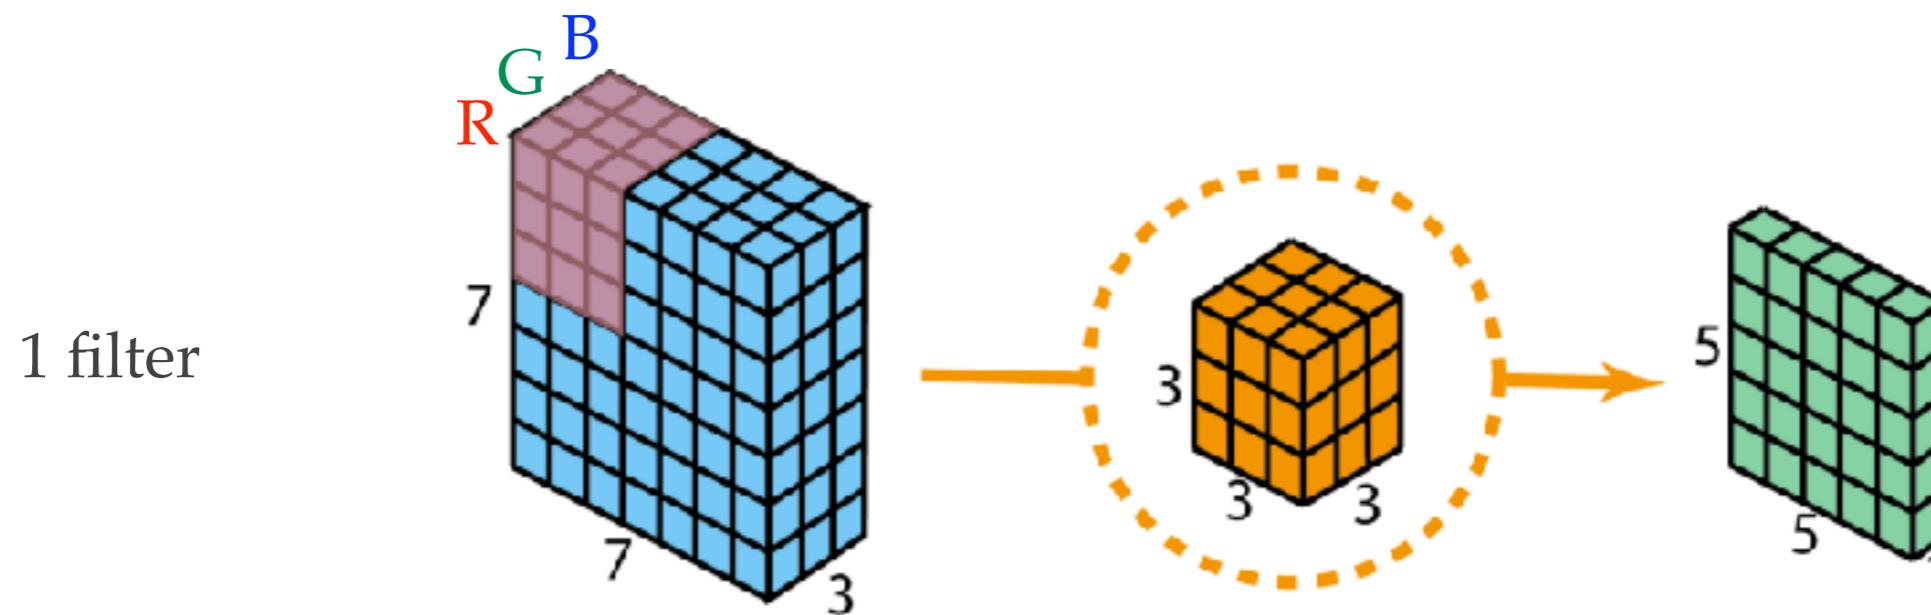


Filter/kernel
defines the relationship
between neighboring pixels

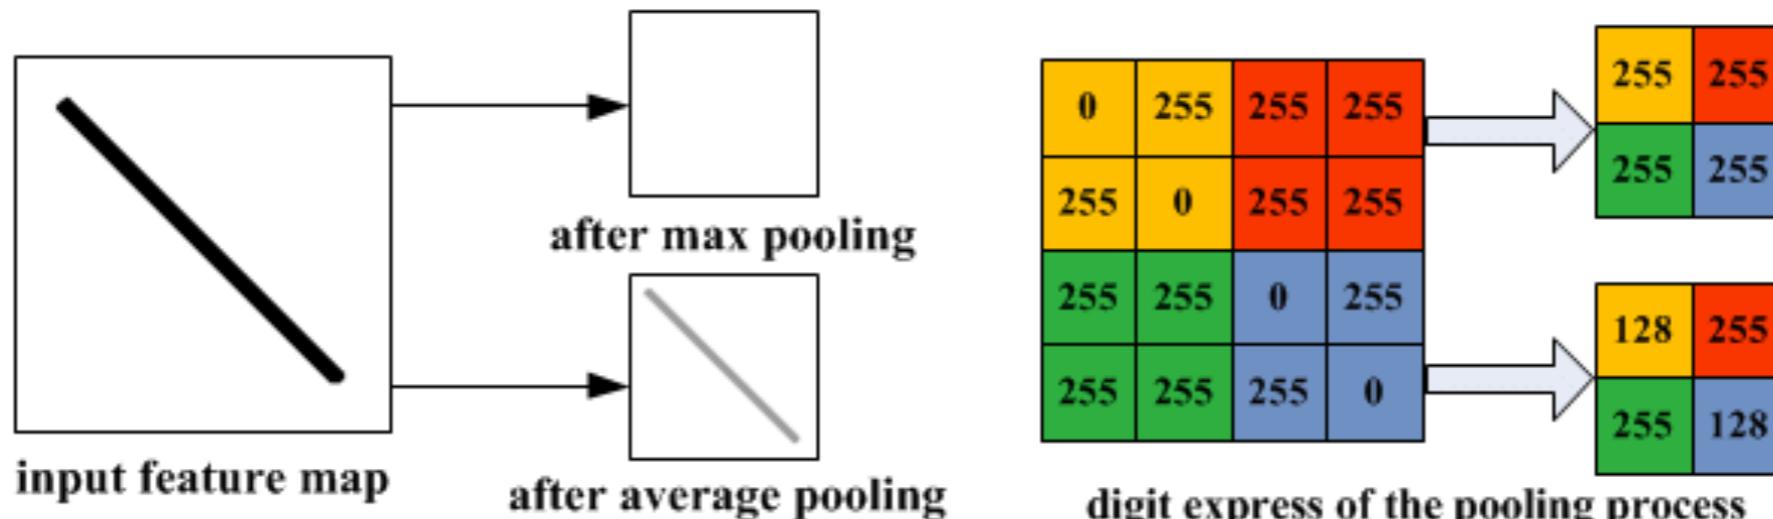
```
model.add(Conv2D(32, (3, 3), padding="same", activation="relu"))  
# of filters      filter size
```

2D Convolution

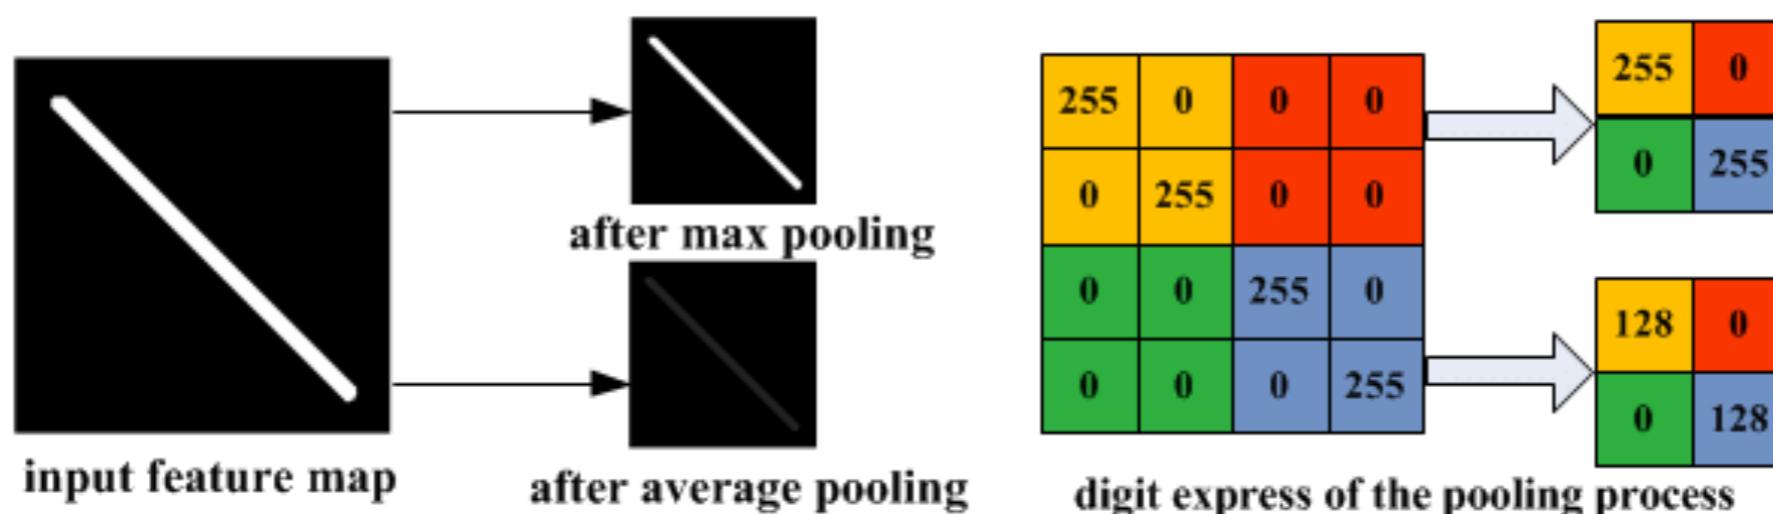
- ❖ 2D convolution with more than one channels



Pooling Layers



(a) Illustration of max pooling drawback

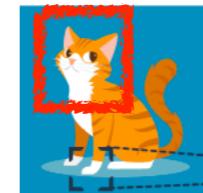
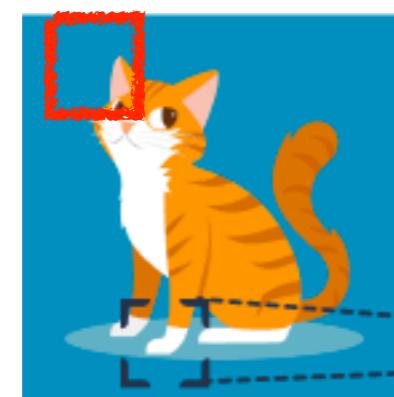
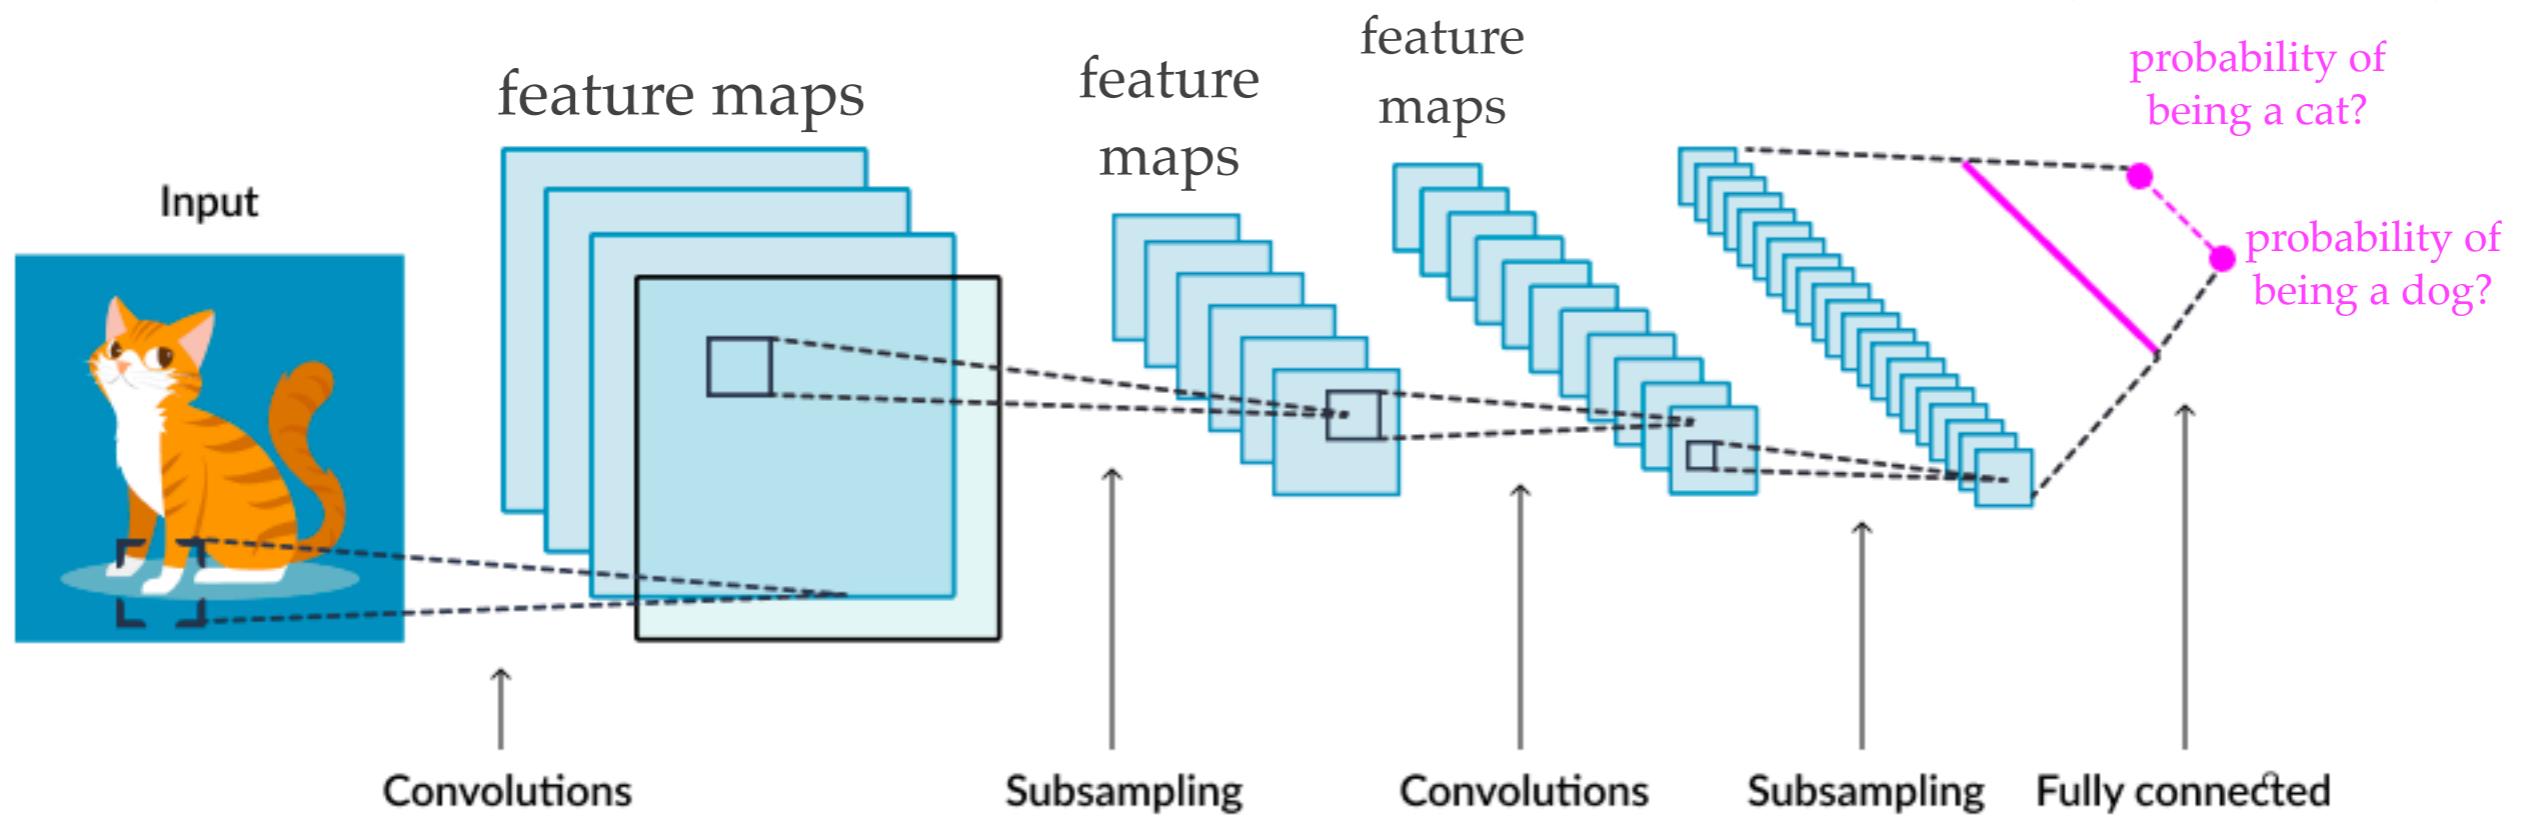


(b) Illustration of average pooling drawback

```
model.add(layers.MaxPooling2D((2,2)))
```

keywords: feature maps, activation maps, receptive field, backbone

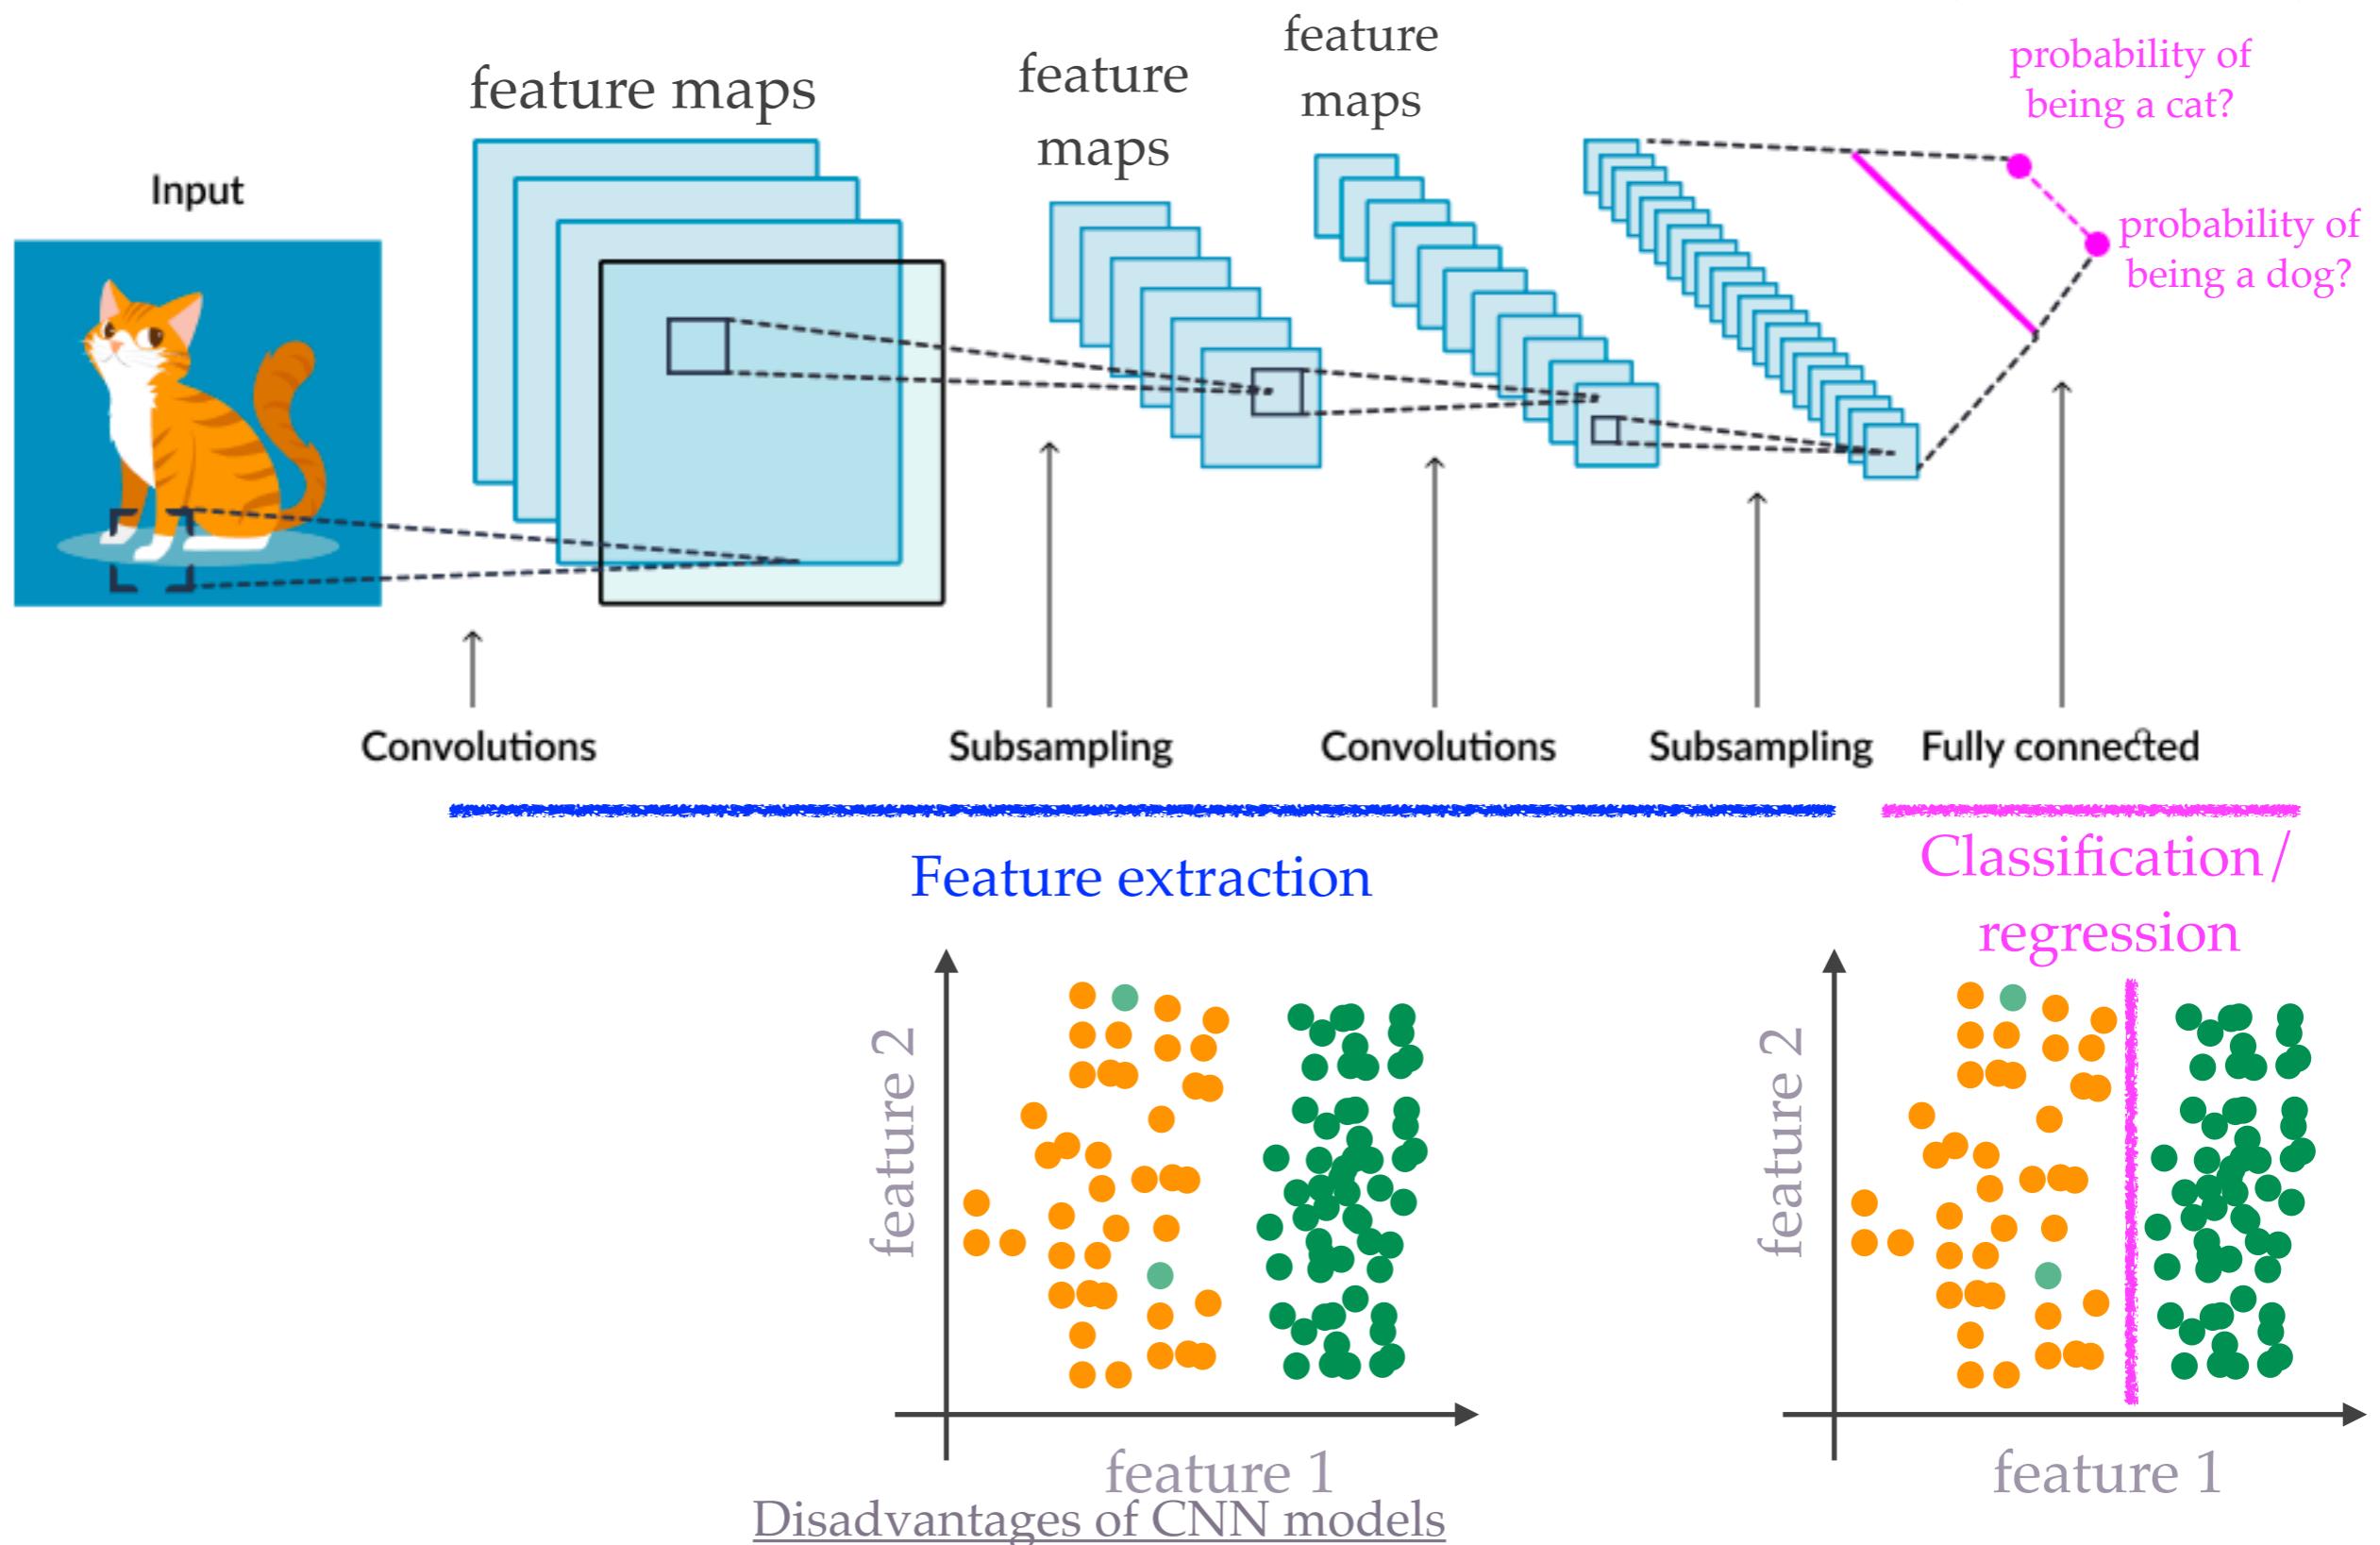
Convolutional Neural Network (CNN)



Same filter size (e.g., 3×3 filter), but different coverages for different “resolutions”

keywords: feature maps, activation maps, receptive field, backbone

Convolutional Neural Network (CNN)





Fully connected layer (Dense)

Optimizer
SGD
Adam
RMSprop

Evaluation metric
accuracy
F1-score
AUC
confusion matrix

Loss function
categorical crossentropy
binary crossentropy
mean squared error
mean absolute error

Regularization
Dropout
Data augmentation
 l_1, l_2 regularizations

Convolutional layer
Conv1D, 2D, 3D, ...
separable Conv

Pooling layer
max-pooling
average-pooling

Activation function
sigmoid
softmax
ESP (swish)
ReLU

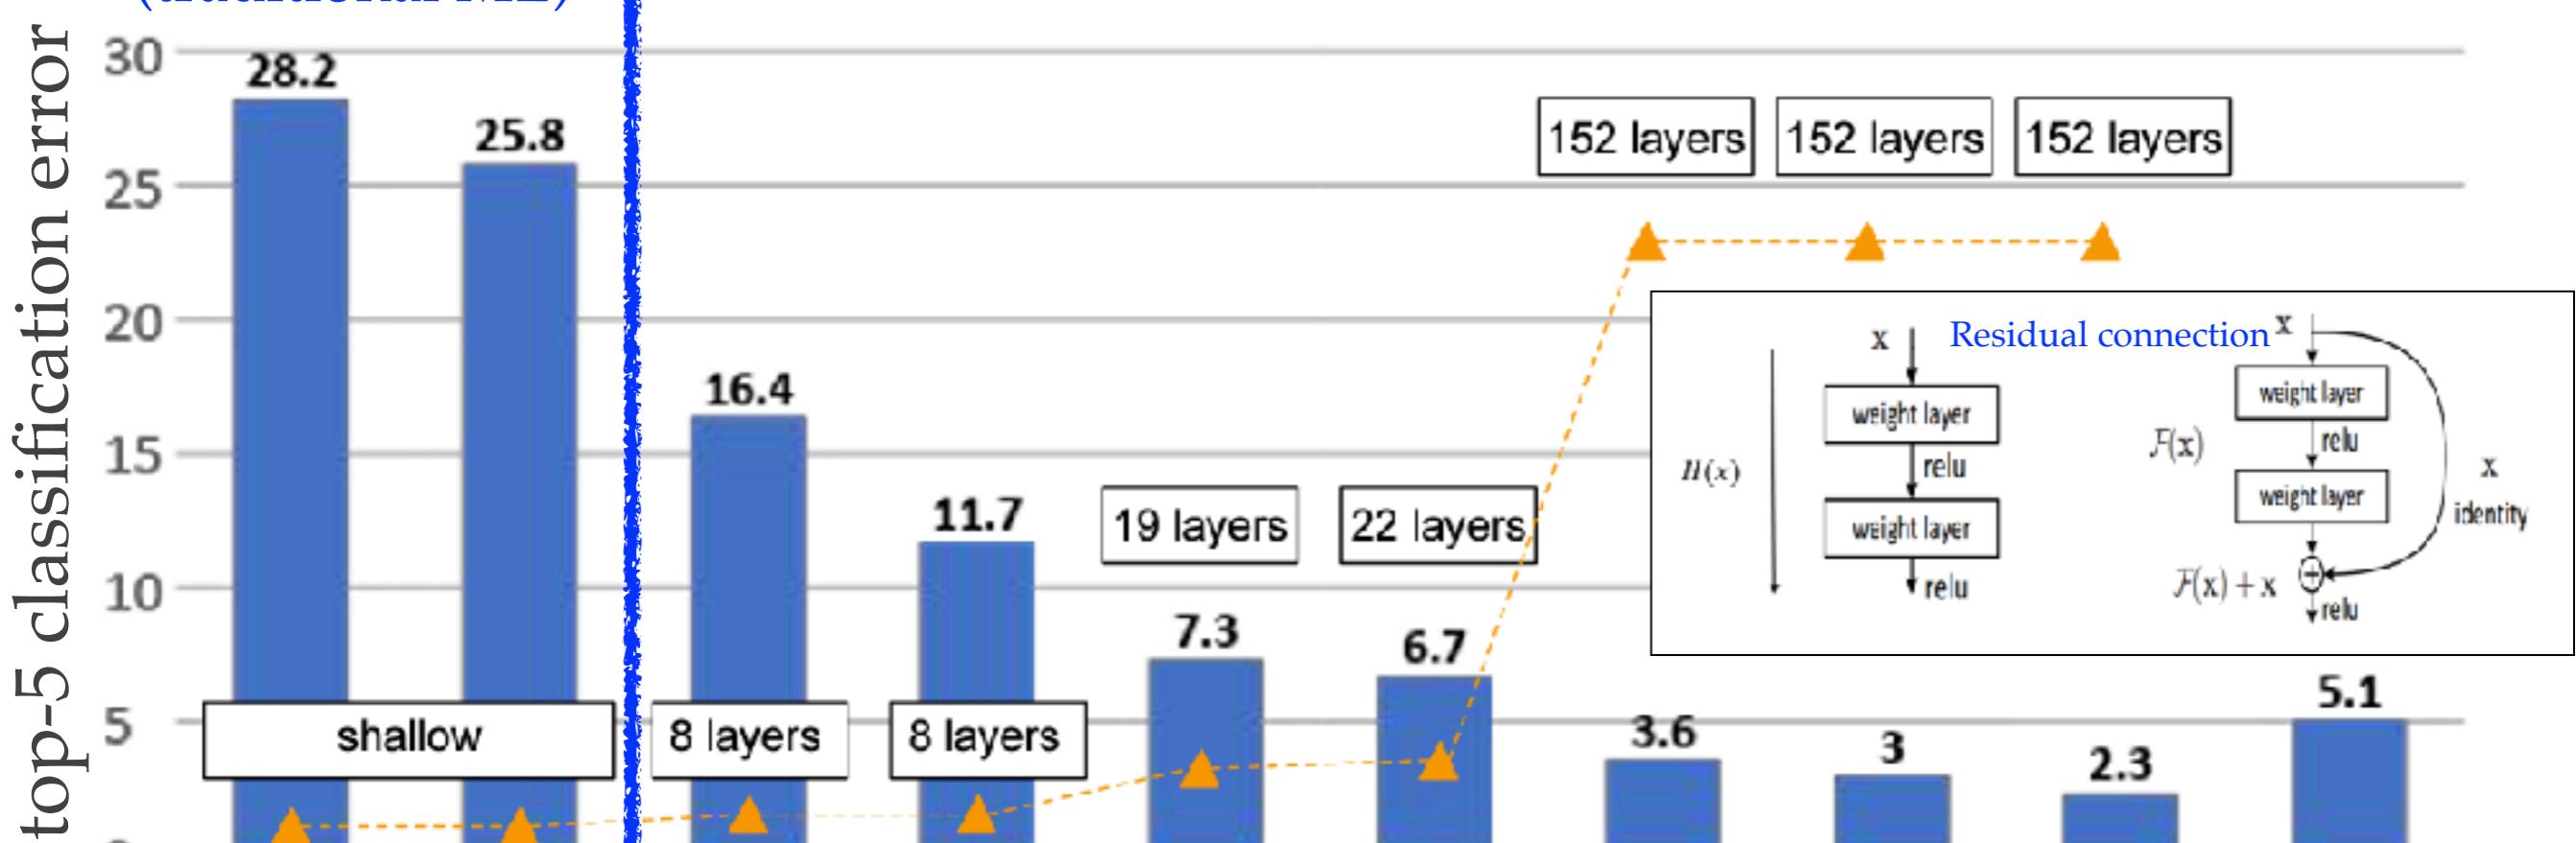
- ❖ **Combine basic components to build a neural network**
 - More components → “More” representative power

More networks: Xception, ResNeXt, DenseNet, MobileNet, NASNet, EfficientNet, SqueezeNet, Feature Pyramid Network

ImageNet

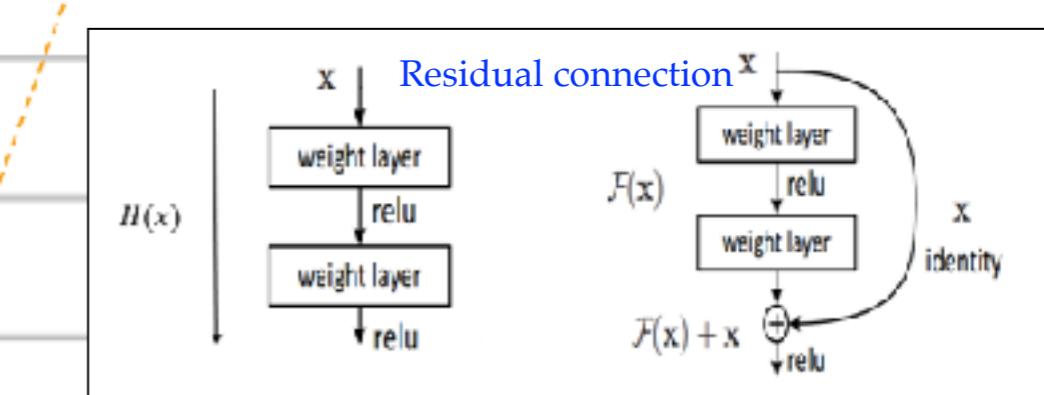
ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

The SVM era
(traditional ML)



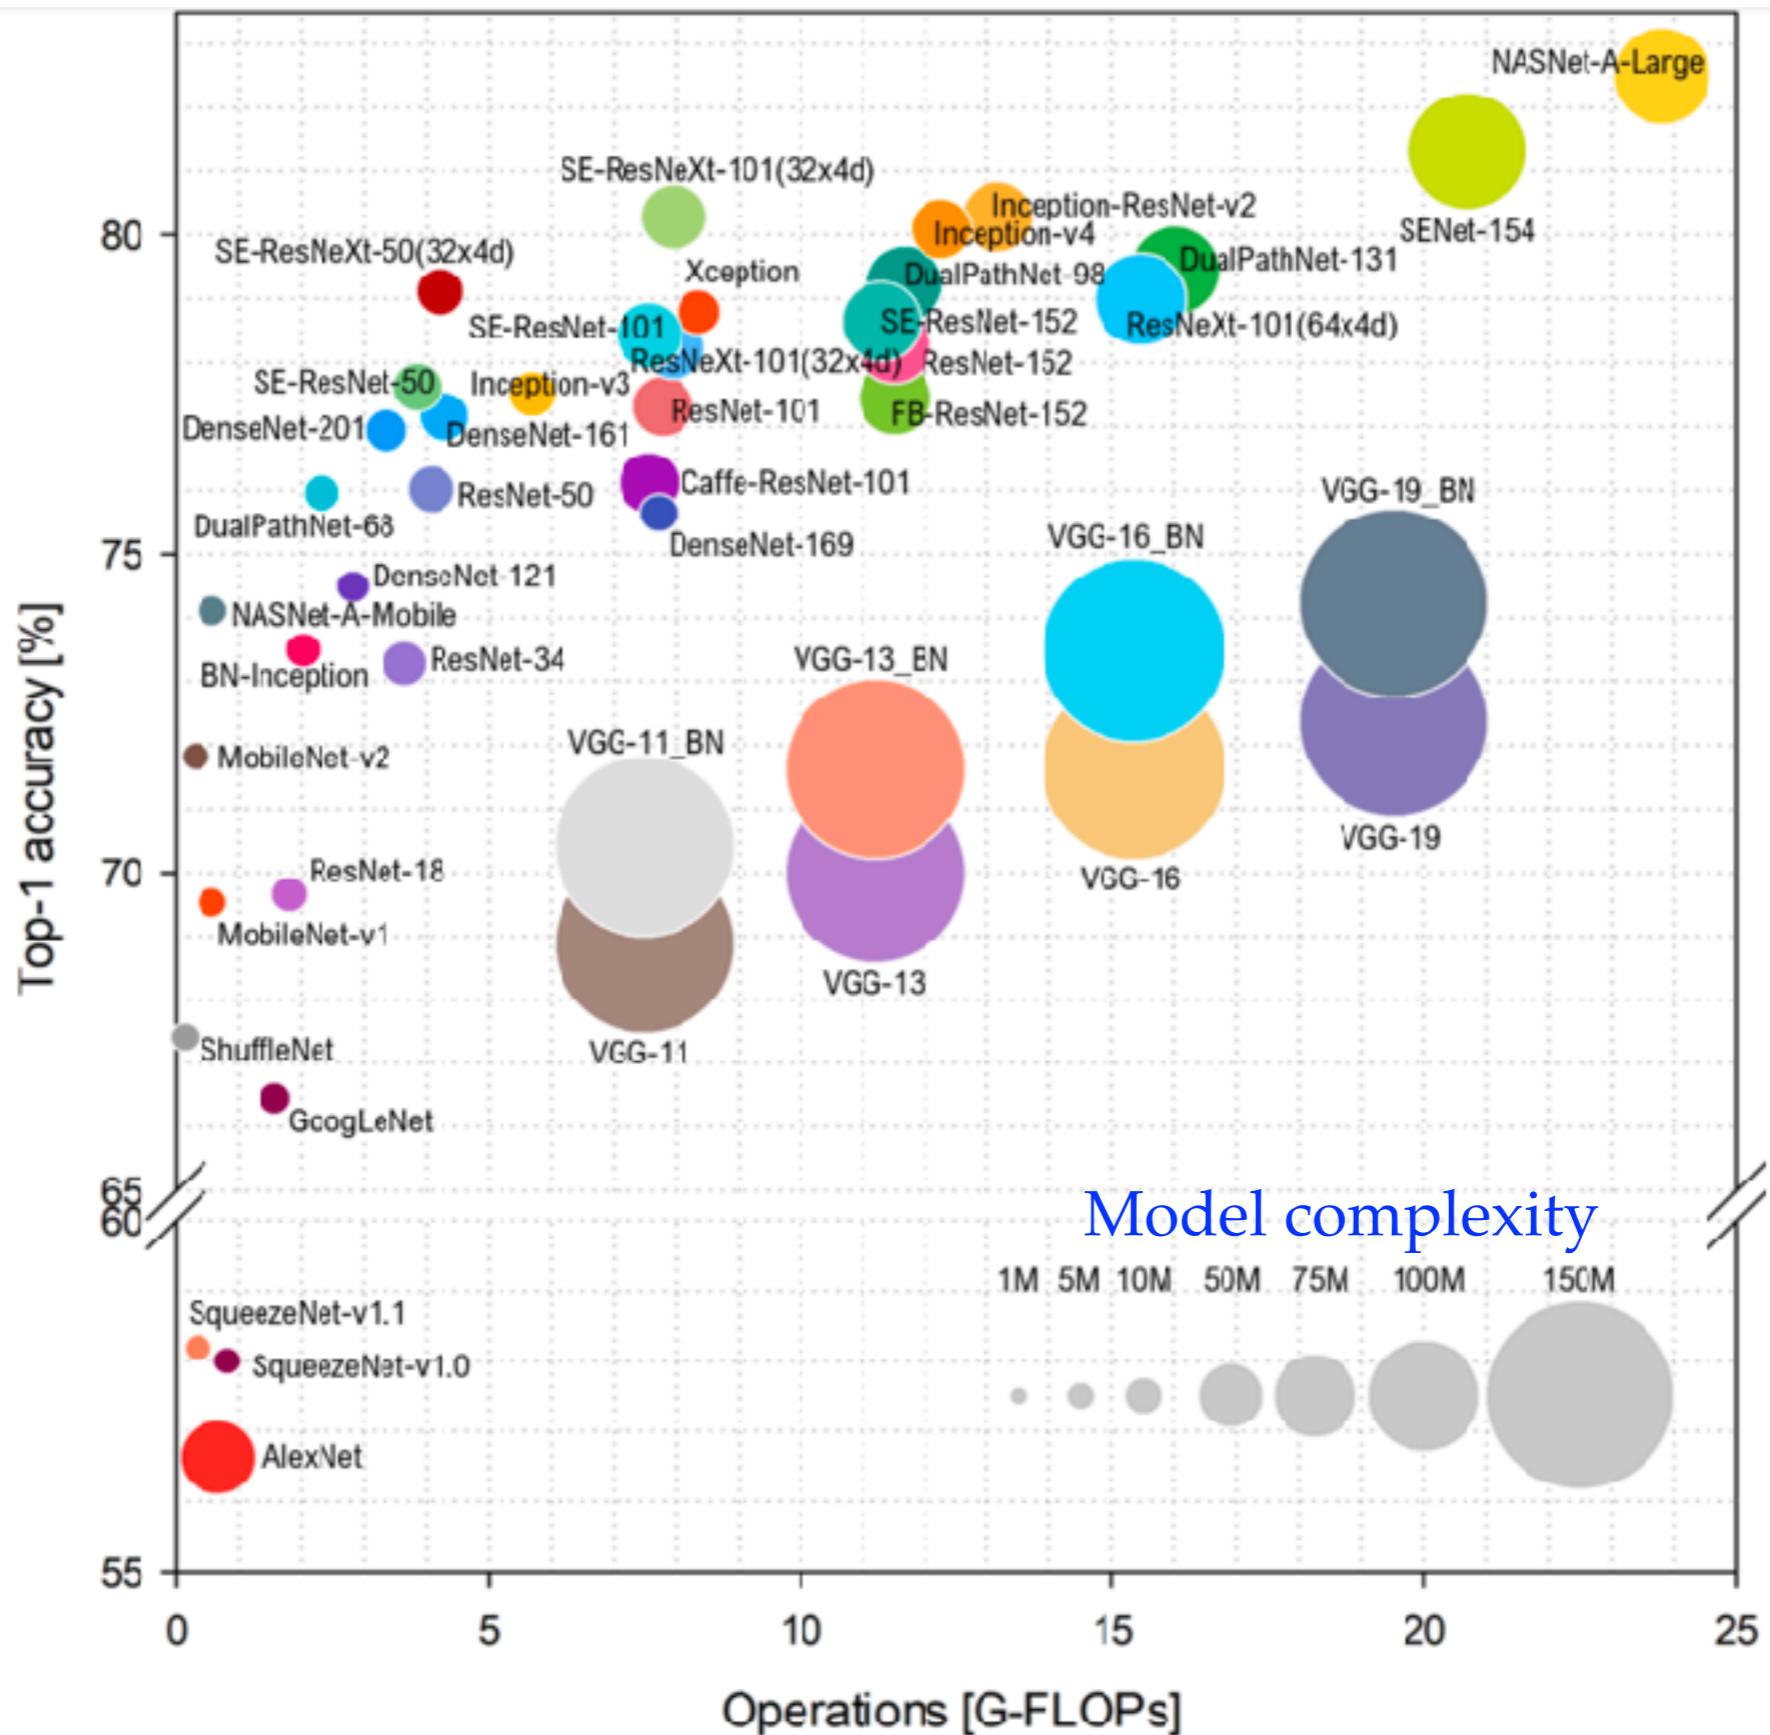
The deep learning era

152 layers 152 layers 152 layers



- Conv
- Pooling
- Dense
- Dropout
- ReLU
- SGD with momentum
- Optimized "AlexNet" in terms of # filter, kernel sizes etc.
- 19 layers with filter size = 3x3
- Introduce 1x1 Conv to reduce the dimensions
- Multiple losses to help with gradients
- Inception module
- Skip connection to help with the vanishing gradient problem
- Network ensemble
- Batch normalization
- Squeeze-and-Excitation (SE) block
- Adaptive feature map reweighting

The last annual ImageNet competition was held in 2017

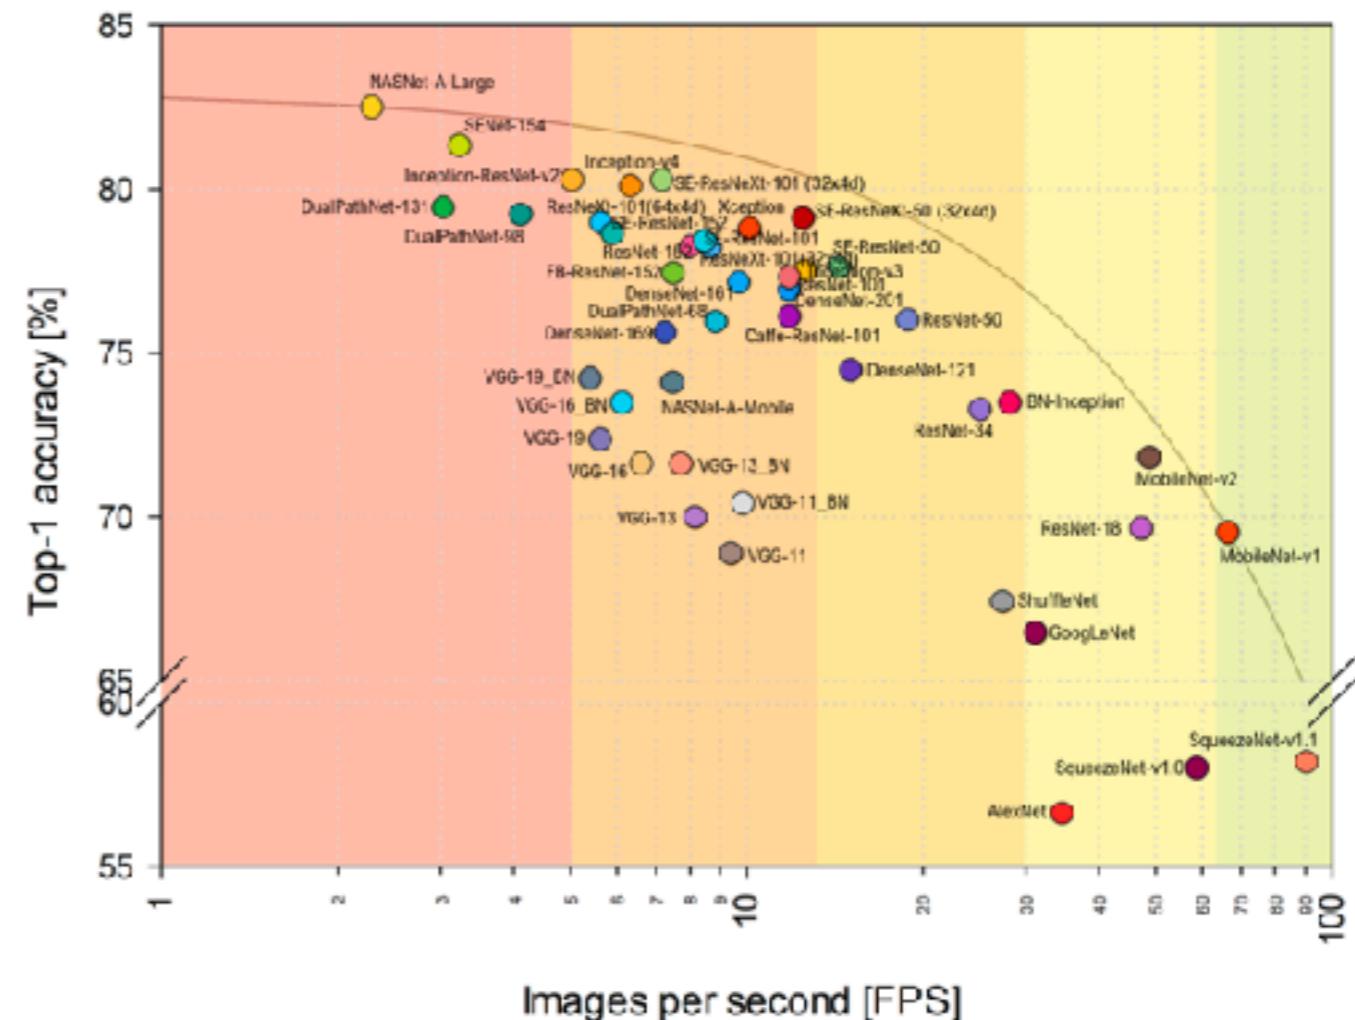
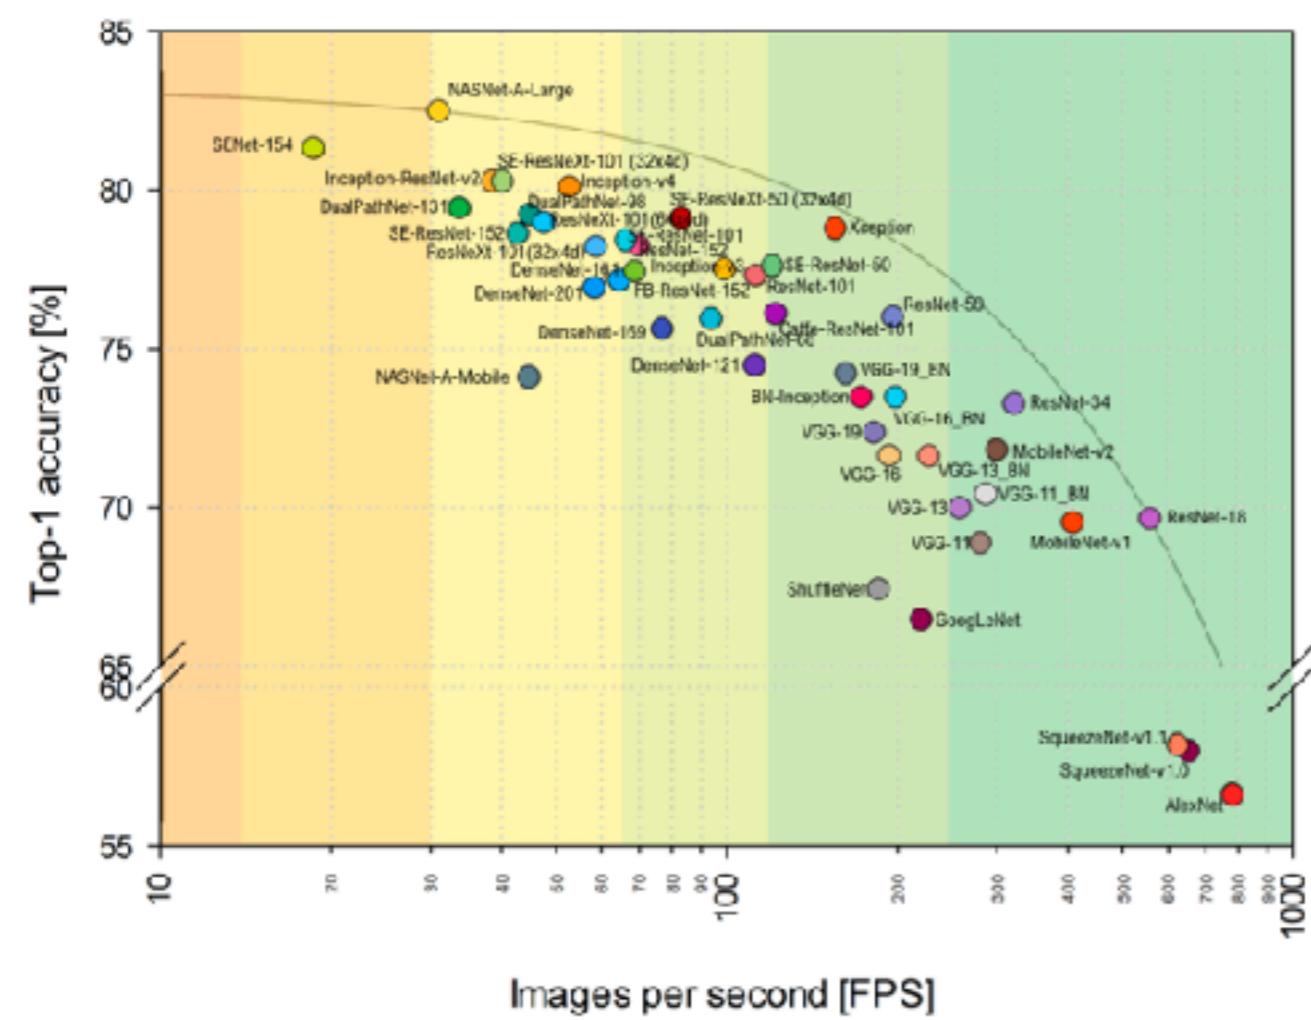
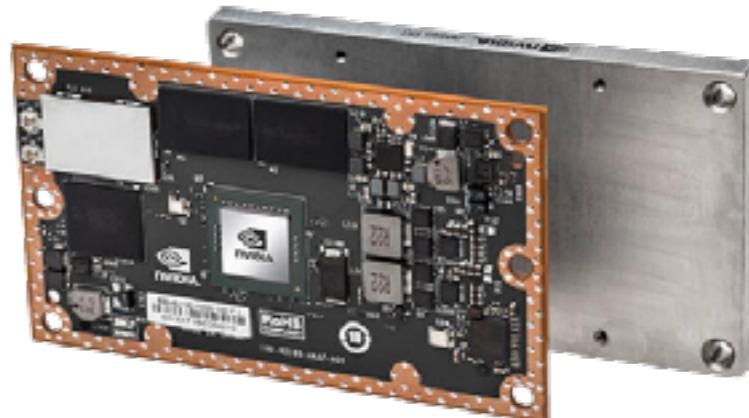


Floating-point operations (FLOPs) required for a single forward pass

Titan Xp

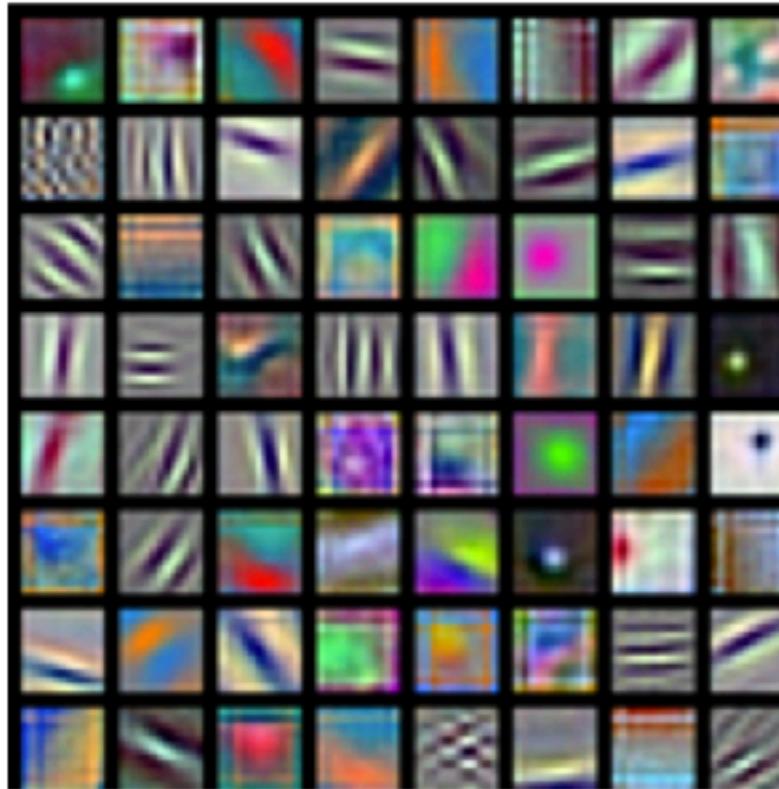


Jetson TX1



Understanding CNN - Filter Visualization

Learned weights of the filters at the first layer



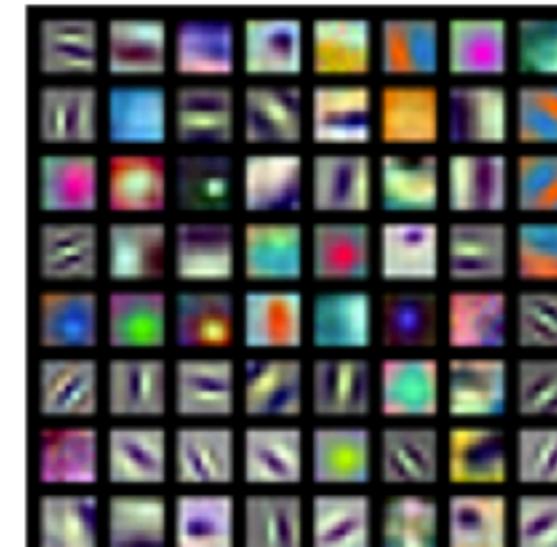
AlexNet:
64 x 3 x 11 x 11

64 filters filter size = 11 x 11
with 3 channels

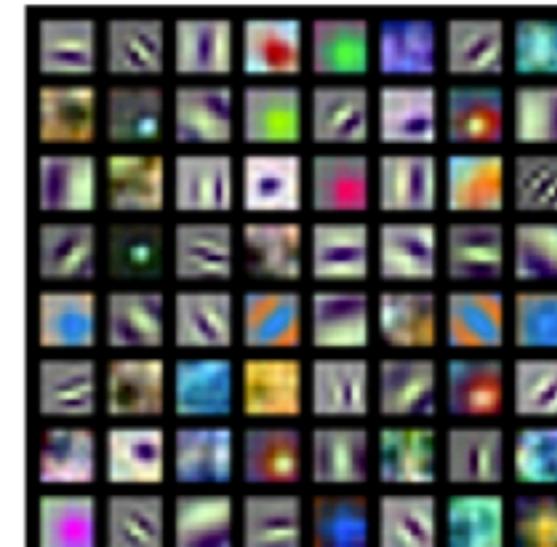
What these filters look for (through template matching/inner product)

- oriented edges
- opposing colors

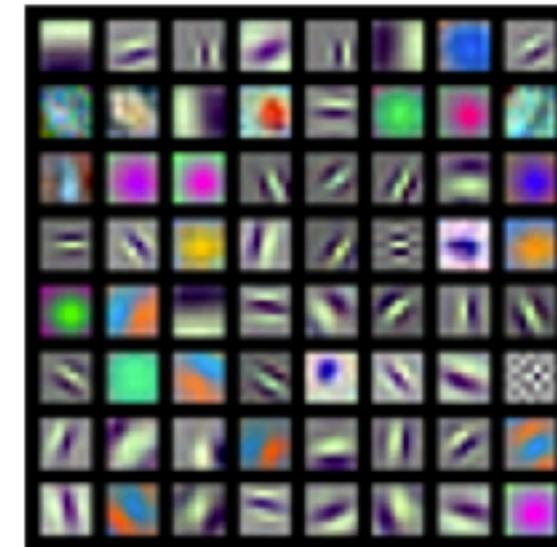
Similar to human visual system!



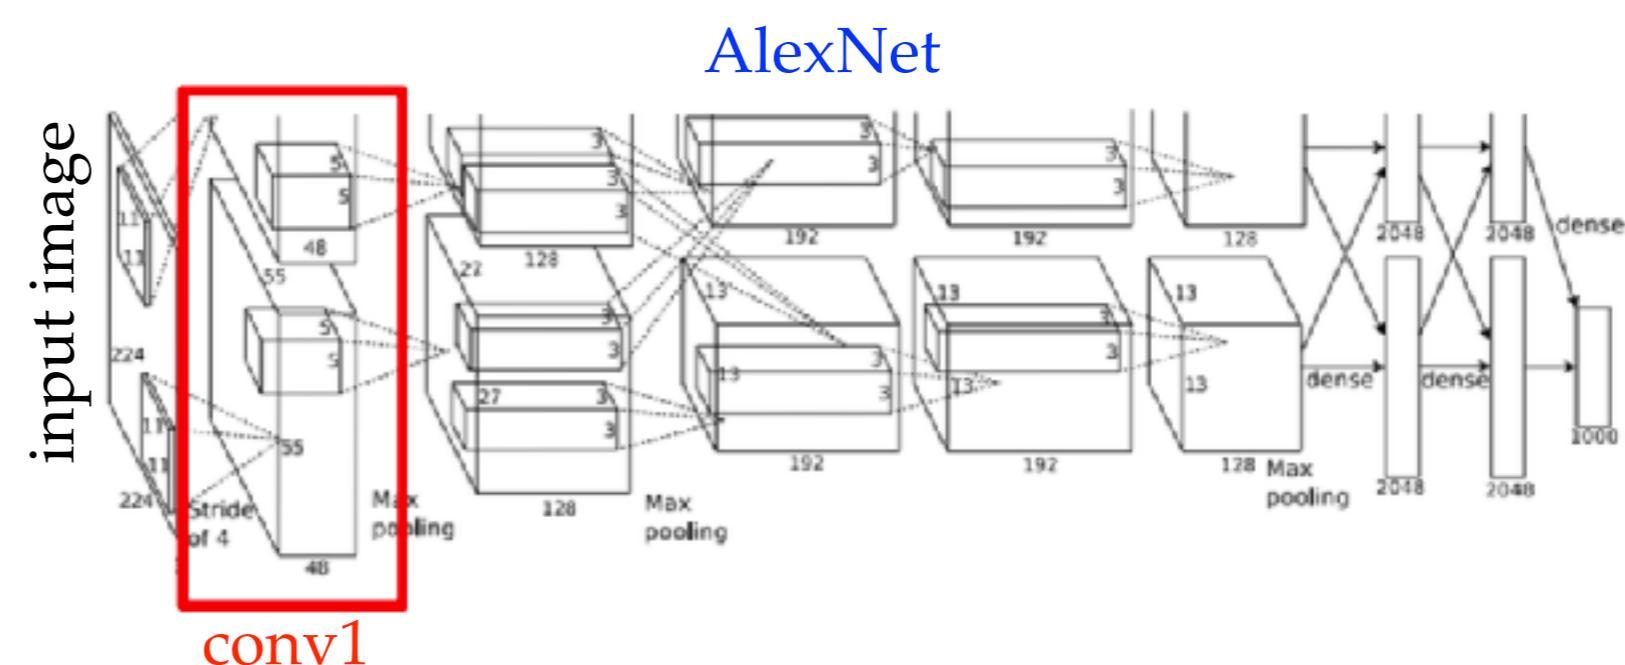
ResNet-18:
64 x 3 x 7 x 7



ResNet-101:
64 x 3 x 7 x 7



DenseNet-121:
64 x 3 x 7 x 7



Understanding CNN - Filter Visualization

Visualize the filters/kernels (raw weights)

We can visualize filters at higher layers, but not that interesting

(these are taken from ConvNetJS CIFAR-10 demo)



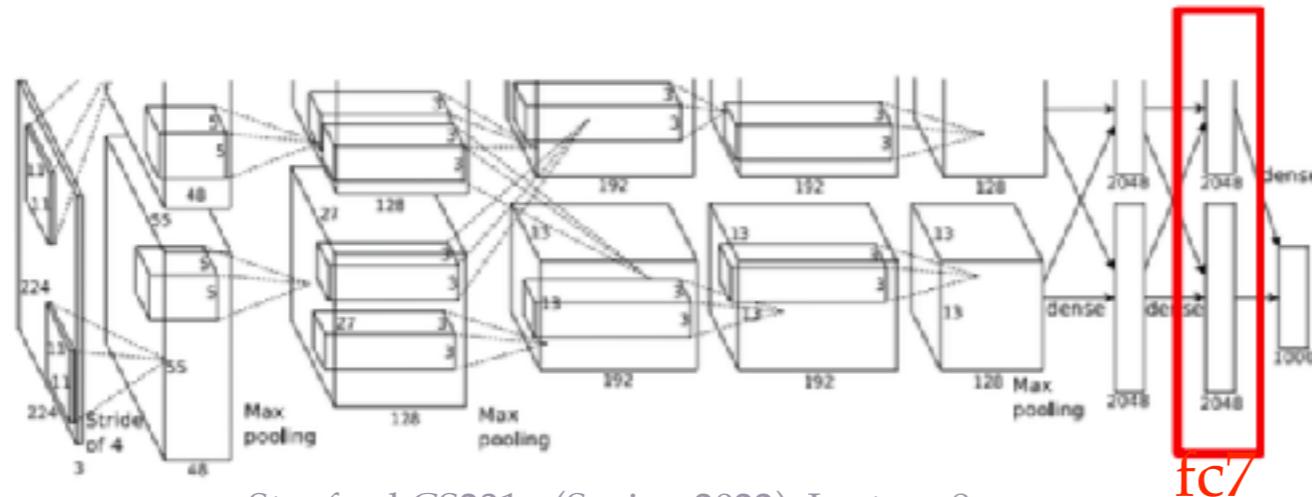
Stanford CS231n (Spring 2022): Lecture 8.

The weights of conv2 tell us what type of the activation patterns after conv1 that would cause conv2 to maximally activate. These filters are not connected directly to the input image, so it is not directly interpretable.

We need more complicated techniques to get a sense of what is going on.

Understanding CNN - Last Layer

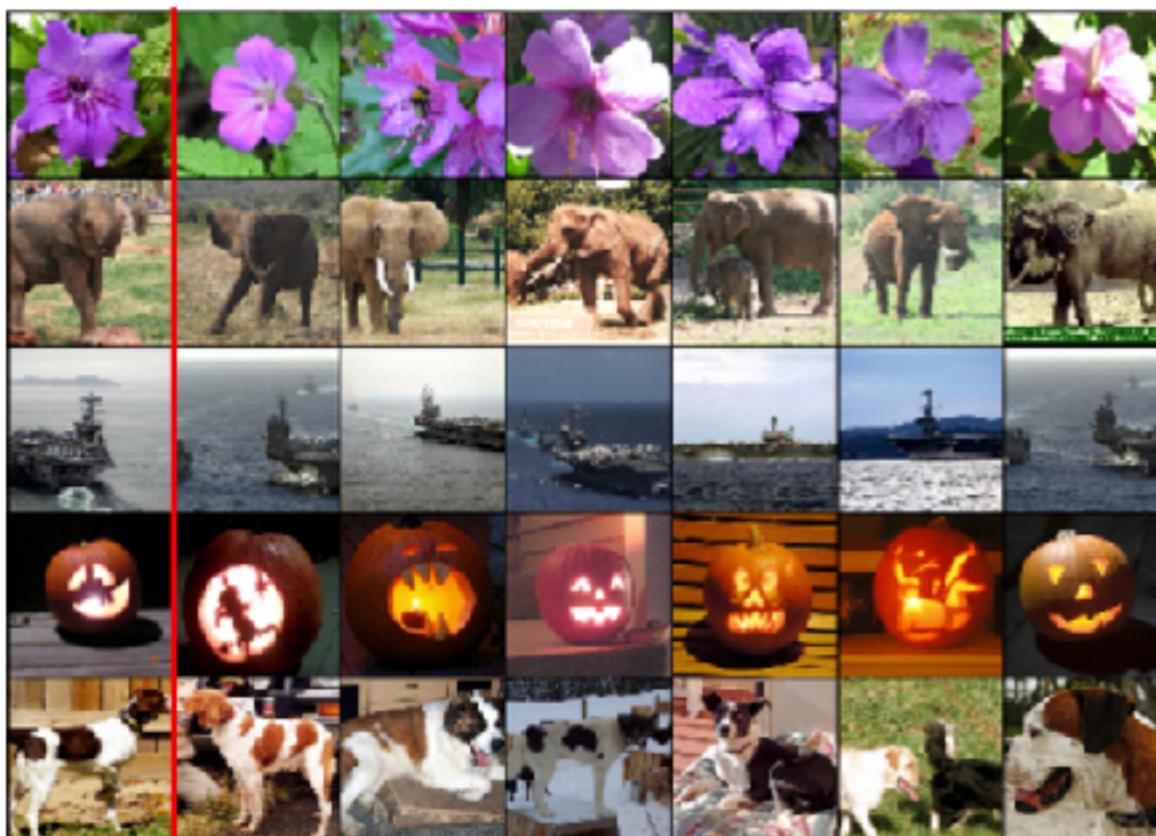
input image



Stanford CS231n (Spring 2022): Lecture 8.

output of fc 7 = feature vector (last layer)
of size 4096

Test image L2 Nearest neighbors in feature space



Test
image

L2 Nearest neighbors in pixel space

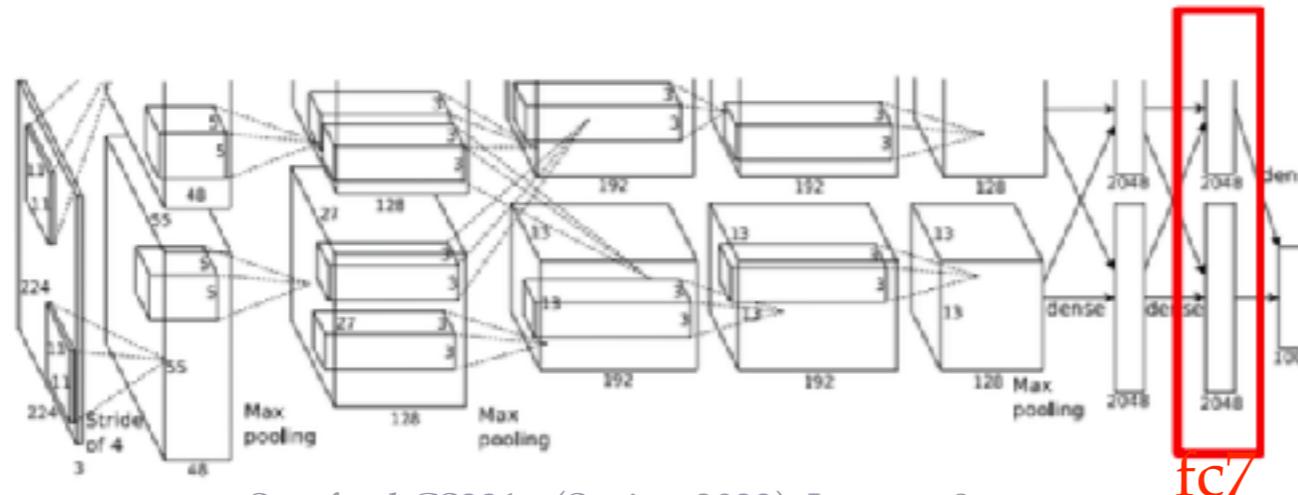


Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems 25 (2012).

Stanford CS231n (Spring 2022): Lecture 8.

Understanding CNN - Last Layer

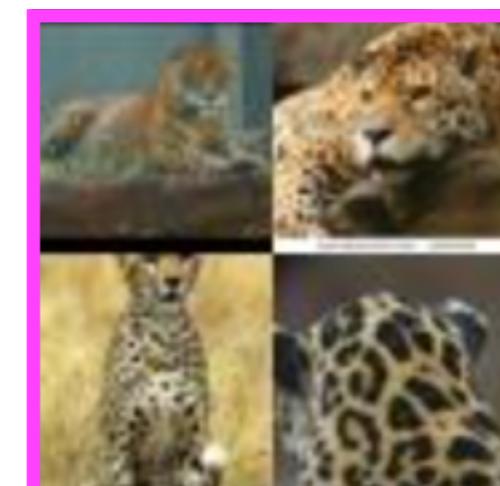
input image



Stanford CS231n (Spring 2022): Lecture 8.

output of fc 7 = feature vector (last layer)
of size 4096

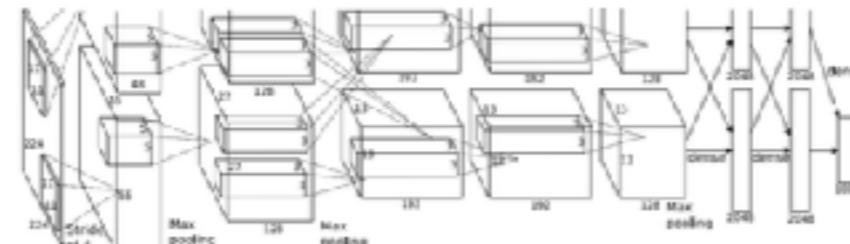
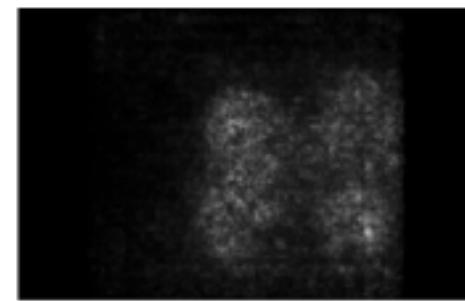
Perform a machine learning algorithm to the output of fc7
Barnes-Hut t-SNE



Understanding CNN - Other Techniques

Saliency Maps

Forward pass: Compute probabilities



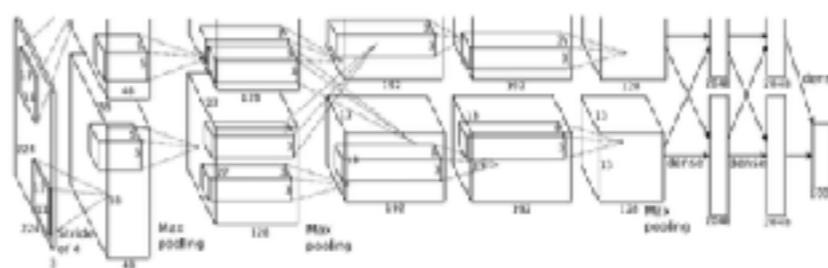
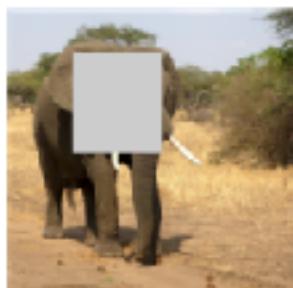
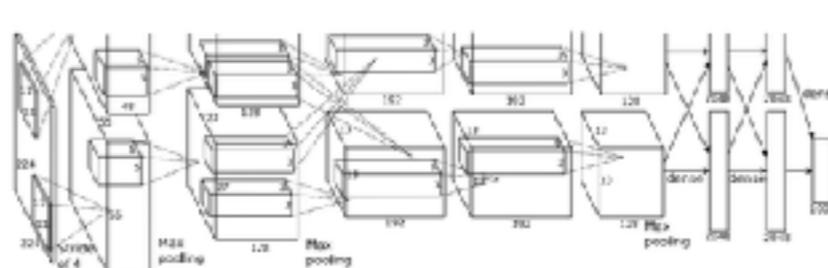
Dog

gradient of class score w.r.t. image pixels

Simonyan, Karen, Andrea Vedaldi, and Andrew Zisserman. "Deep inside convolutional networks: Visualising image classification models and saliency maps." arXiv preprint arXiv:1312.6034 (2013).

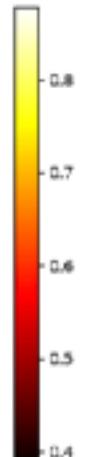
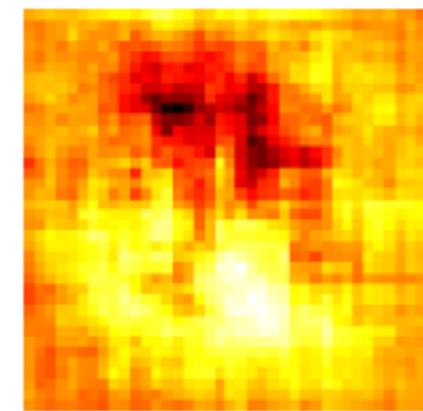
Occlusion Experiments

Mask part of the image before feeding to CNN,
check how much predicted probabilities change



Zeiler, Matthew D., and Rob Fergus. "Visualizing and understanding convolutional networks." European conference on computer vision. Springer, Cham, 2014.

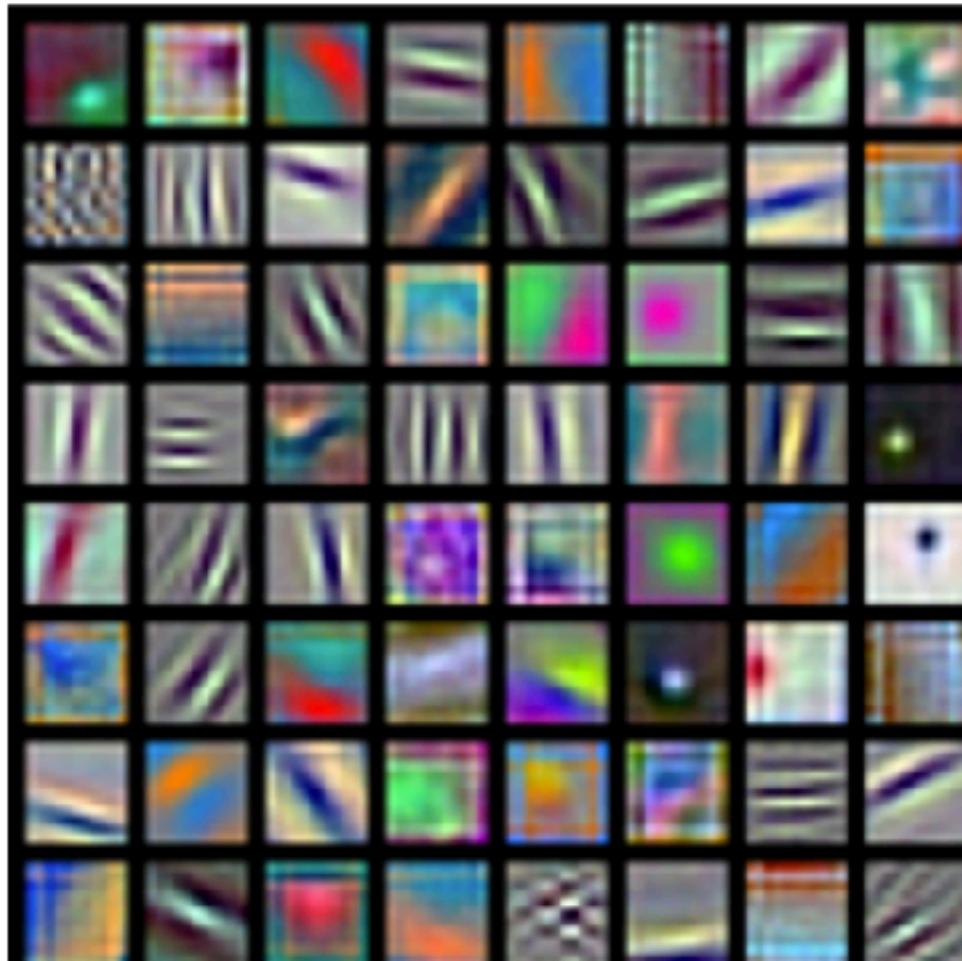
African elephant, Loxodonta africana



...and many more such as

- ❖ Guided Backpropagation
- ❖ Grad-CAM, Score-CAM, HiResCAM
- ❖ Attention map visualization

Trained CNN as a Feature Extractor



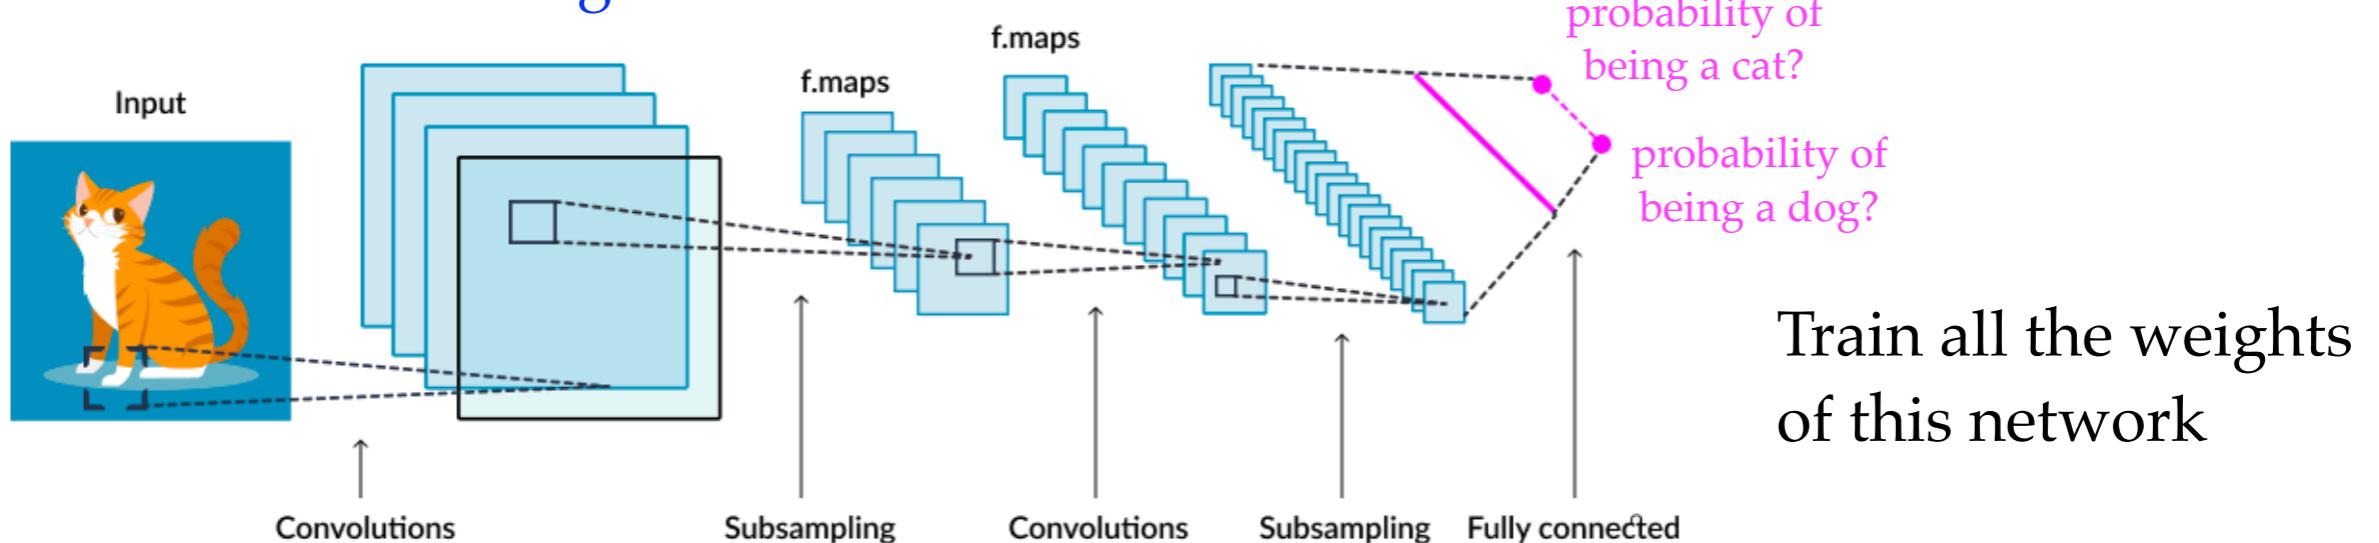
What these filters look for (through template matching/inner product)

- oriented edges
- opposing colors

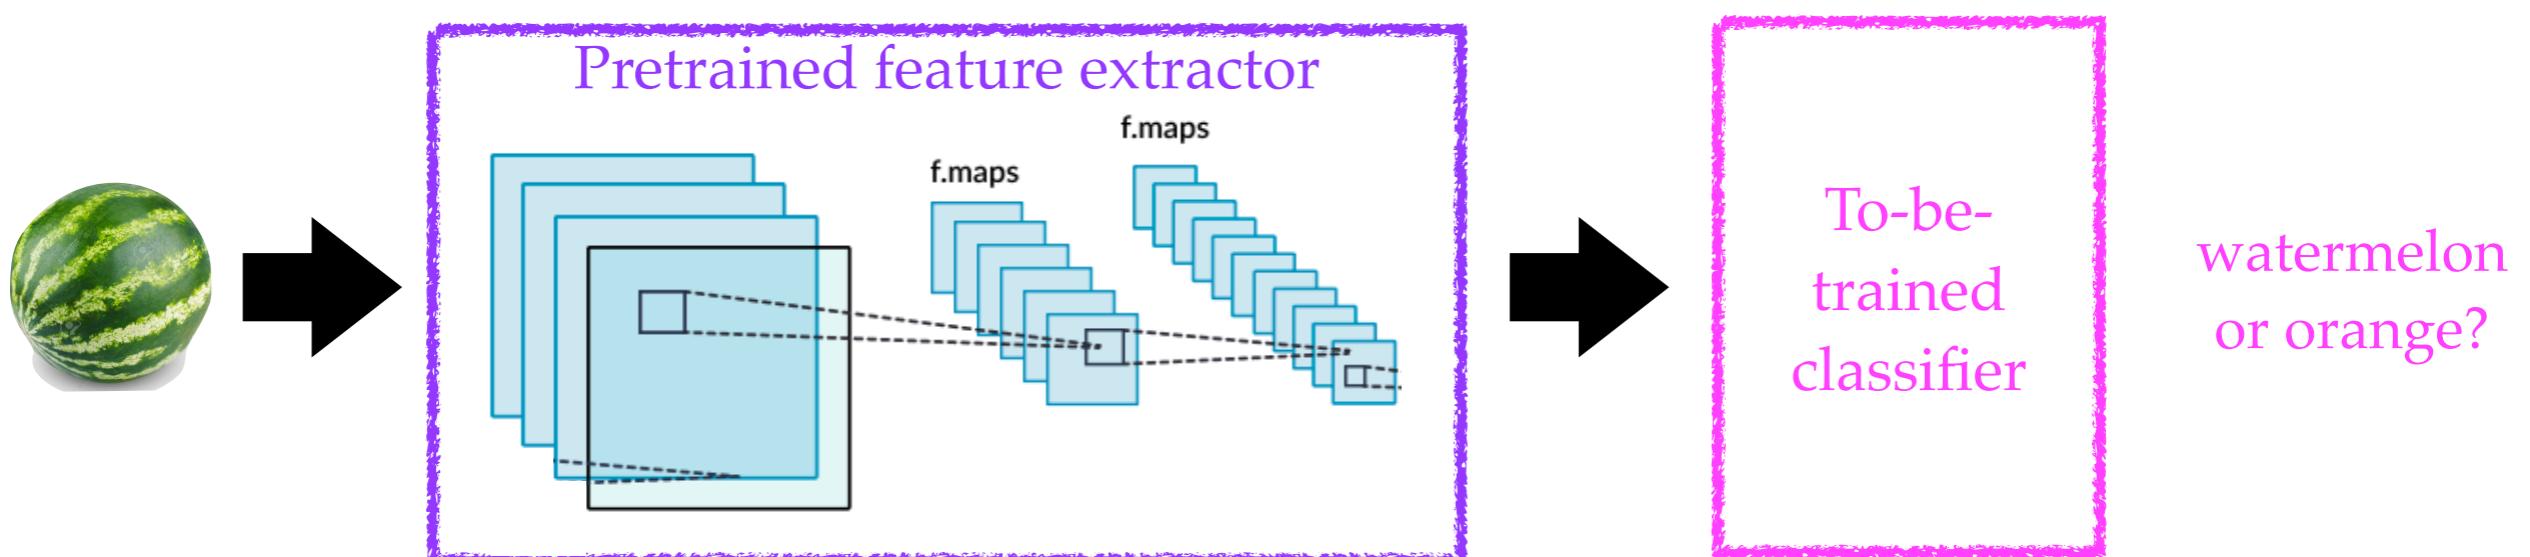
These patterns are important for processing other types of images as well, so we can reuse trained weights when we encounter a related task or the same task but with a different dataset

Trained CNN as a Feature Extractor

Previous Task: Cat-vs-dog classification



New Task: Watermelon-vs-orange classification



Take parts of the trained network
and use them as a feature extractor

Only train the
weights here

Transfer Learning and Fine-Tuning

ImageNet (1000 classes)

14M images



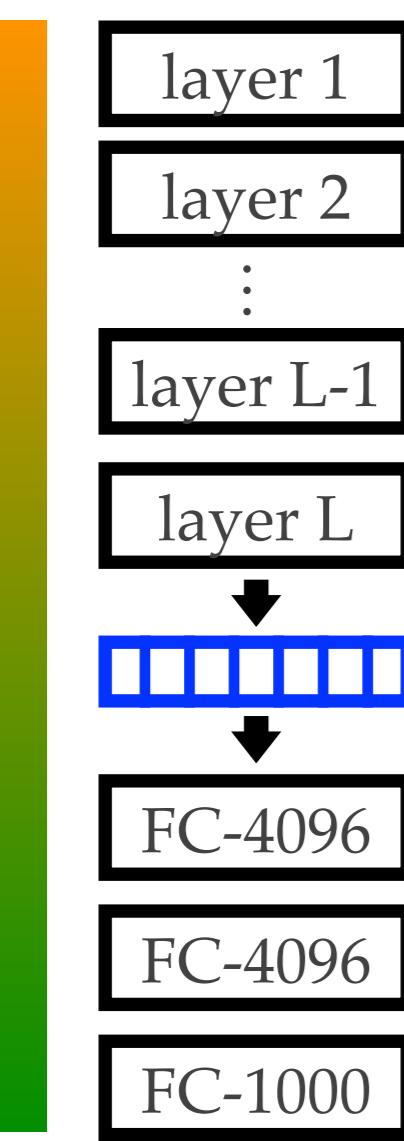
AI-vs-Artists (2 classes)

Assume small dataset



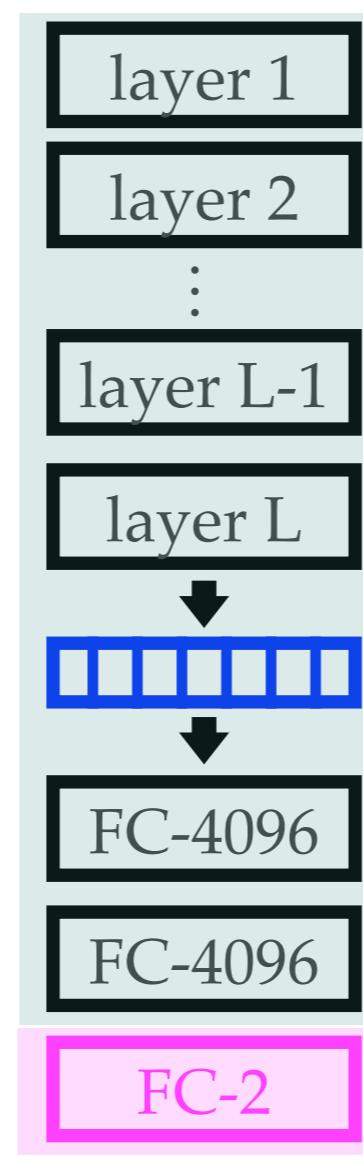
More
“reusable”

Less
“reusable”



Clone these layers
along with their
weights and do not
modify the weights

Replace the
classification layer
and train it



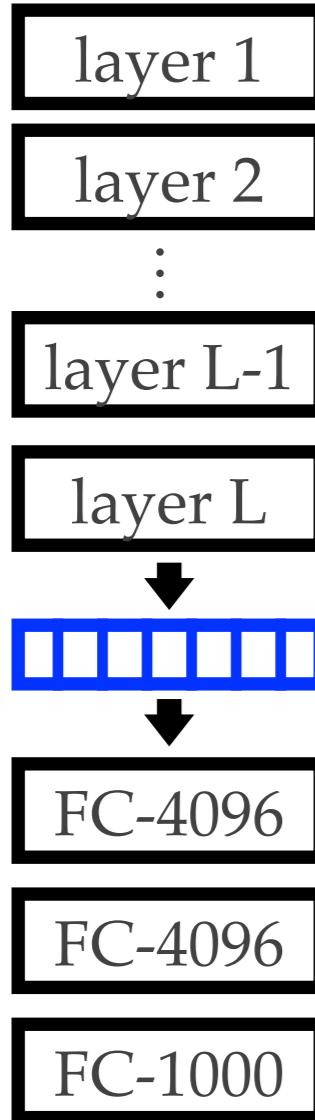
Transfer Learning and Fine-Tuning

ImageNet (1000 classes)

14M images



More
“reusable”



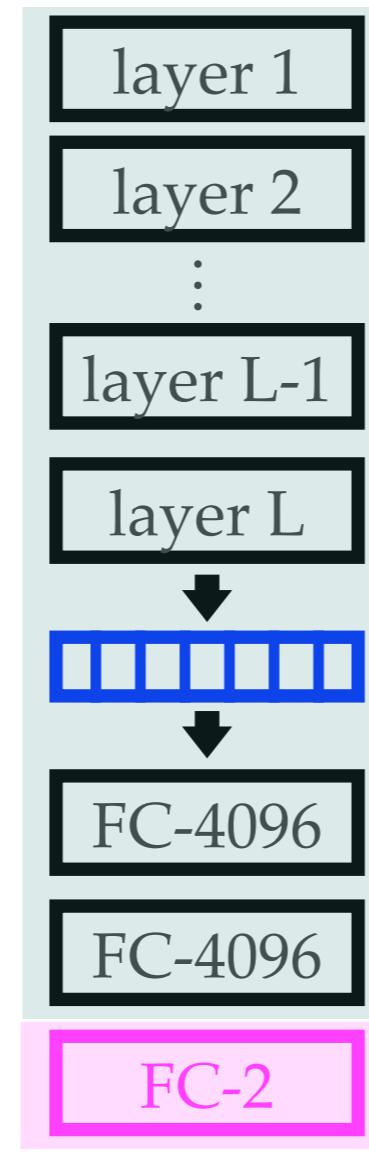
Clone these layers along with their weights and do not modify the weights
Replace the classification layer and train it

AI-vs-Artists (2 classes)

Assume small dataset



Less
“reusable”



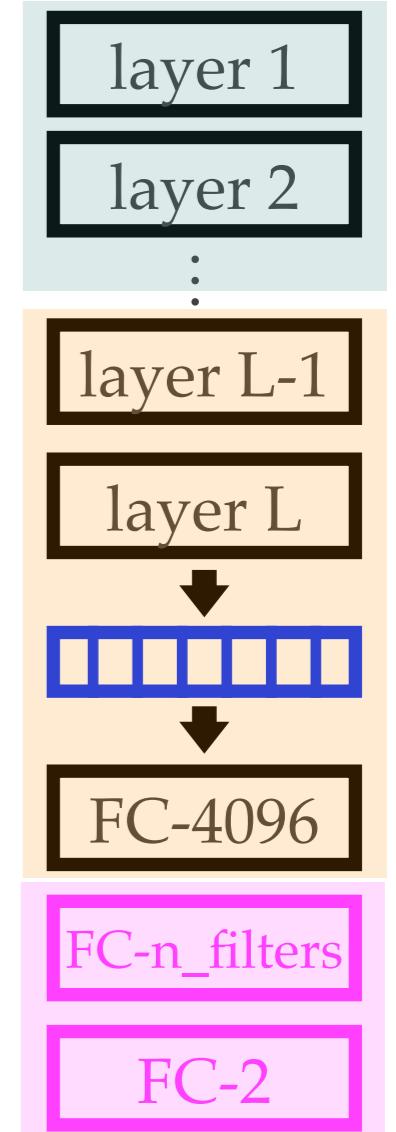
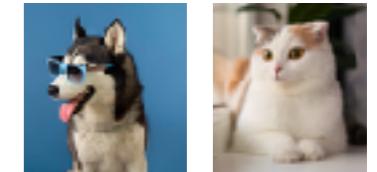
Clone these layers along with their weights and do not modify the weights

Clone these layers along with their weights and retrain them with an initially low learning rate

Replace these layers and train them

Dog-vs-cat (2 classes)

Assume dataset of moderate size



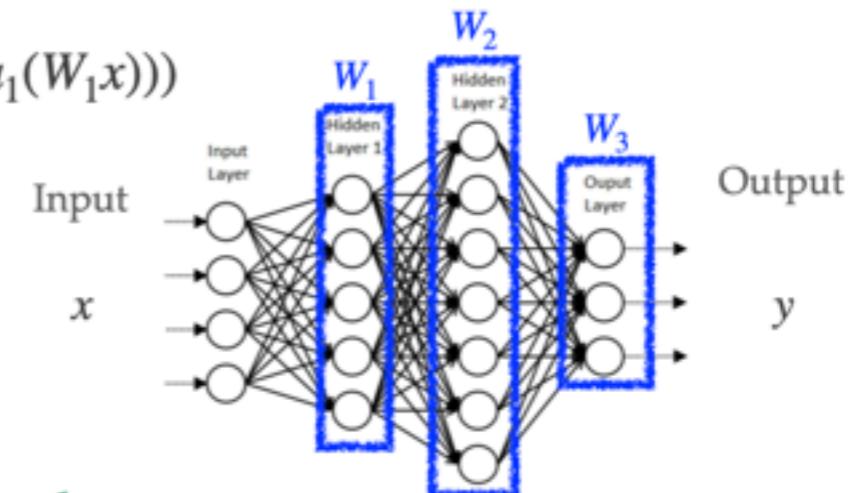
Clone these layers along with their weights and retrain them with an initially low learning rate

Replace these layers and train them

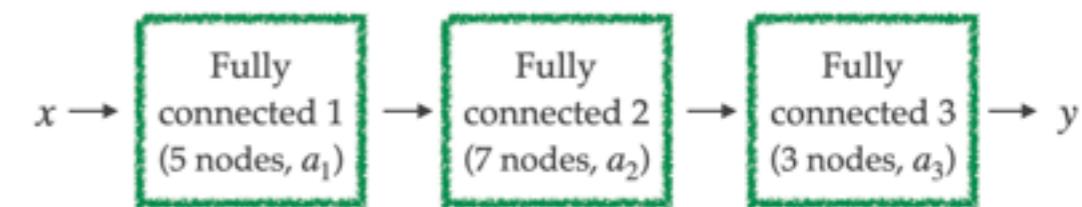
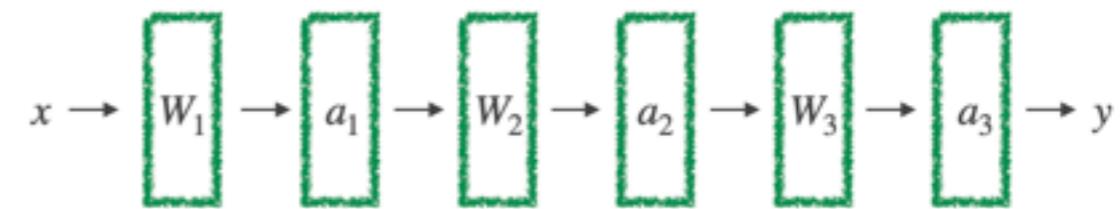
Summary

- ❖ Deep Learning Components
 - ❖ Fully connected layer
 - ❖ Computation graph
 - ❖ Activation functions
 - ❖ Loss function
 - ❖ Model optimization
 - ❖ Gradient descent
 - ❖ Backpropagation
 - ❖ Stochastic gradient descent (SGD)
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)

$$y = a_3(W_3 a_2(W_2 a_1(W_1 x)))$$



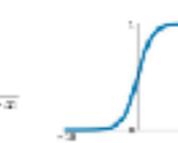
Computation graph



Sigmoid
 $\sigma(x) = \frac{1}{1+e^{-x}}$

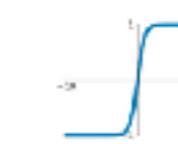
tanh
 $\tanh(x)$

ReLU
 $\max(0, x)$

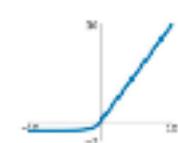


Leaky ReLU
 $\max(0.1x, x)$

Maxout
 $\max(w_1^T x + b_1, w_2^T x + b_2)$

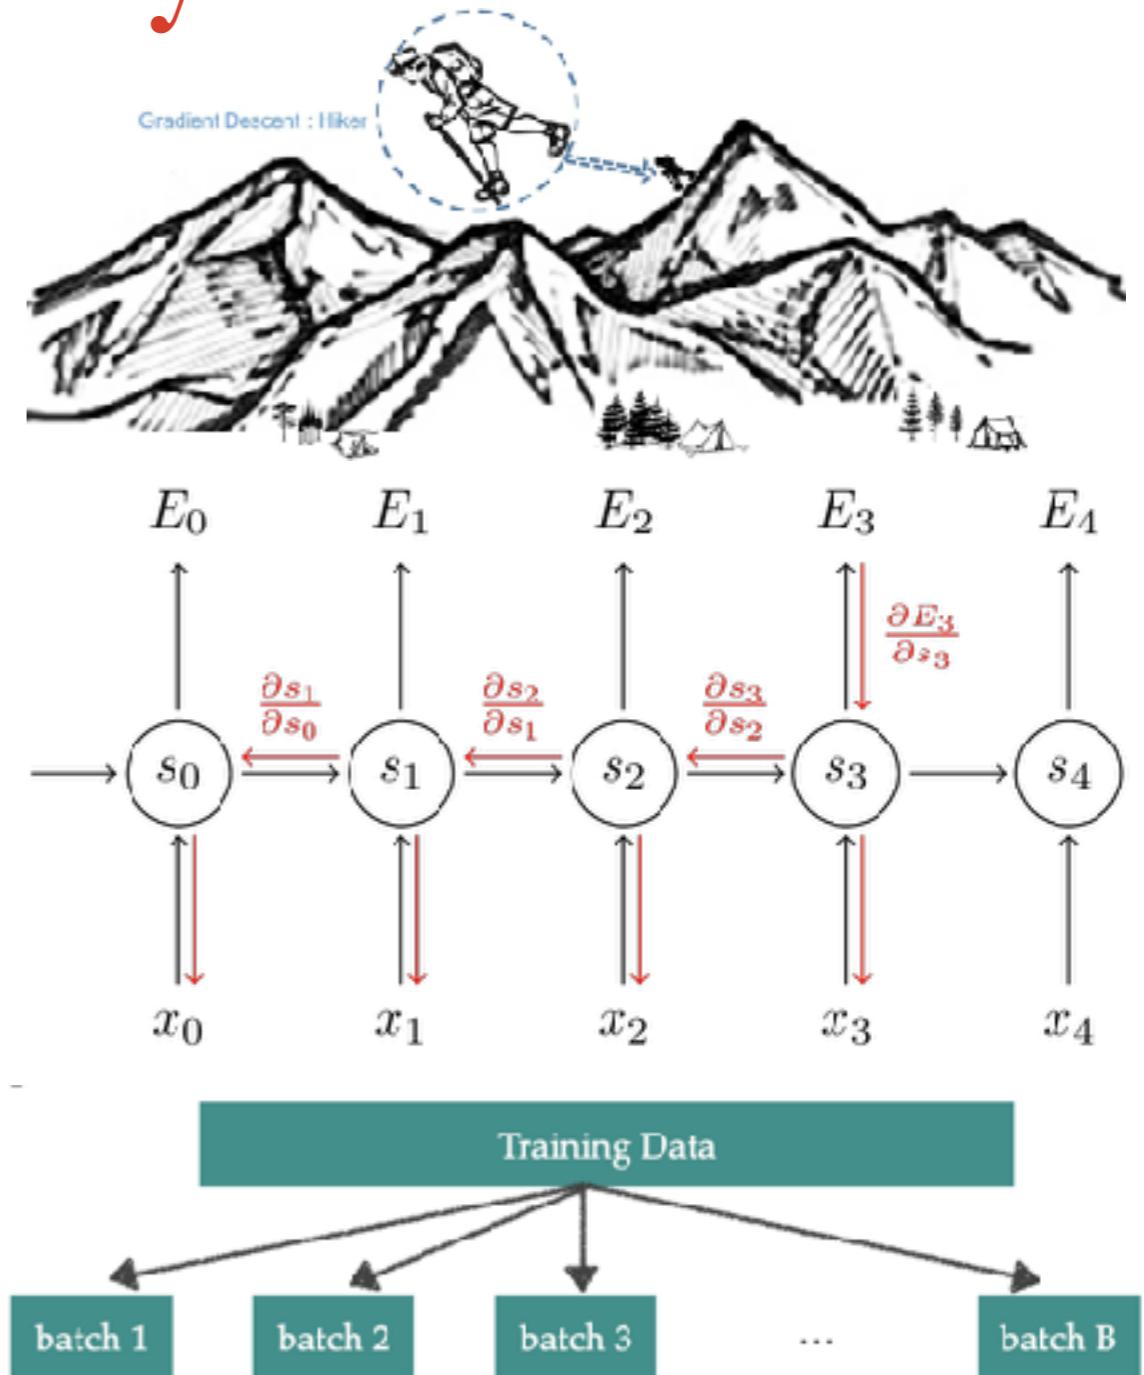


ELU
 $\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$



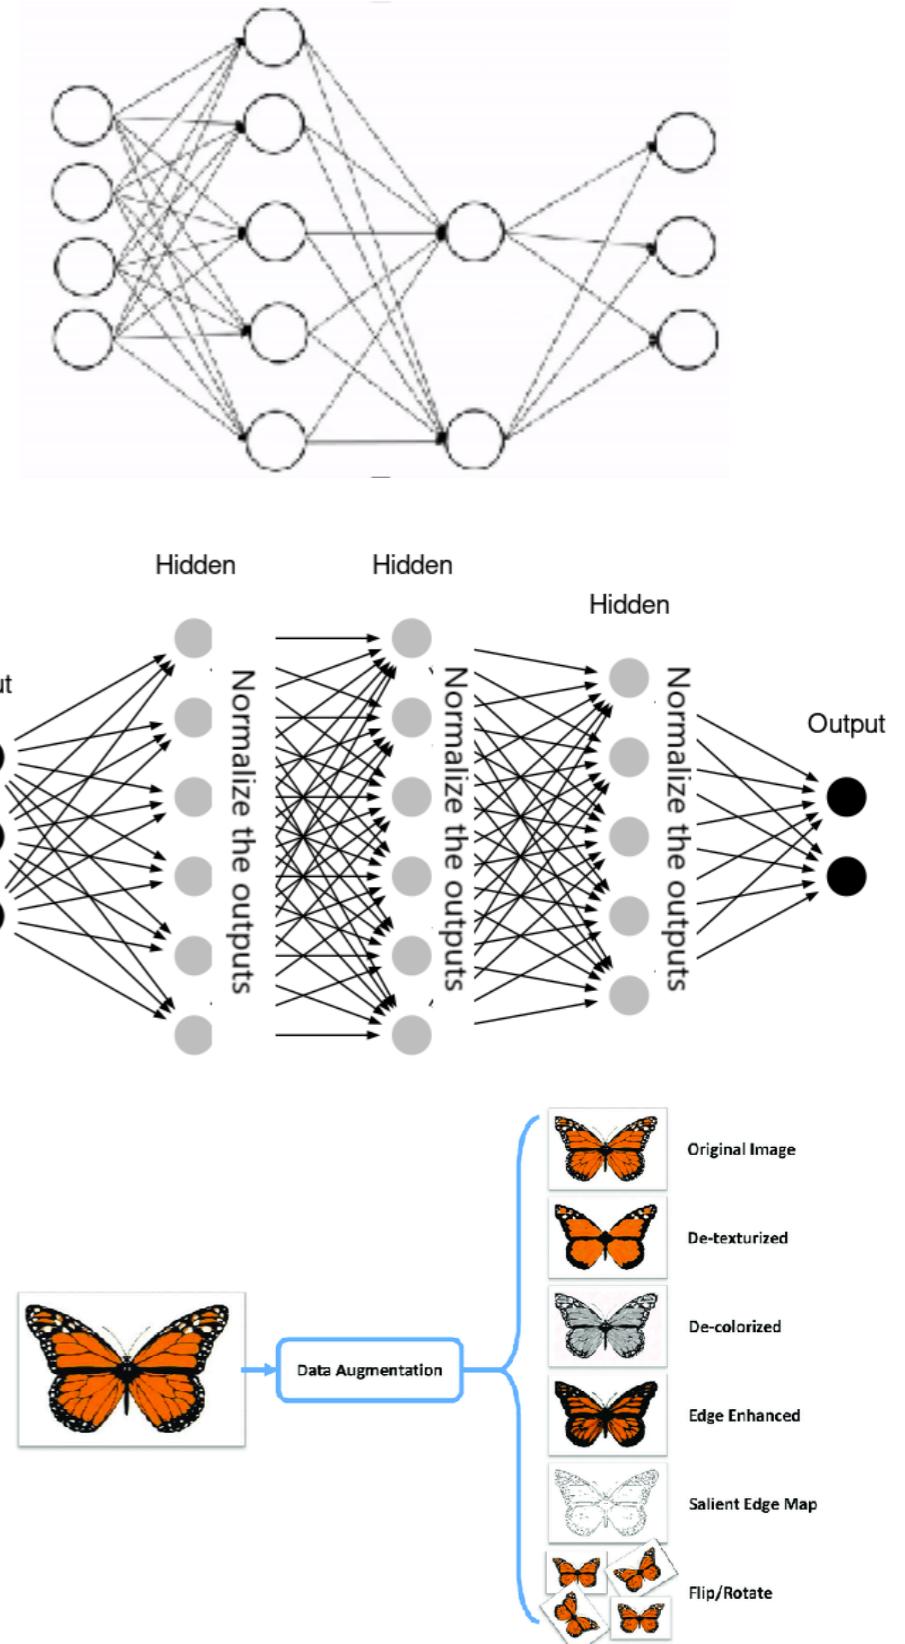
Summary

- ❖ Deep Learning Components
 - ❖ Fully connected layer
 - ❖ Computation graph
 - ❖ Activation functions
 - ❖ Loss function
 - ❖ Model optimization
 - ❖ Gradient descent
 - ❖ Backpropagation
 - ❖ Stochastic gradient descent (SGD)
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)



Summary

- ❖ Deep Learning Components
 - ❖ Fully connected layer
 - ❖ Computation graph
 - ❖ Activation functions
 - ❖ Loss function
 - ❖ Model optimization
 - ❖ Gradient descent
 - ❖ Backpropagation
 - ❖ Stochastic gradient descent (SGD)
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)

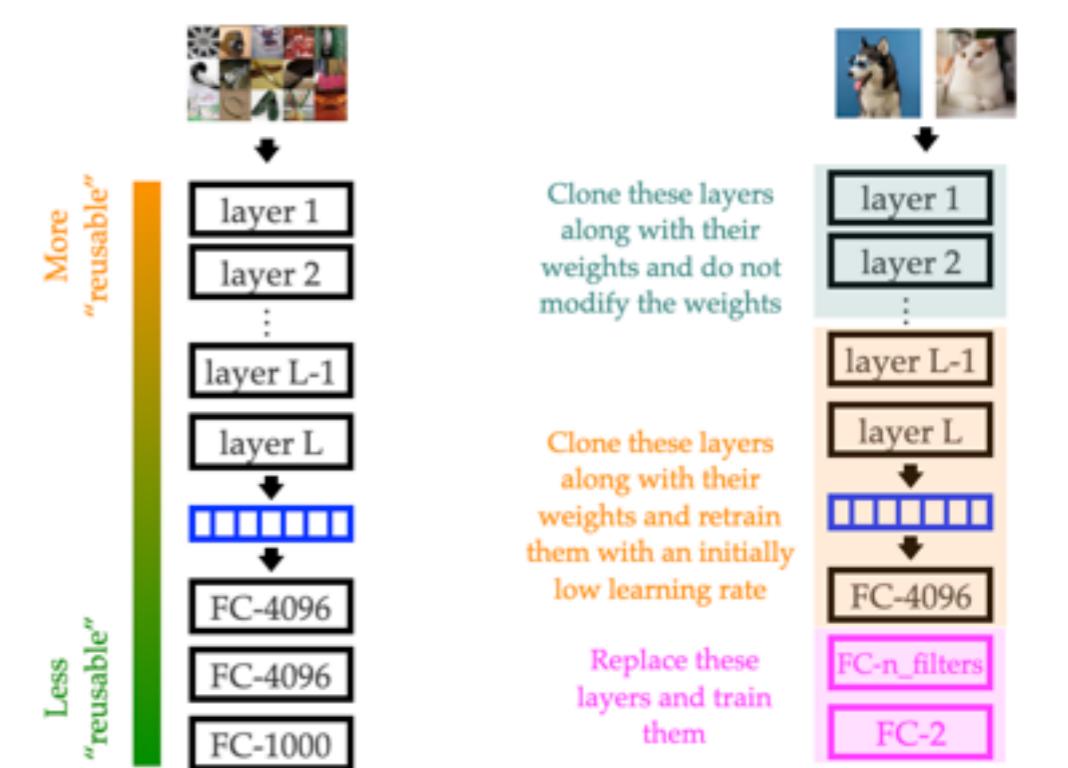
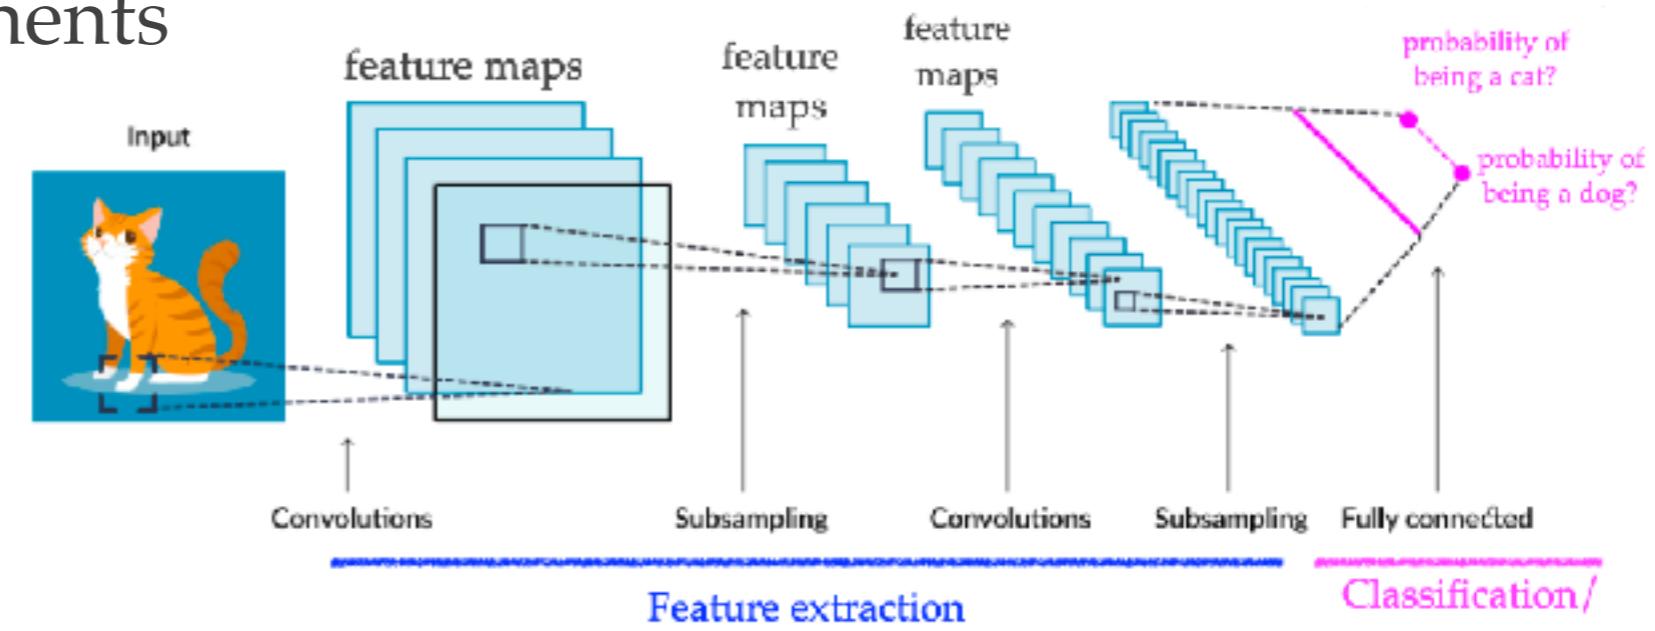


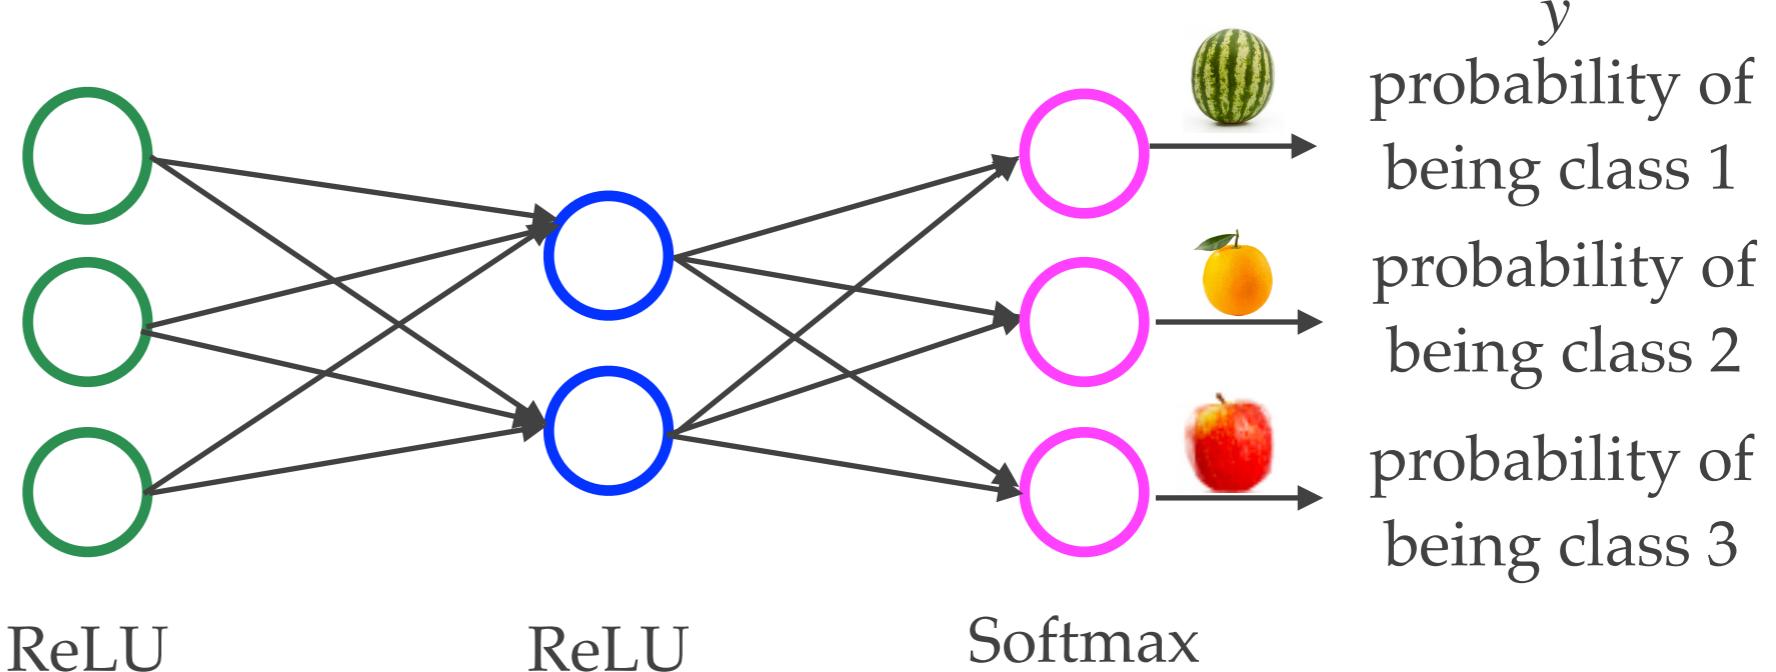
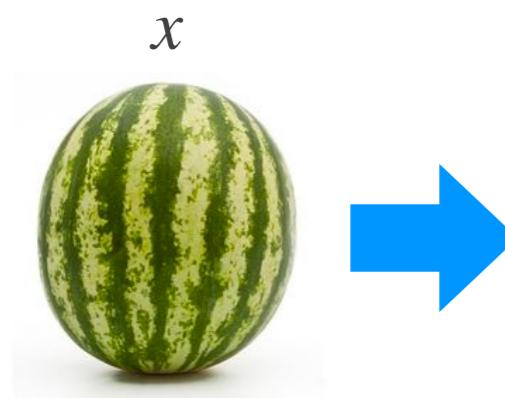
Summary

❖ Deep Learning Components

- ❖ Fully connected layer
- ❖ Computation graph
- ❖ Activation functions
- ❖ Loss function
- ❖ Model optimization
 - ❖ Gradient descent
 - ❖ Backpropagation
 - ❖ Stochastic gradient descent (SGD)
- ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
- ❖ Convolutional layer
- ❖ Pooling layer

❖ Convolutional Neural Network (CNN)





activation functions

ReLU

ReLU

Softmax

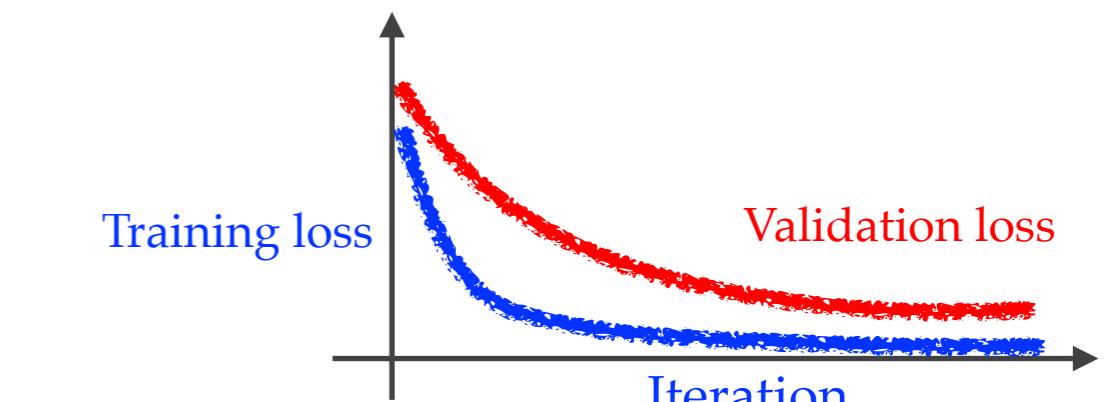
y
probability of
being class 1
probability of
being class 2
probability of
being class 3

```
# Import necessary modules
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

# Create the model
model = keras.Sequential()
model.add(layers.Dense(3, activation="relu"))
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(3, activation="softmax"))

# Compile the model
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(),
)

# Train the model for 100 epochs with a batch size of 32
model.fit(x_train, y_train, batch_size=32, epochs=100, validation_data=(x_val,y_val))
```



Model

Loss function
and optimizer

Resources

Conferences

arXiv

Attention Is All You Need

Cybersecurity

Deep Residual Learning For Image Recognition

Cornell University

arXiv

COVID-19 Quick Links

We gratefully acknowledge support from the Simons Foundation and member institutions.

Login

Search... All fields Search

Help | Advanced Search

Courses

CS231n Home Course Schedule Office Hours Final Projects 9d

CS231n: Deep Learning for Computer Vision



Course Description

Computer vision has become ubiquitous in our society with applications in search, image understanding, object mapping, medicine, games, and self-driving cars. Over 10 million of these applications are visual recognition tasks, such as image classification, localization and detection. Recent developments in neural networks (aka "deep learning") approaches have greatly advanced the performance of these state-of-the-art visual recognition systems. This course is a deep dive into the details of deep learning architectures with a focus on learning end-to-end models. For these goals, particularly image classification, during the 11 weeks novice students will learn how to implement and train their own neural networks and gain a detailed understanding of cutting-edge research in computer vision. Additionally, the final assignment will give them the opportunity to train and apply multi-million parameter networks on real-world vision problems of their choice. Through multiple hands-on assignments and the final course project, students will acquire the toolkit for setting up deep learning codes and practical engineering tricks for training and fine-tuning deep neural networks.

INTRODUCTION

2110573 Pattern Recognition



Pattern Recognition 2019 - II (W) Introduction
last updated: 11/18/2019 11:58:00 AM

- 1.00 news. Streamed on Jan 13, 2020 0:00 pre-lesson.
- 1.000 Introduction & course logistics
- 1.001 Technology trends and machine learning
- 1.002 Machine learning technologies
- 1.003
- 1.004 Introduction to supervised learning
- 1.005 Typical workflow in machine learning
- 1.014 Feature extraction
- 1.015 Evaluation & metrics
- 1.016 Implementing convolutional neural network (convolutional layers)
- 1.017 class exercise

Textbooks

Journals

Blog posts

YouTube videos



Pattern 2022

EkapolC

20 videos 6,456 views Last updated on Sep 12, 2022

Play all Shuffle

Machine Learning workshops for material scientists

Lecture 1: Statistics review and warm up

October 5, 2022

Sira Sriwasdi, PhD

- Research Affairs
- Center of Excellence in Computational Molecular Biology (CCMB)
- Center for Artificial Intelligence in Medicine (C3-AIM)

MTEC Machine Learning Workshop

Sira Sriwasdi

18 videos 169 views Last updated on Dec 15, 2022

Play all Shuffle

Machine learning workshop series hosted by MTEC that I and colleague (@ItthiC) gave in 2022.

Contents are in Thai.

ເງິນຕົກສອບເຖິງຂ້າກັນຄວາມຮູ່ແລະເຫດຜົນດ່າງຈາກທາງດ້ານການຮັບຮູ່ຂອງເຄື່ອງ ພຣັດທັວຂ່າງຄໍາສົ່ງການໄພທອນທີ່ມີມແລະເພື່ອນ (@ItthiC) ກັບທີ່ມີນັ້ອງຜູ້ຂ່າຍສອນຮ້ານກັນບຽນຢ່າງໃຈຈາກທີ່ສັນສົນໂດຍ MTEC ໃນປີ

Introduction to Machine Learning and Deep Learning

Itthi Chatnuntawech

Nanoinformatics and Artificial Intelligence (NAI)
National Nanotechnology Center (NANOTEC)
National Science and Technology Development Agency (NSTDA)

September 5, 2022

MTEC Machine Learning Mini-Lecture

by Sira Sriwasdi

5 videos 193 views Last updated on Dec 15, 2022

Play all Shuffle

Introduction to machine learning mini-lecture series hosted by MTEC that I and colleague (@ItthiC) gave in September 2022.

Contents are in Thai.

ເຄື່ອງສົ່ງ ເຖິງຂ້າກັນແນວທີ່ກິ່ນຮູ່ານ ແຫດນິກ ແລະກາຮອກແນບຄໍາຄາມທາງດ້ານການຮັບຮູ່ຂອງເຄື່ອງ ທີ່ມີມແລະເພື່ອນ (@ItthiC) ຈຳກັນບຽນຢ່າງໃຈທີ່ສັນສົນໂດຍ MTEC ເຊື້ອນກັນຫອຍນ 2565

Outline

- ❖ Deep Learning Components
 - ❖ Regularization
 - ❖ ℓ_2 and ℓ_1 regularization
 - ❖ Dropout
 - ❖ Batch normalization
 - ❖ Convolutional layer
 - ❖ Pooling layer
- ❖ Convolutional Neural Network (CNN)
- ❖ Hands-on: MRI Tumor Classification Using CNN
 - ❖ <https://github.com/ichatnun/CMU-medical-imaging-deep-learning>